

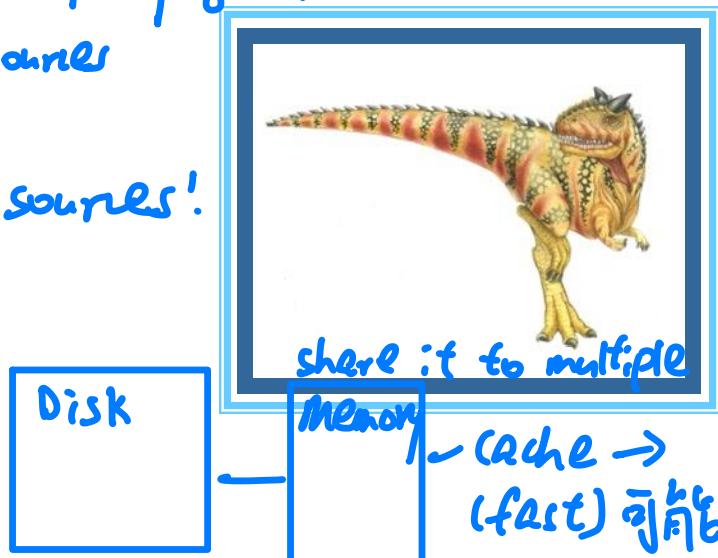
# Spring 2025

# COMP 3511 Operating Systems



One of the hardest ?<

- security
- enable multiple program
- manage resources
- allocate
- hardware sources!



more convenient,

OS  
↓  
Do CPU  
Scheduling  
interface between  
users and hardware

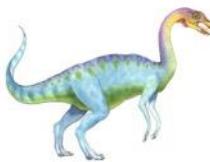


# Lectures and Labs/Tutorials

---

- **Lectures (3 February – 10 May 2025):** Basic concepts
  - L1 Monday 01:30PM - 02:50PM, Friday 09:00AM - 10:20AM, LTG
  - L2 Wednesday, Friday 01:30PM - 02:50PM, LTL-CYT
- **Lab Tutorials**
  - LA1 Thursday, 12:30PM - 02:20PM, LTG
  - LA2 Tuesday, 06:00PM - 07:50PM, LTC
- **Course Website:** <https://course.cse.ust.hk/comp3511/>
- **Instructors:** Bo Li (L1) and Kai CHEN (L2)

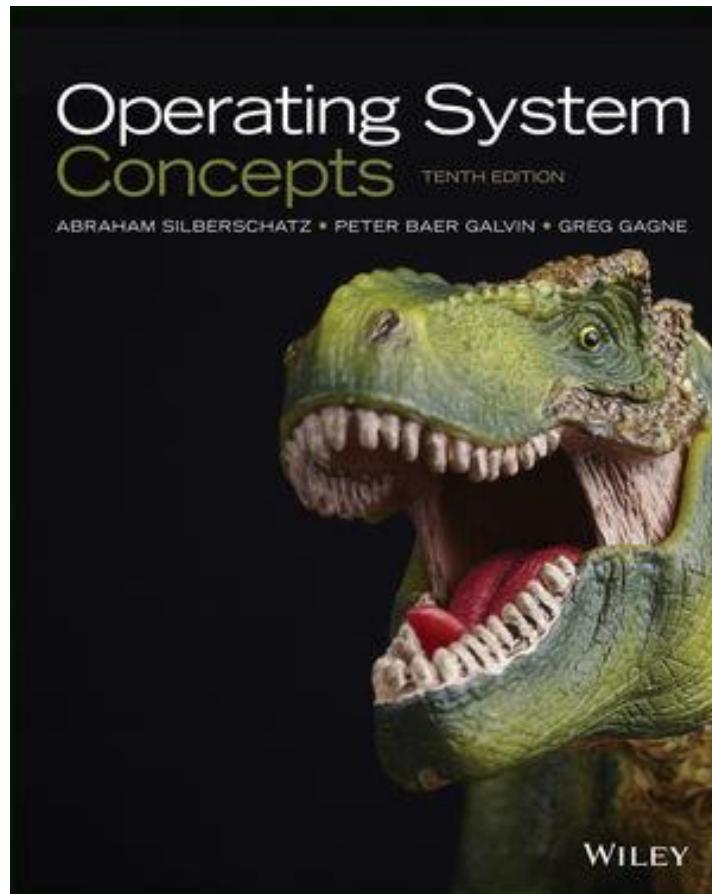


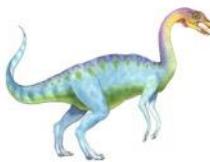


# Textbook

---

- **Operating System Concepts, A. Silberschatz, P. B. Galvin and G. Gagne, 10th Edition**

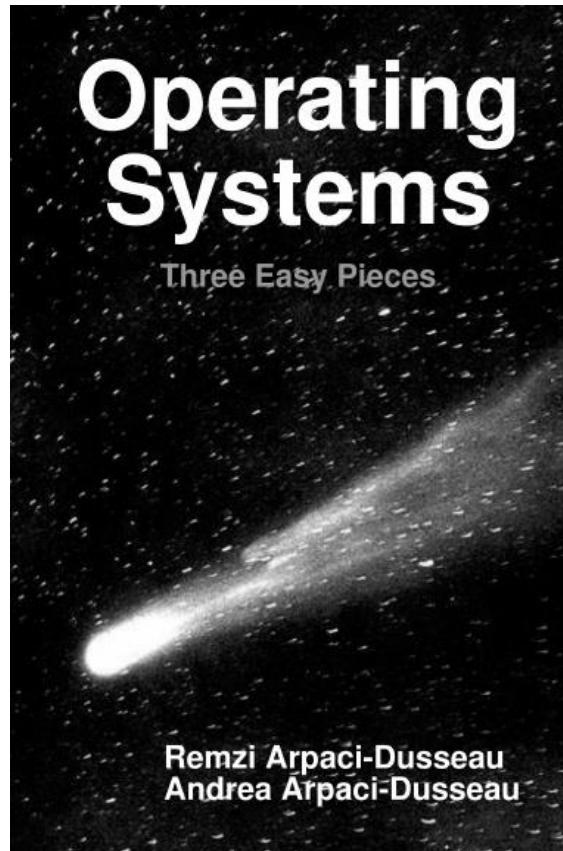


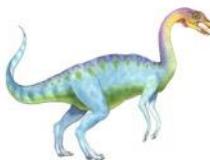


# Reference Book

---

- Operating Systems: Three Easy Pieces, Remzi H. Arpaci-Dusseau and Andrea C. Arpaci-Dusseau
- Online (free access): <http://pages.cs.wisc.edu/~remzi/OSTEP/#book-chapters>



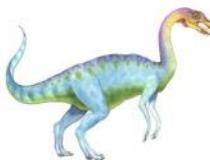


# Course Prerequisite

---

- **COMP 2611 or ELEC 2300 or ELEC 2350 (Computer Organization)**
    - Computer organization – von Neumann machine, CPU, pipelining, caching, memory hierarchy, I/O systems, interrupt, storage and hard drives
  - **COMP 2011 or COMP 2012H (C programming)**
    - UNIX/Linux basic
    - Programming requirement - C programming
- ↓  
*fundamental computer architecture*
- ↓  
*load / store in I/O*





# Labs and Tutorials

---

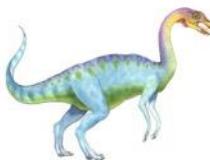
- **9 Labs and Tutorials** – Tentative schedule subject to lecture progress

- No Labs on week 1
- Lab #1 (week 2): Introduction to Linux
- Lab #2 (week 3): C/C++ programming
- Lab #3 (week 4): Linux process, pipe(), and Project #1
- Lab #4 (week 5): Review
- Lab #5 (week 6): Project #2
- Lab #6 (week 7): Review
- No Labs on week 9 (Midterm week)
- Lab #7 (week 10): Review
- Lab #8 (week 11): Project #3
- No Labs on week 12 (affected by Labor Day)
- Lab #9 (week 13) Review

↗ Easy

harder





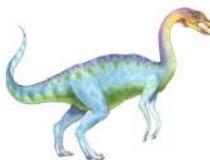
# Grading Scheme

---

- **4 Homework - written assignments – 20% (5% each)**
  - HW #1 (week 2-4)
  - HW #2 (week 5-7)
  - HW #3 (week 8-10)
  - HW #4 (week 11-13)
- **3 Projects - programming assignments – 30%**
  - Project #1 (week 4-6) (10%)
  - Project #2 (week 7-9) (10%)
  - Project #3 (week 10 -12) (10%)
- **Midterm Exam (week #9) - 20%**
- **Final Exam - 30%**

Apply knowledge!



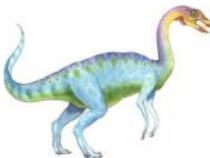


# Plagiarism Policy

---

- There are differences between collaborations, discussions and copy!
  
- First time: all involved get ZERO marks, and will be reported to ARR
- Second time: need to terminate (**Fail** grade)
- Any cheating in midterm or final exam results in **automatic Fail** grade
- Reference: [HKUST Academic Honor Code and Academic Integrity](#)





# Lecture Format

---

- Lectures:
  - Lecture notes are made available before lectures
- Tutorials and Labs
  - Unix environment, editor (vim), compile and run programs, Makefile
  - C++ and C programming basic
  - Tutorials on programming assignments
  - C programming APIs and interfaces
  - Supplement materials with more examples and exercises
- Reading the corresponding materials in the textbook and reference book
  - **Lecture notes do not and can not cover everything**
- Chapter Summaries
  - Comprehensive summary for each chapter

*Read textbook*



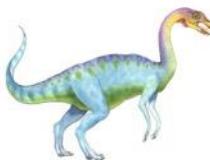


# Assignments

---

- Written assignments
  - Due by time specified
  - Contact the corresponding TA for any disputes on the grading
  - Regrading requests be granted within **two weeks** after the homework grades are released
  - Late policy: **10% reduction, only one day delay is allowed**
  
- Programming assignments - **individual project**
  - Due by time specified
  - **Run on a CS Lab 2 Linux Machines**
  - Submit it using Canvas
  - Regrading requests be granted within **two weeks** after the grades are released
  - Late policy : **10% reduction, only one day delay is allowed**



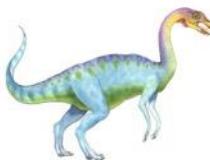


# Midterm and Final Examinations

---

- **Midterm Exam**
    - Time: 11 April (Friday) (week #9) 7:00 pm – 9:00 pm
    - Venues: LT-A and Room 2304
  - **Final Exam**
    - TBD
- No make-up exam*
- **All exams are open-book and open-notes (hard copies)**
    - **NO electronic devices are allowed**
  - **No make-up exams will be given unless**
    - Under special circumstances, e.g., sickness, with **letters of proof**
    - The instructor must be informed **before the exam**



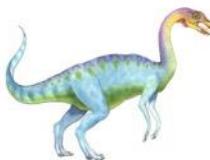


# Tips for Learning

---

- Attend lectures and lab tutorials
  - Download lecture/lab notes prior to lectures
  - Important concepts are explained, with examples
- Complete homework and projects independently
  - This is to test your knowledge and how much you comprehend
- **Spend 30 minutes or so each week to review the content**
  - Chapter summary helps
  - This can save you lots of time later when you prepare for exams
  - You can not expect to learn everything 2-3 days before exams
  - Knowledge is accumulated incrementally
- Start your project earlier
  - Have a plan for the project
- **Raise questions during or after lectures !**
  - Do not delay your questions until close to the exams





# What you are supposed to learn

- Define the fundamental principles, strategies and algorithms used in the design and implementation of operating systems
- Analyze and evaluate operating system functions
- Understand the basic structure of an operating system **kernel**, and identify the relationship between the various subsystems
- Identify the typical events, alerts, and symptoms indicating potential operating system problems
- Design and implement programs for basic operating system functions and algorithms
- **Advanced OS course – COMP 4511 System and Kernel Programming in Linux**

*definition?*

*no universe definition*

*next year*

*Manage and schedule the resources*

*some algorithm, useful but not practical*

*OS can have different goals  
e.g. lower power consumption*





# Course Outline

## □ Overview (4 lectures)

- Basic OS concept (2 lectures)

## AAA □ System architecture (2 lectures)

## Process and Thread (12 lectures)

- Process and thread (4 lectures)

- CPU scheduling (4 lectures)

- Synchronization and synchronization examples (2 lectures)

- Deadlock (2 lectures)

## Memory and storage <sup>↗ 4 weeks</sup> (8 lectures)

- Memory management (2 lectures)

- Virtual memory (3 lectures)

- Secondary storage (1 lectures)

- File systems and implementation (2 lectures)

## Protection (1 lectures)

- Protection (1 lecture)

- Security (1 lecture) - optional

*very important!!!*

*Algorithms and programs*

*Different jobs*

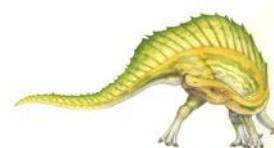
*Do sth within  
deadlines*

*Fit program t.t.  
不能违例!*

*Move the disk to the memory*

*How to allocate?*

*if have time  
(not main focus)*





# Course Coverage

---

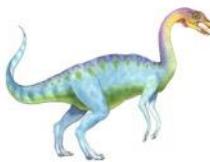
## □ Overview

- [Chapter 1](#) – high-level description of OS, basic components in computer systems including multi-processor and parallel systems, virtualization
- [Chapter 2](#) – OS services including APIs and system calls, and common OS design approaches (monolithic, layered, microkernel, modular)

## □ Process and Thread

- [Chapter 3](#) (Process) – concept of a process capturing a program execution, creating and terminating a process, process communications
- [Chapter 4](#) (Thread) – concept of a thread and multi-threaded process for concurrent execution of a program
- [Chapter 5](#) (CPU scheduling) – A variety of common CPU scheduling algorithms including real-time scheduling, and issues associated with multiprocessor scheduling and thread scheduling
- [Chapter 6-7](#) (Synchronization) – critical section problem, synchronization tools (hardware and software), and synchronization examples
- [Chapter 8](#) (Deadlock) – deadlock characterization, resource allocation graph, deadlock prevention, avoidance and detection algorithms





# Course Coverage (Cont.)

## □ Memory and Storage

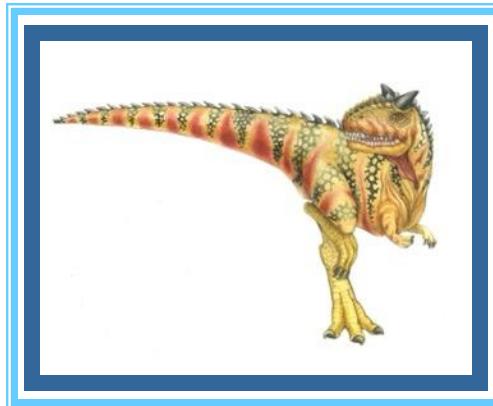
- Chapter 9 (Memory) - contiguous memory allocation, segmentation, paging including hierarchical paging
- Chapter 10 (Virtual memory) – virtual vs. physical memory, demand paging, page replacement algorithm, thrashing and frame allocation
- Chapter 11 (Secondary storage) – hard drive, disk structure, disk scheduling algorithms and RAID (disk array) structure
- Chapter 13-14 (File systems) – file access methods, directory structure and implementation, basic file system data structure (on-disk and in-memory), disk space management including disk block allocation

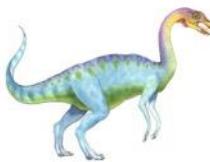
## □ Protection

- Chapter 17 (Protection) – basic protection principles, protection rings, protection domain and implementation (access matrix)
- Chapter 16 optional (Security) - security threats and attacks, countermeasures to security attacks



# Chapter 1: Introduction





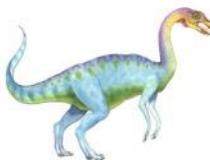
# Chapter 1: Introduction

---

- What Operating Systems Do
- Computer System Organization and Architecture
- Multiprocessor and Parallel Systems
- Definition of Operating Systems
- Virtualization and Cloud Computing
- Free and Open-Source Operating Systems

Apple, Windows close-source



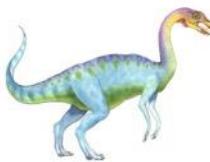


# Objectives

---

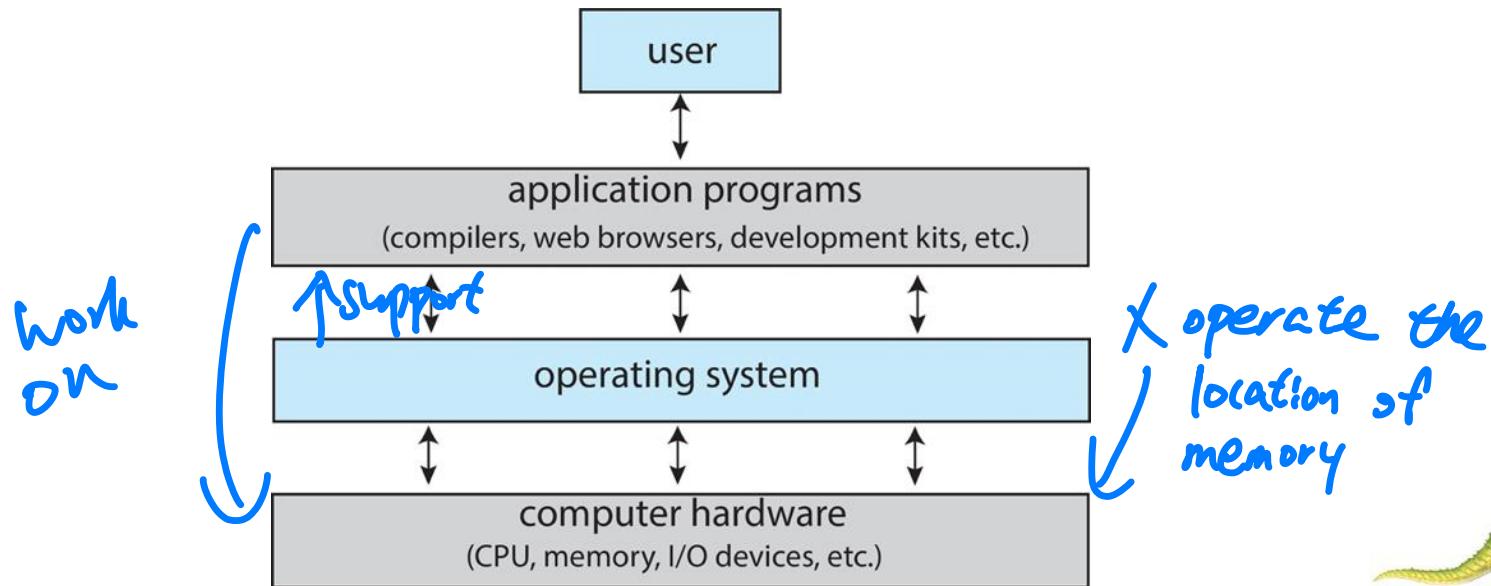
- Describe the general organization of a computer system and the role of interrupts. *多任务处理，如何切换！*
- Illustrate the components in a modern multiprocessor computer system.
- Discuss how operating systems are used in various computing environments
- Provide examples of free and open-source operating systems





# What is an Operating System?

- **Users** - people, machines, other computers or devices
- **Application programs** – define the ways how system resources are used to solve user problems
  - Editors, compilers, web browsers, database, video games, etc.
- **Operating system** – controls and coordinates use of computing resources among various applications and among different users
- **Hardware** – basic computing resources, CPU, memory, I/O devices



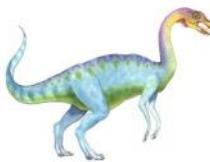


# What is an Operating System?

- OS is a **program** (extremely complex) that acts as an intermediary between users or applications and computer hardware
  - Microsoft window, MacOS, iOS, Android, Linux ...
- Operating system goals: *Multiple* *(Example: Graphics)*
  - Execute user programs and make solving user problems easier
  - Make the computer system convenient to use
  - Manage and use the computer hardware in an efficient manner
- **User view**
  - Convenience, ease of use, good performance and security
  - Users do not care about resource utilization, efficiency
- **System view** *利用*
  - OS as a resource allocator and a control program

*use the memory (resource)  
finish more jobs, to do more efficiency!*





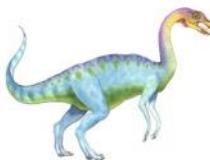
# What Operating Systems Do

*Goal*

- It depends on the point of view (user or system) and **target devices**
- Shared computers such as **mainframe** or **minicomputer** *Share time, resources*
  - OS tries to keep all users satisfied – performance vs. fairness
- Individual systems (e.g., **workstations**) have dedicated resources,
  - performance rather fairness, may use shared resources from **servers**
- Mobile devices (e.g., smartphones and handheld devices) are resource constrained
  - Target specific user interfaces such as **touch screen**, voice control such as Apple's **Siri**, and optimized for usability and battery life *os target*
- Computers or computing devices **with little or no user interface**
  - **Embedded systems** - present within home devices (AC, toasters), automobiles, ships, spacecraft, run real-time operating systems
  - Designed to run primarily without user intervention – some may have numeric keypads and indicator lights to show status

*Depending the computer devices!*





# Operating System Definition

---

- There is no universally accepted definition on OS
  - “Everything a vendor ships when you order an operating system” is a good approximation, but it varies a great deal
- OS is a resource allocator
  - Manages all resources – hardware and software
  - Decides between conflicting requests for efficient and fair resource use
- OS is a control program → *控制多台计算机，其他人随便用！*
  - Controls execution of programs, prevent errors and improper use of the computer
- In a nutshell, OS manages and controls hardware and helps to facilitate programs to run on computers.





# Operating System Definition

- **Kernel** *Always running, as long as the computer is on!*
  - “The one program that is running at all time on a computer”
  - The essential functionalities - discussed in this introductory course
    - ↳ how to allocate CPU time
    - ↳ how to do resource allocation
    - ↳ how to do processor communication
- **Middleware** *Develop different tools*
  - A set of software frameworks that provide additional services to application developers such as databases, multimedia, graphics
  - Popular in mobile OSes - Apple's iOS and Google's Android
- Everything else
  - **System programs** (ships with the operating system, but not part of the kernel), such as word processors, browsers, compilers
  - **Application programs**, not associated with the operating system – e.g., apps from appstores *→ not coming from os!*
- OS includes the always running kernel, middleware frameworks that ease application development and provide additional features, as well as **system programs** that aid in managing the system while it is running

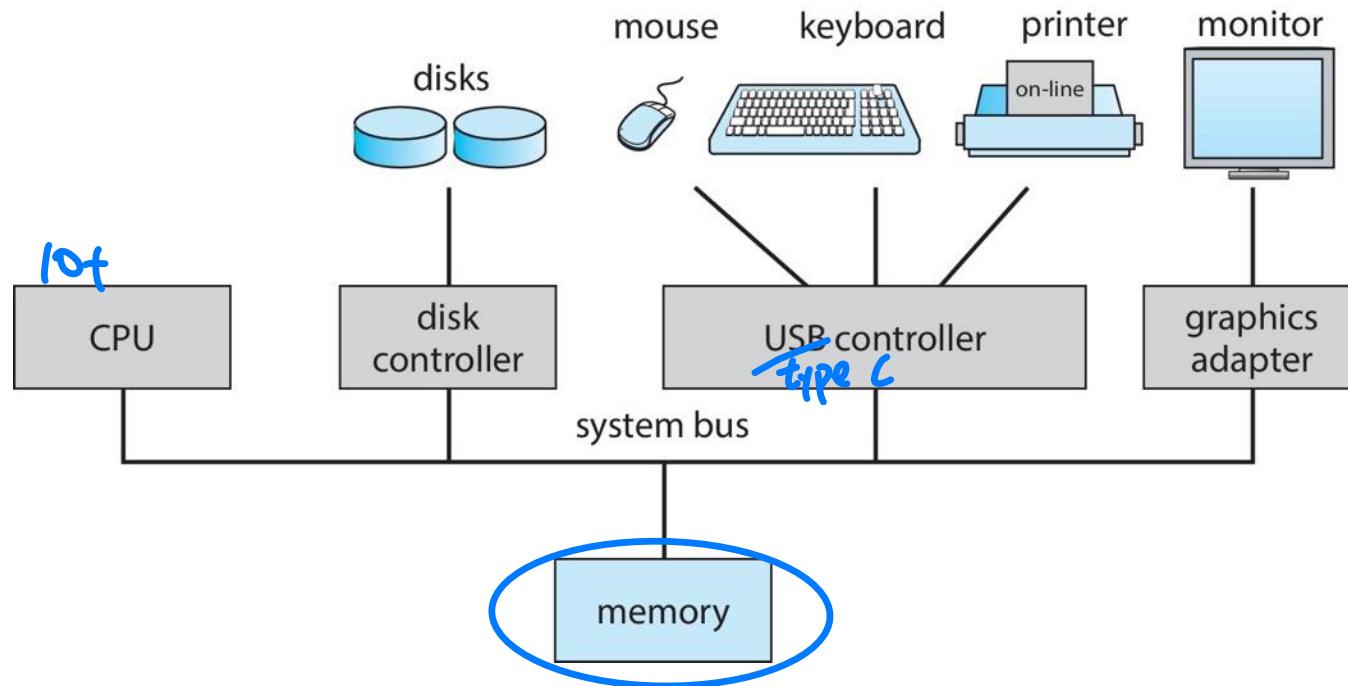




# Computer System Organization

## □ Computer-system operation

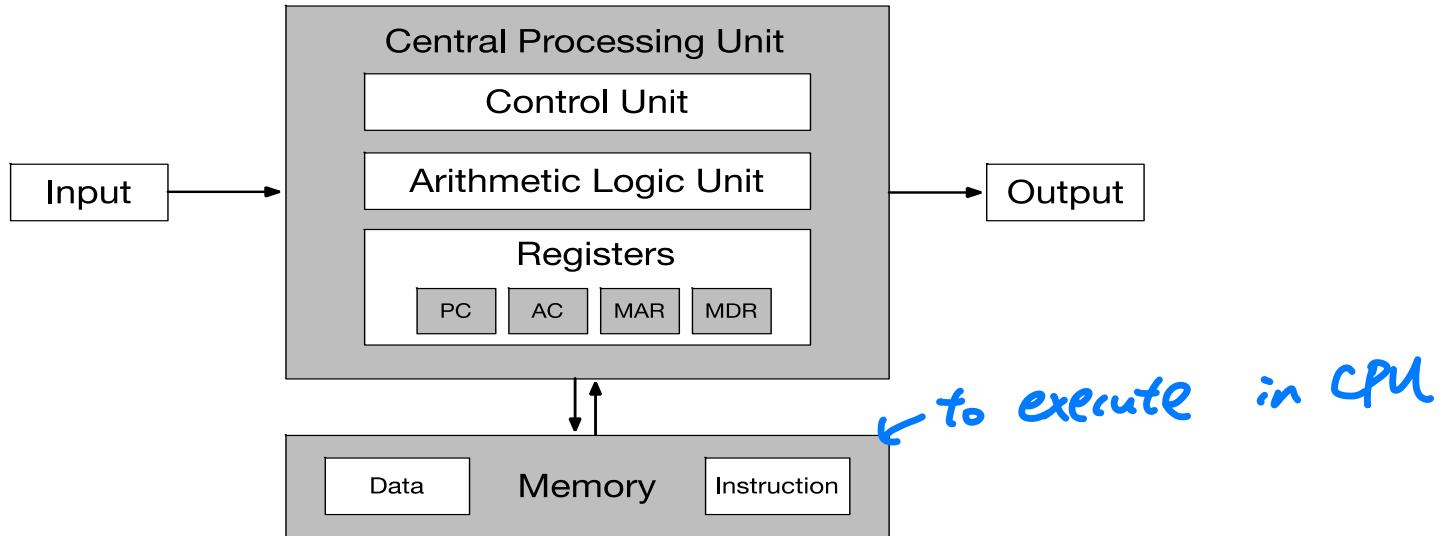
- One or more CPU cores, device controllers connected through common bus providing access to **shared memory**
- **Concurrent execution** of CPUs and devices - competing for memory cycles through **shared bus**





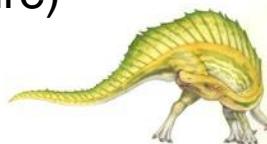
# Learn in Comp 2611

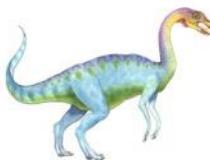
## A von Neumann Architecture



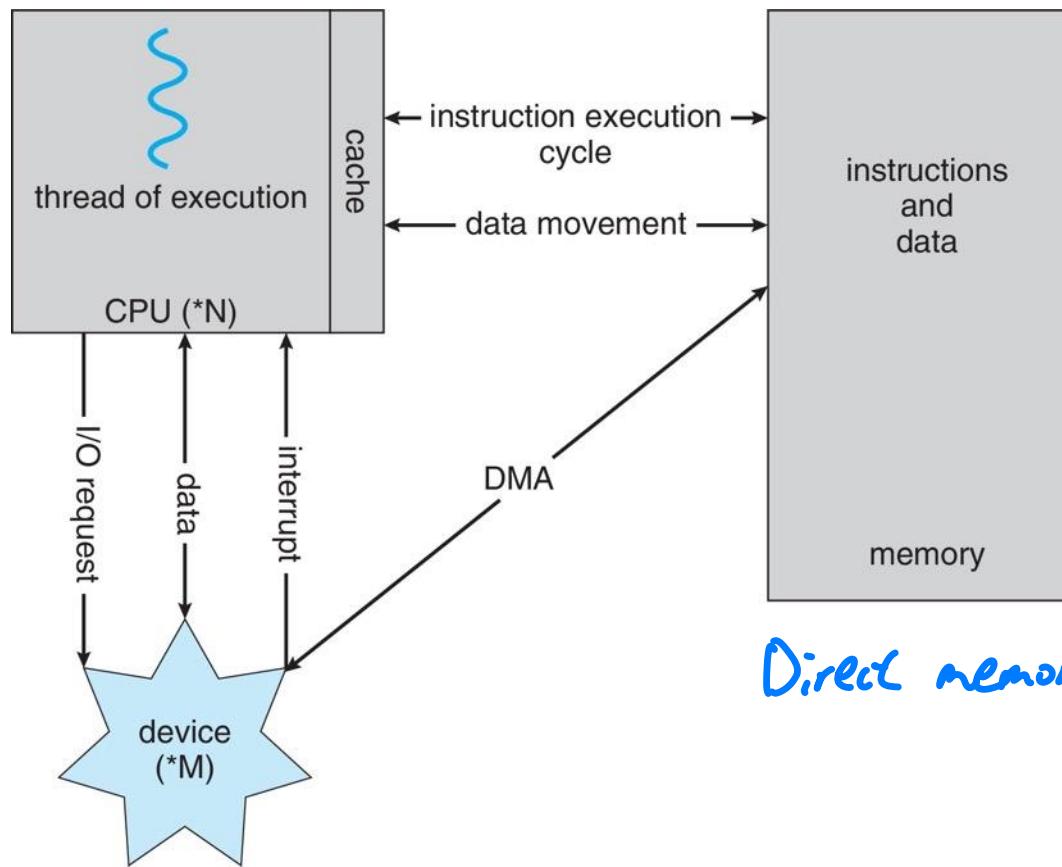
- A central processing unit that contains an arithmetic logic unit (ALU) and processor registers – Program Counter (PC), Accumulator (AC), Memory Address Register (MAR), Memory Data Register (MDR)
- A control unit that contains an instruction register (IR) and program counter (PC)
- Memory stores data and instructions – along with caches
- External mass storage – secondary storage (not shown in the figure)
- Input and output mechanisms

*I/O devices!*





# How a Modern Computer Works

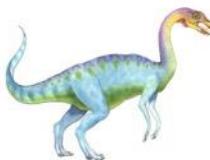


Steps in executing an instruction:

- Fetch instruction
- Decode instruction
- Fetch data
- Execute instruction
- Write back if any

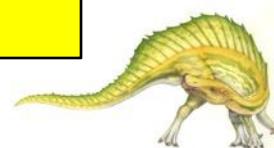
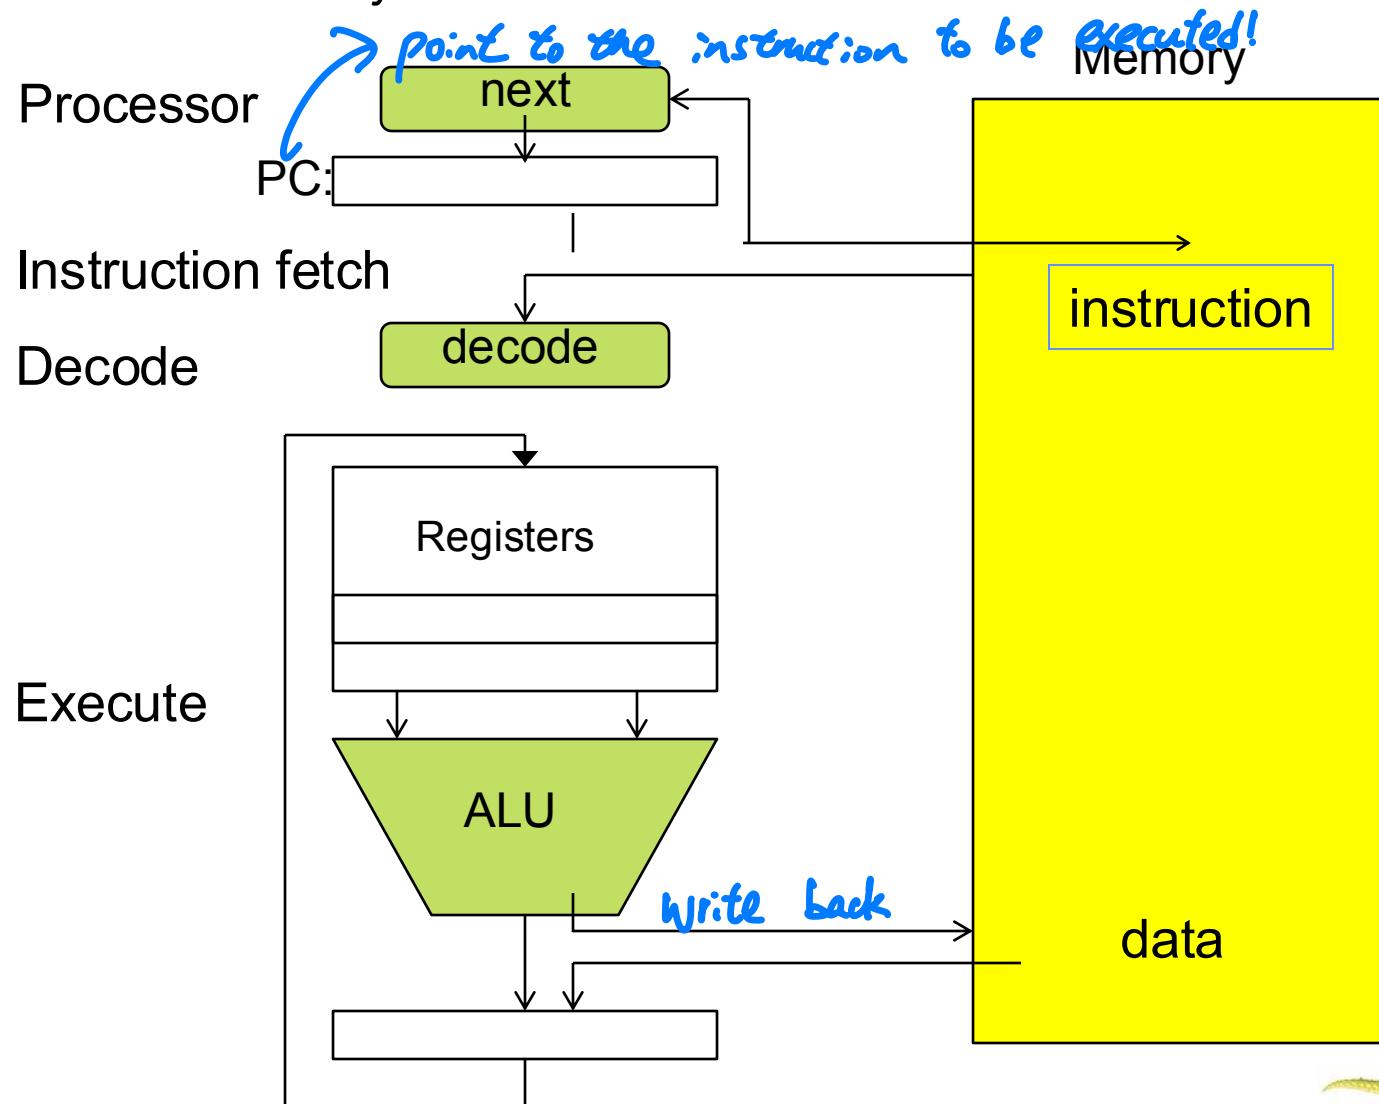
The von Neumann architecture





# Instruction Fetch/Decode/Execute

The instruction cycle



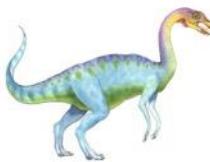


# Computer-System Operation – I/O

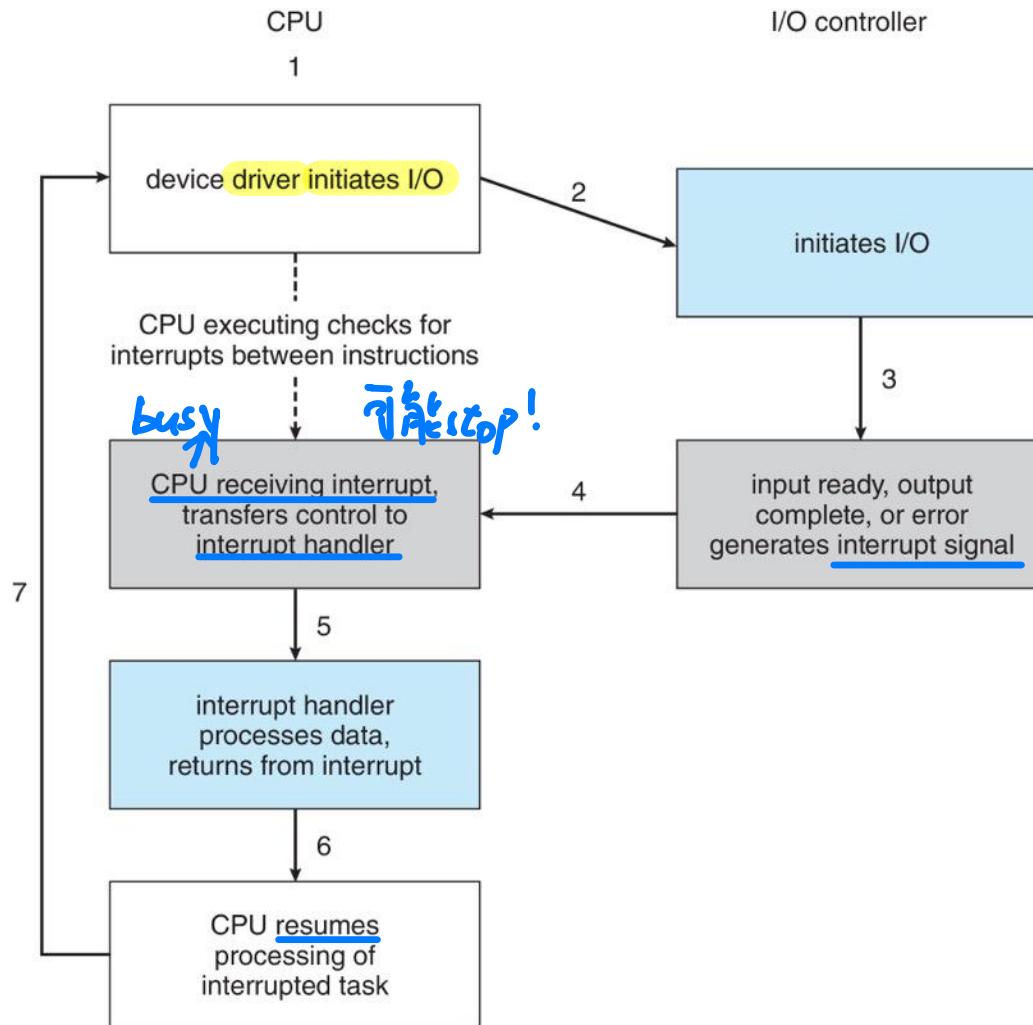
---

- I/O devices and CPU execute **concurrently** and **asynchronously**
- Each **device controller** is in charge of a particular device
- Each device controller has a local buffer
- The device controller is responsible for moving data between the peripheral devices that it controls and its local buffer storage
  - **I/O operations** are from the device to local buffer of the controller
- CPU moves data from/to main memory to/from local buffers, typically for slow devices such as keyboard and mouse
- **DMA** controller is used for move the data for fast devices like disks
  - *for more data of fast devices*
- The device controller informs CPU that it has finished an operation by causing an **interrupt** – requiring CPU attention
  - For input devices, this implies that data is available in local buffer
  - For output devices, it informs CPU that an I/O operation is completed





# Interrupt-Driven I/O Cycle





# Interrupt Timeline

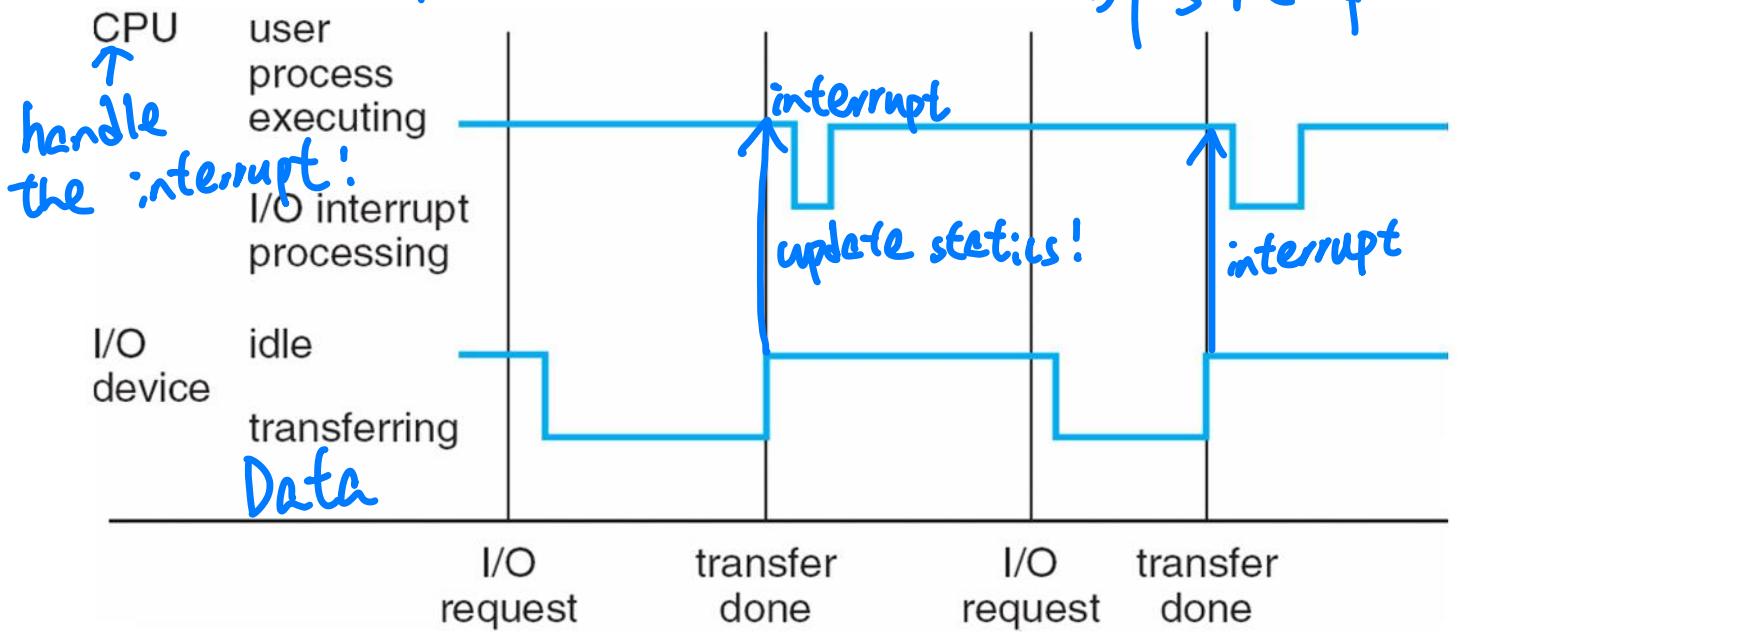
何时何故？为什么？CPU，多个设备控制

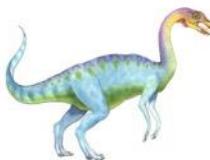
device control

can work concurrently

- CPU and devices execute concurrently
- An I/O device may trigger an **interrupt** by sending a signal to the CPU
- CPU handles the interrupt, and then returns to the interrupted instruction

work concurrently

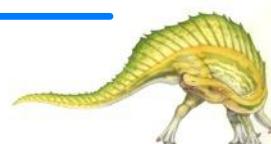


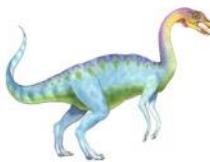


# Common Functions of Interrupts

- **Interrupts** are widely used in modern operating systems to handle asynchronous events - device controllers and hardware faults  
*隨時隨地*
- Interrupt transfers control to an **interrupt service routine** or **interrupt handler** – part of kernel codes, which OS runs to handle a specific interrupt  
*Maybe stop → handle the interrupt!*
- The interrupt mechanism also implements a system of interrupt priority levels, making it possible for a high-priority interrupt to preempt the execution of a low-priority interrupt *Example: divide by 0*
- A trap or exception is a software-generated interrupt caused either by an error (e.g., arithmetic errors) or a user request (e.g., a system call requesting OS services – to be discussed)
- All modern operating systems are **interrupt-driven**
- In a modern computer system, hundreds of interrupts occur per second – as CPU runs extremely fast in fraction of a nanosecond

*時間很短!*



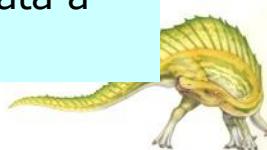


The basic unit of computer storage is the **bit**. A bit can contain one of two values, 0 and 1. All other storage in a computer is based on collections of bits. Given enough bits, it is amazing how many things a computer can represent: numbers, letters, images, movies, sounds, documents, and programs, to name a few. A **byte** is 8 bits, and on most computers it is the smallest convenient chunk of storage. For example, most computers don't have an instruction to move a bit but do have one to move a byte. A less common term is **word**, which is a given computer architecture's native unit of data. A word is made up of one or more bytes. For example, a computer that has 64-bit registers and 64-bit memory addressing typically has 64-bit (8-byte) words. A computer executes many operations in its native word size rather than a byte at a time.

Computer storage, along with most computer throughput, is generally measured and manipulated in bytes and collections of bytes.

- A **kilobyte**, or **KB**, is 1,024 bytes
- a **megabyte**, or **MB**, is  $1,024^2$  bytes
- a **gigabyte**, or **GB**, is  $1,024^3$  bytes
- a **terabyte**, or **TB**, is  $1,024^4$  bytes
- a **petabyte**, or **PB**, is  $1,024^5$  bytes

Computer manufacturers often round off these numbers and say that a megabyte is 1 million bytes and a gigabyte is 1 billion bytes. Networking measurements are an exception to this general rule; they are given in bits (because networks move data a bit at a time).

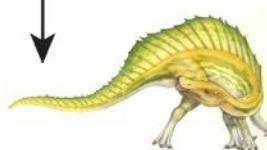
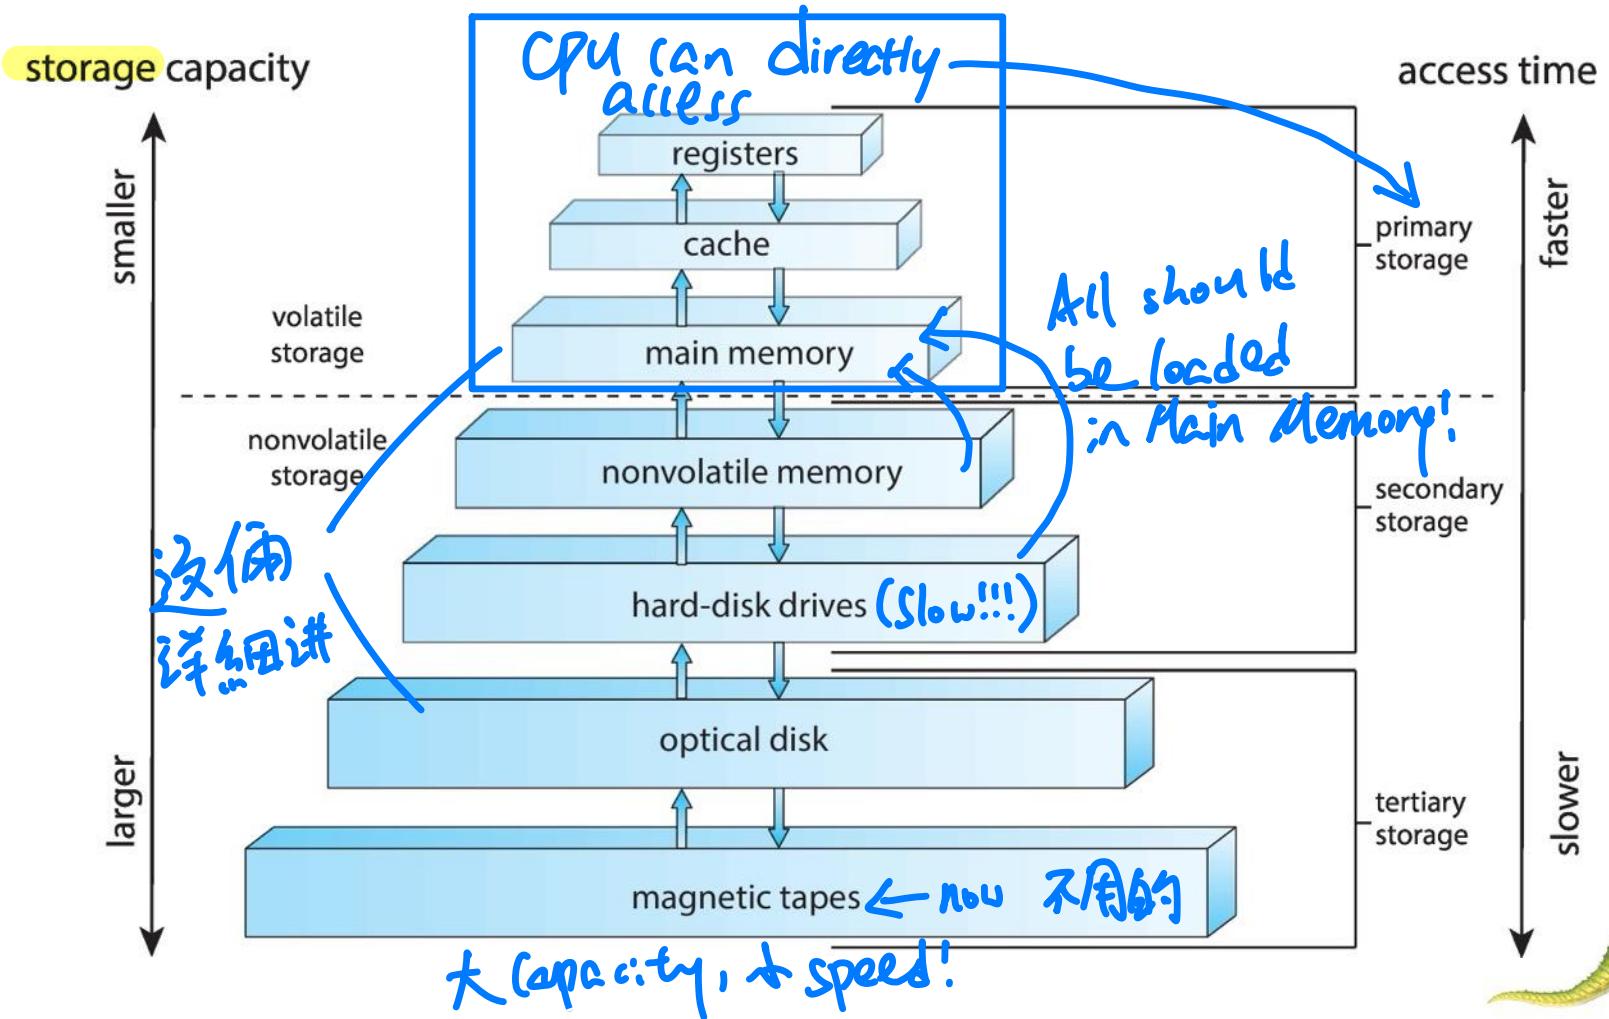


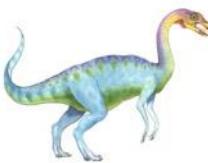


# Storage Hierarchy

Level 书度!

- Storage systems organized in **hierarchy**, varied with **speed**, **cost per unit**, **capacity (size)** and **volatility** (non-volatile disk vs. volatile memory)





# Memory

---

- Main memory – the only large storage media that CPU can access directly
  - **Volatile**, and typically random-access memory in the form of Dynamic Random-Access Memory (DRAM)
  - The basic operations **load** and **store** instructions to specific memory addresses, which is **byte-addressable** – each address refers to one byte in memory  
*from memory to register/memory*
- Computers use other forms of memory as well. For example, the first program to run on computer power-on is a **bootstrap program**, which is stored on electrically erasable programmable read-only memory (**EEPROM**)





# Second Storage

permanent

- The secondary storage – extension of main memory providing large non-volatile storage capacity, which can hold large quantities of data permanently. *Data still exists  
→ don't care the power on or off*
- The most common secondary-storage devices are **hard-disk drives** (HDDs) and **nonvolatile memory (NVM)** devices, which provide storage for both programs and data.
- There are generally two types of secondary storage

- Electrical faster than Mechanical*
- **Mechanical**, such as HDDs, optical disks, holographic storage, and magnetic tape
  - **Electrical**, such as flash memory, SSD, FRAM, NRAM. Electrical storage is usually referred to as **NVM** *(costly, smaller, but much faster)*
  - Mechanical storage is generally **larger** and less expensive per byte than electrical storage. Conversely, **electrical storage is typically costly, smaller, more reliable, and faster** than mechanical storage.





# Main memory is faster than hard disk, slower Caching than register & cache ...



- **Important principle** - performed at many levels in computers
  - Cache for memory, address translation, file blocks, file names (frequent used), file directories, network routes, etc. *many different chapters*
- **Fundamental idea:** A subset of information copied from a slower to a faster storage temporarily → *Speed up yr program*
  - Make frequently used case faster and less frequent case less dominant
- The access first checks to determine if information is inside the cache
  - Hit: if it is, information used directly from the cache (fast) *=access directly*
  - Miss: if not, data copied from slower storage to cache and used there
- Cache usually much smaller than storage (e.g., memory) being cached
  - Cache management: cache size and replacement policy *algorithm*
  - Major criteria – cache hit ratio; percentage *content found in cache* *the data* *which part will be cached*
- Important measurement

$$\text{Average Access time} = (\text{Hit Rate} \times \text{Hit Time}) + (\text{Miss Rate} \times \text{Miss Time})$$

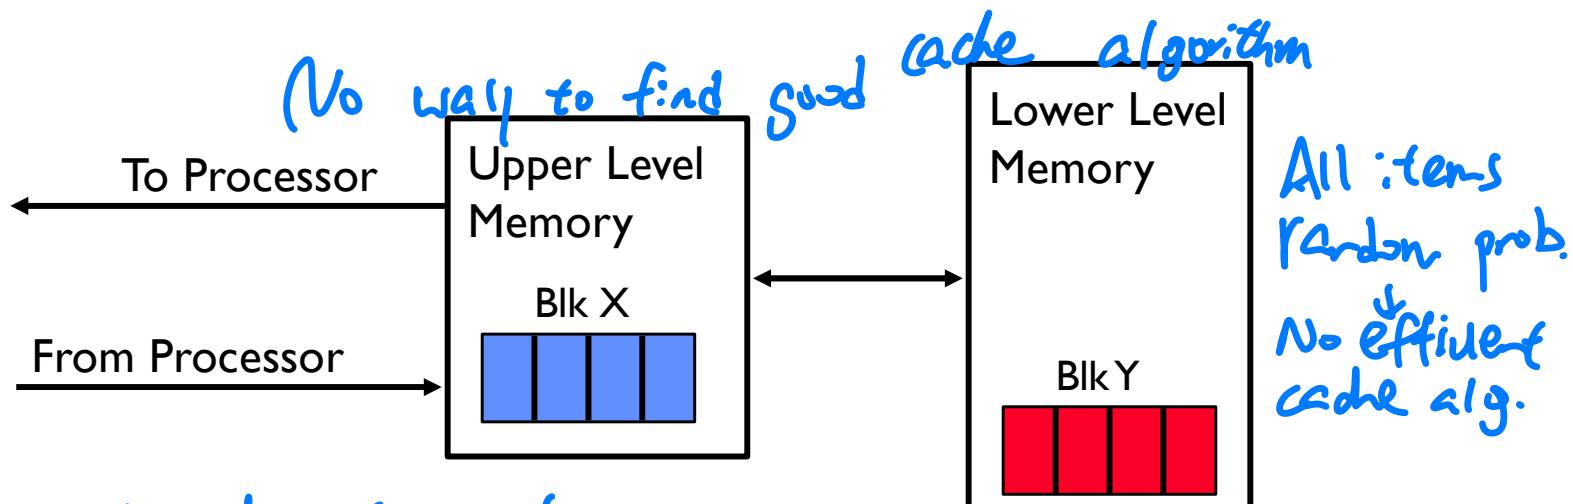




# Why Does Caching Work? - Locality

Possible make efficient cache

- Temporal locality (Locality in Time)
  - The recently accessed items likely to be accessed again → improve the performance
- Spatial locality (Locality in Space)
  - The contiguous blocks (i.e., those near the recently accessed items) likely to be accessed shortly (both data and program)
- Without access locality pattern, for instance If all items are accessed with equal probability, cache would never work!



No a good algorithm to  
find what things is good to cache





# Characteristics of Various Types of Storage

*much larger!!!*

*Small, fast*



*Speed much slower*

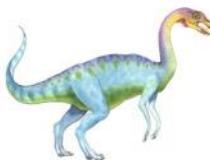
| Level                     | 1                                      | 2                             | 3                | 4                | 5                |
|---------------------------|--|-------------------------------|------------------|------------------|------------------|
| Name                      | registers                              | cache                         | main memory      | solid-state disk | magnetic disk    |
| Typical size              | < 1 KB                                 | < 16MB                        | < 64GB           | < 1 TB           | < 10 TB          |
| Implementation technology | custom memory with multiple ports CMOS | on-chip or off-chip CMOS SRAM | CMOS SRAM        | flash memory     | magnetic disk    |
| Access time (ns)          | 0.25-0.5                               | 0.5-25                        | 80-250           | 25,000-50,000    | 5,000,000        |
| Bandwidth (MB/sec)        | 20,000-100,000                         | 5,000-10,000                  | 1,000-5,000      | 500              | 20-150           |
| Managed by                | compiler                               | hardware                      | operating system | operating system | operating system |
| Backed by                 | cache                                  | main memory                   | disk             | disk             | disk or tape     |

*<1ns*

*i  
bit slower*

Movement between levels of storage hierarchy can be explicit or implicit





# Range of Timescales

Jeff Dean: “Numbers Everyone Should Know”

|                                    | <i>Access time</i> |
|------------------------------------|--------------------|
| L1 cache reference                 | 0.5 ns             |
| Branch mispredict                  | 5 ns               |
| L2 cache reference                 | 7 ns               |
| Mutex lock/unlock                  | 25 ns              |
| Main memory reference              | 100 ns             |
| Compress 1K bytes with Zippy       | 3,000 ns           |
| Send 2K bytes over 1 Gbps network  | 20,000 ns          |
| Read 1 MB sequentially from memory | 250,000 ns         |
| Round trip within same datacenter  | 500,000 ns         |
| Disk seek                          | 10,000,000 ns      |
| Read 1 MB sequentially from disk   | 20,000,000 ns      |
| Send packet CA->Netherlands->CA    | 150,000,000 ns     |

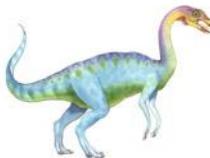




# I/O Subsystem

- OS needs to accommodate a wide variety of devices, each with different capabilities, control-bit definitions, and protocols for interacting with host
  - OS enables I/O devices to be treated in a **standard, uniform way** – that involves abstraction, encapsulation, and software layering, like for any complex software engineering design *very complex software design*
  - I/O system calls encapsulate device behaviours in a few generic classes, each is accessed through a standardized set of functions - interface
  - One purpose of OS is to hide peculiarities of hardware devices from users
  - I/O subsystem responsible for *complex* *need I/O subsystem*
    - Memory management of I/O including buffering (storing data temporarily while it is being transferred), caching (storing parts of data in faster storage for performance), spooling (the overlapping of output of one job with input of other jobs, typically used in printers)
    - General device-driver interface *- provide, high complexity and variety*
    - Drivers for specific hardware devices
- Hardware technology, not software*





# Direct Memory Access

~~inefficient data transfer every byte~~

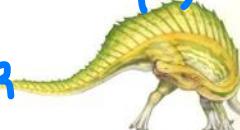
~~when lots of~~

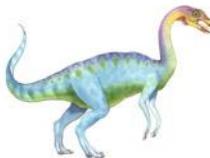
'briefly' in last class

- **Programmed I/O** - CPU runs special I/O instructions to move one byte at a time between memory and slow devices, e.g., keyboard and mouse
- To avoid programmed I/O, for fast devices and for large amount of data transfer, it uses direct memory access or DMA controller - bypasses CPU to transfer data between I/O device and memory directly - CPU or OS initializes DMA controller, and DMA controllers are responsible for moving the data between devices and memory without CPU involved.
- This relieves the CPU from slow data movement (I/O operations)
- OS writes DMA command block into memory
  - Source and destination addresses
  - Read or write mode
  - Number of bytes to be transferred
  - Writes location of command block to DMA controller
  - Bus mastering of DMA controller – grabs bus from CPU
  - When done, send interrupt to CPU for signaling completion

transfer large  
data from  
memory to fast  
device without  
involvement of  
CPU (不需 CPU  
interrupt)

only initializes the DMA → can move data  
between devices and memory CPU



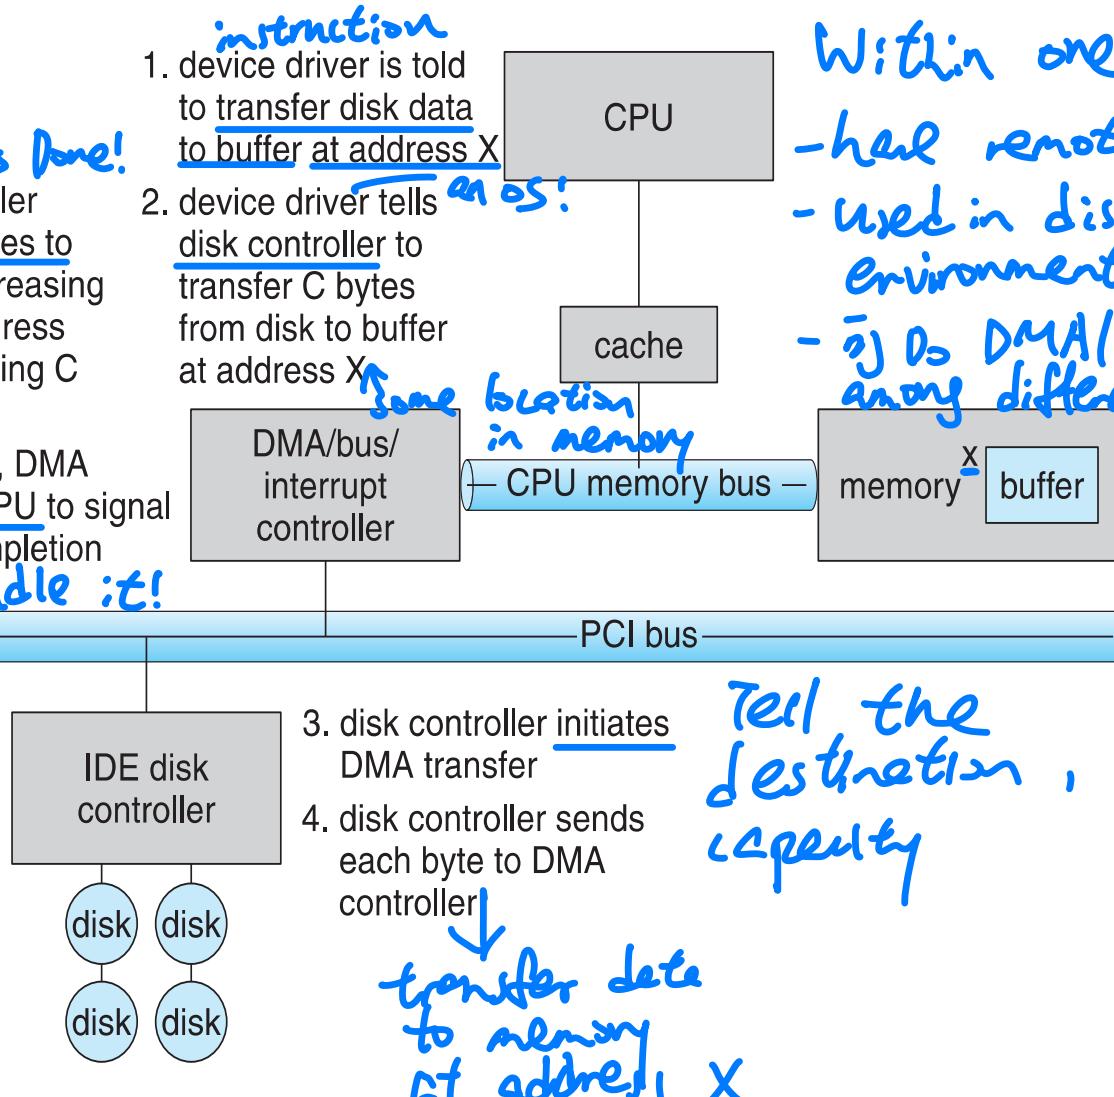


# Six Step Process to Perform DMA Transfer

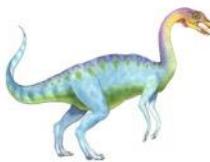
Transfer is Done!

5. DMA controller transfers bytes to buffer X, increasing memory address and decreasing C until  $C = 0$
6. when  $C = 0$ , DMA interrupts CPU to signal transfer completion

↓  
transfer  
finished  
↓  
raised interrupt!  
handler will handle it!  
resume its execution



Within one computer  
- local memory  
- used in distributed environment  
- ⇒ Do DMA/DMT among different computers  
↓ train model NL



# Single-Processor Systems

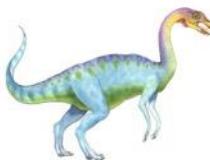
Physical chip (CPU)

- In the past, most computer systems used a single processor containing one CPU with a single processing core
  - The core executes instructions and registers for storing data locally.
  - The processing core or CPU core is capable of executing a general-purpose instruction set
- Such systems have other special-purpose processors - device-specific processors, such as disk, and graphics controllers (GPU).
  - They run a limited instruction set, usually do not execute instructions from user processes

I  
Can only run user CPU

Outdated





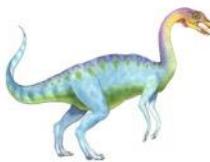
# Multiprocessor Systems

- On modern computers, from mobile devices to servers, multiprocessors systems now dominate the landscape of computing
  - Traditionally, such systems have two (or more) processors, each with a single-core CPU *more capability, more instructions in parallel*
  - The speed-up ratio with N processors is less than N, because of overhead, e.g., contention for shared resources (bus or memory)  
*Not linear increase!!!*
- Multiprocessors systems growing in use and importance, advantages are
  - Increased throughput – more computing capability  
*Buying 2 computers > Buying 2 processors!*
  - Economy of scale – share other devices such as I/O devices  
*false tbh*
  - Increased reliability – graceful degradation or fault tolerance
- Two types of multiprocessor systems
  - All computer system!  
Asymmetric Multiprocessing – often master-slave manner, the master processor assign specific tasks to slaves, and the master handles I/O  
*only one → single reliability  
Still one can compute other core!*
  - Symmetric Multiprocessing – each processor performs all tasks, including operating-system functions and user processes  
*cpu equals!*

*will handle I/O*

*Use NUMA Architecture*

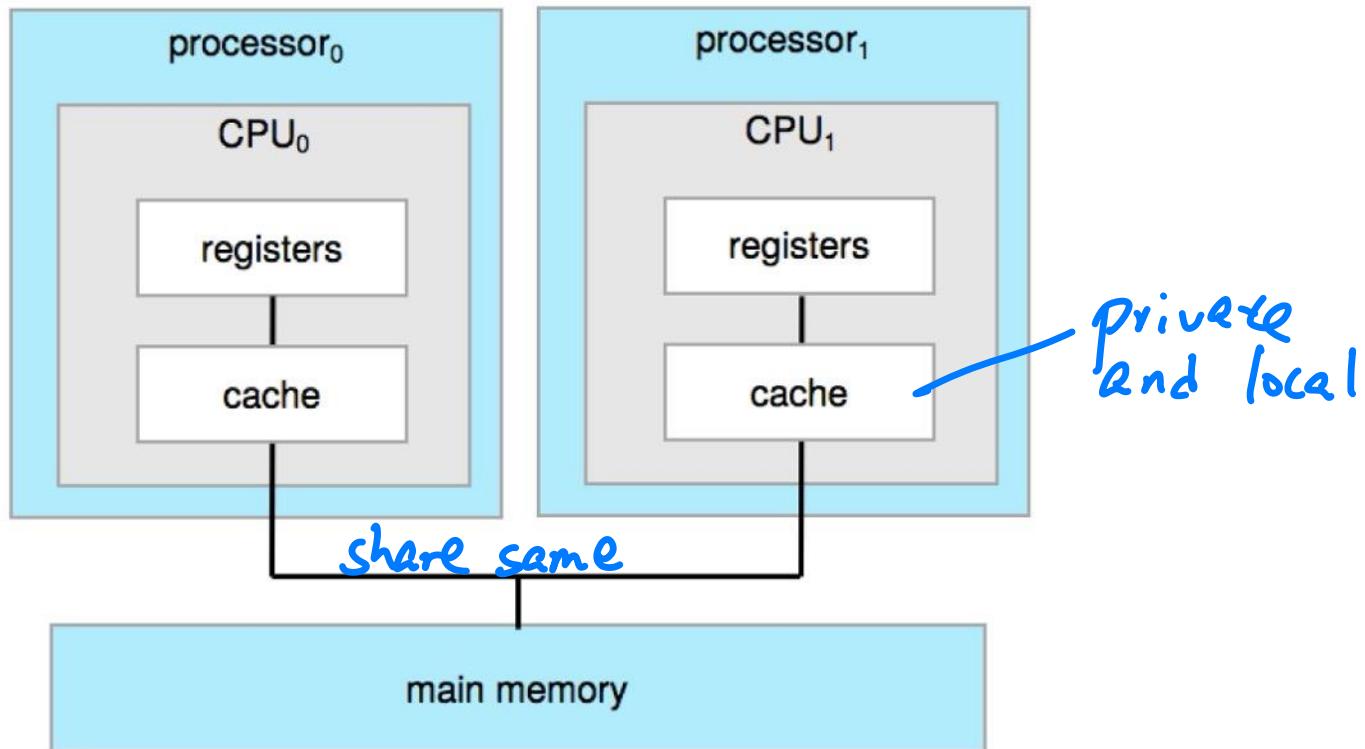
*one master processor,  
one master node*



# Symmetric Multiprocessor Systems

- Symmetric Multiprocessing or **SMP** – each CPU processor has its own set of registers, as well as a private or local cache. However, all processors share physical memory through system bus.

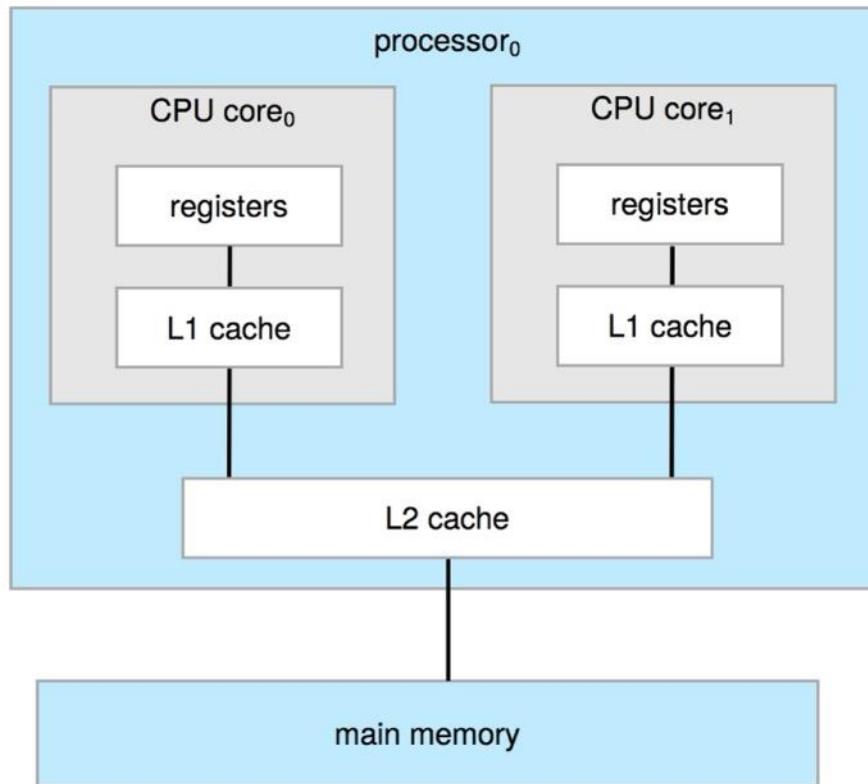
speed up  $< N$

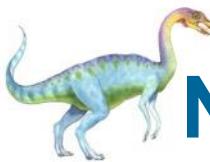




# A Multi-Core Design

- The **multicore**, multiple computing cores reside on **a single physical chip**
  - Faster on-chip communication than between-chip communication
  - Uses significantly less power - important for mobile devices and laptops

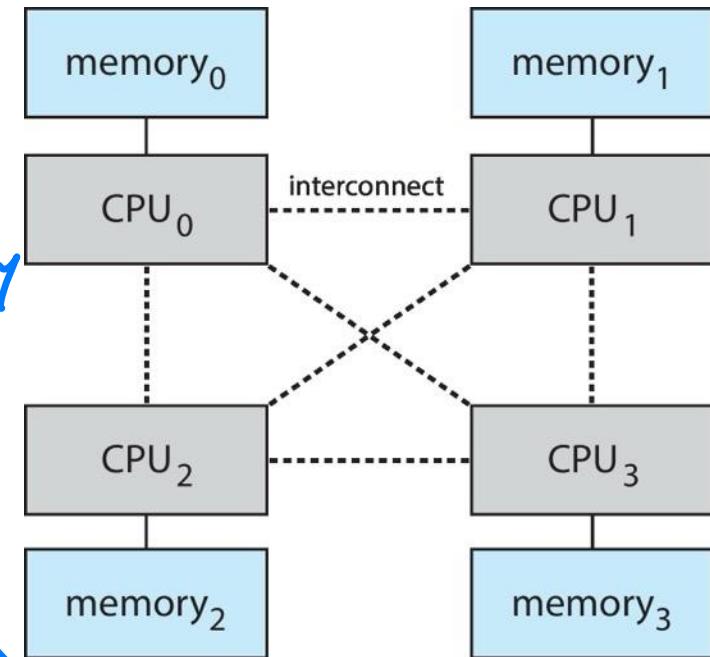




# Non-Uniform Memory Access (NUMA)

- Adding more CPUs to a multiprocessor system may not **scale**, due to the contention for **system bus**, which can become a bottleneck
- An alternative is to provide each CPU (or group of CPUs) with its own local memory that is accessed via a small, fast local bus.  
*(can access directly)*
- The CPUs are connected by a **shared system interconnect**, and all CPUs share one physical memory address space.
- This approach—known as **non-uniform memory access** or **NUMA**
- The potential drawback with a NUMA system is increased latency when a CPU must access remote memory across the system interconnect – scheduling and memory management implication

*more CPU  
more contention*



*solve the contention problem*

*CPU0 in memory3, will be slower*





# Computer System Component

---

- **CPU** - The hardware that executes instructions
- **Processor** - A physical chip that contains one or more CPUs
- **Core** – The basic computation unit of the CPU or the component that executes instructions and registers for storing data locally
- **Multicore** – Including multiple computing cores on a single physical processor chip
- **Multiprocessor system**– including multiple processors





# Operating System Structure

time sharing

multiple process in main memory

- There are two common characteristics in all modern operating systems
- **Multiprogramming** (batch system) is needed for efficiency *CPU is busy!*
  - In old days, OS loads one program into the memory at a time for execution
  - Single program cannot always keep CPU or I/O devices busy as they become faster and faster— all modern computer systems are multi-programmed
  - Multiprogramming organizes jobs *need job Scheduling* in a way hoping CPU always has one to execute
  - In mainframe computers, jobs are submitted remotely and queued, and jobs are selected and run via job scheduling – load into the memory (discussed later)
- **Timesharing** (multitasking) is logical extension of multiprogramming in which CPU switches “frequently” between jobs that users can interact with each job while it is running, enable interactive computing
  - Response time should be < 1 second
  - Each user has at least one program executing in memory ⇒ process
  - If several jobs ready to run at the same time ⇒ CPU scheduling *decide which process can be used first*
  - If processes do not fit in memory, swapping technique moves them in and out of memory during execution → *chapter 9*
  - Virtual memory allows execution of processes not completely in memory

larger than main memory! part program into the main memory!  
Remaining is in the hard disk

Multiprogramming: load 2 programs together to the main memory, ~~not~~ load one program at one time

- 2 cores  $\rightarrow$  do their own job

Multi-tasking: CPU can split into small slices, each slice can be assigned to different programs, working "parallel". Actually work concurrently, switch so quickly!



# Virtualization

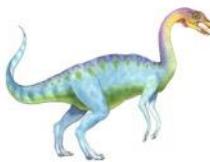
*Run other computers!*

- **Virtualization** abstracts the hardware of a single computer into multiple different execution environment(s) - creating an illusion that each user or program is running on its own “private computer”
  - It creates a virtual system - virtual machine or VM on which operation systems and applications can run over it
  - It also allows an operating system to run as an application within other operating system – this has been a vast and growing industry
- Several components
  - **Host** – underlying hardware system
  - **Virtual machine manager (VMM)** or **hypervisor** – creates and runs virtual machines by providing interface that is identical to the host
  - **Guest** – process provided with virtual copy of the host, usually an operating system – guest OS
- This allows a single physical machine can run multiple operating systems concurrently, each in its own virtual machine

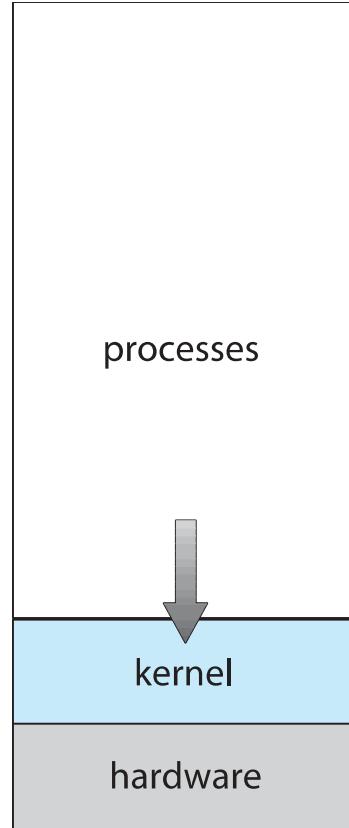
*Run Guest*

*User* ↗ ↘ *Windows prog in Mac computer!*

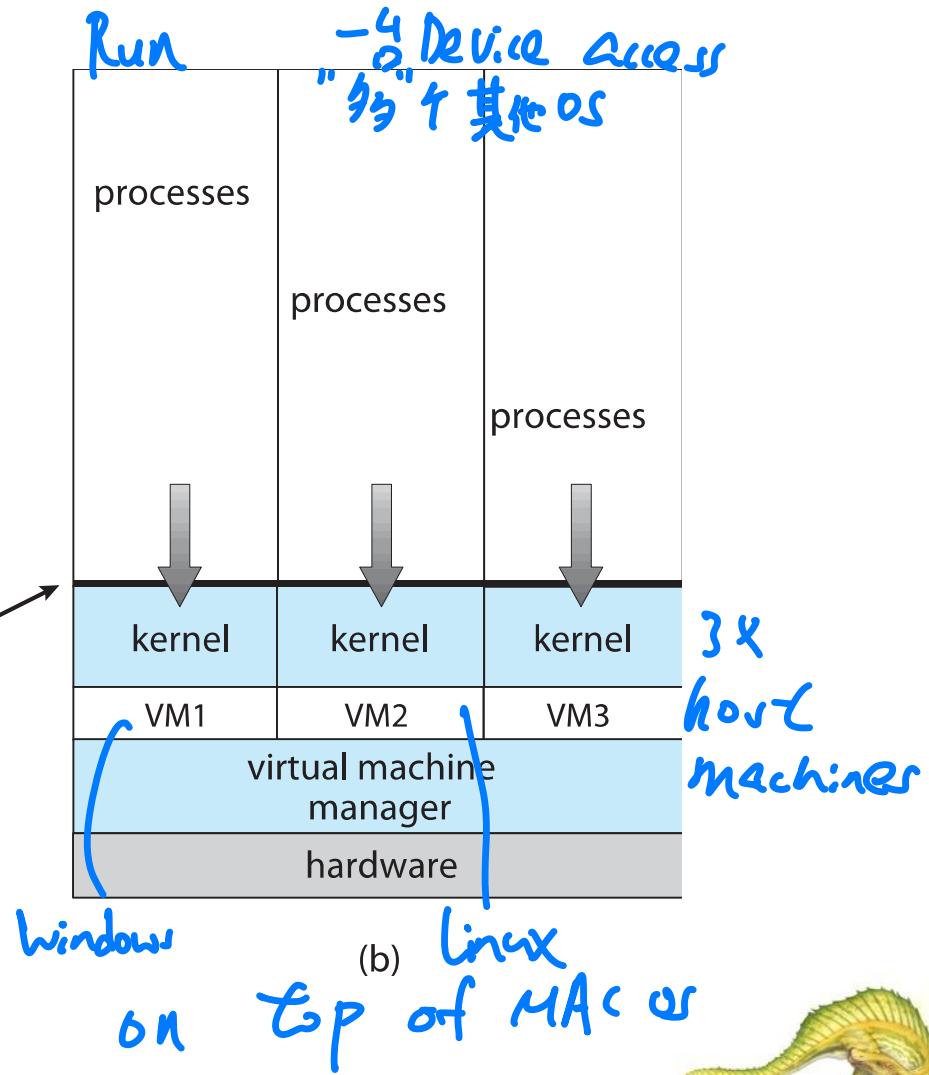




# Virtualization – System Models



(a)

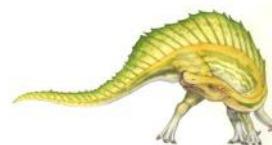




# # virtual barn! Virtualization – a bit history

- **Virtualization** – OS natively compiled for CPU, running guest OSes
  - Virtualization originally designed in IBM mainframes (1972) to allow multiple users to run tasks concurrently in a system designed for a single user or share a batch-oriented system
  - VMware runs one or more guest copies of Windows, each running its own applications, on Intel x86 CPU - 一样
  - A Virtual Machine Manager or VMM provides an environment for programs that is essentially identical to the original machine (interface)
  - Programs running within such environments show only minor performance decreases – passing more layers of software
    - The VMM is in complete control of system resources
- In late 1990s Intel CPUs fast enough - virtualization on general purpose PCs
  - Xen and VMware created technologies, still used today
  - Virtualization has expanded to many OSes, CPUs, VMMs

Amazon: launch virtual machine  
After used, delete virtual machine!



Full virtualization

VMware, KVM

You don't know you are in virtual environment,  
No different

Partial virtualization

Xen

You definitely know in virtualization environment!



# Large computer infrastructure, many CPU, memory, storage Cloud Computing and Virtualization - one to multiple machines

many years (2010 - )

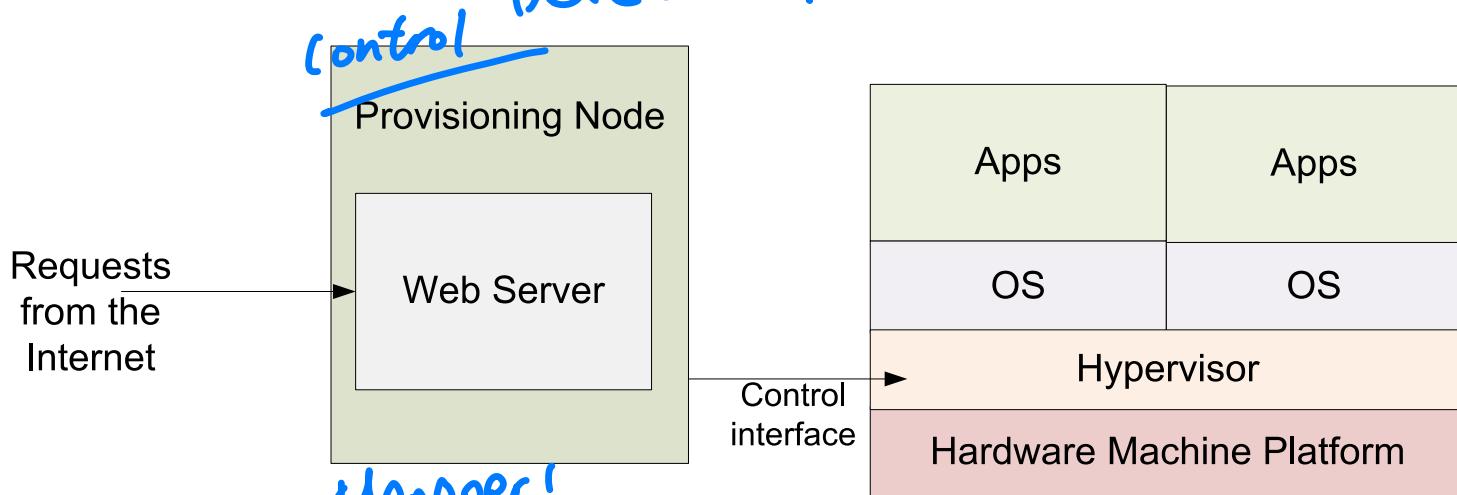
large datacentre

RFI

buy

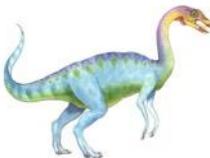
- Delivers computing, storage, and apps as a service over a network
- Logical extension of virtualization because it uses virtualization as the base for its functionality. - Across an internet (very flexible) machines
  - Amazon EC2 has millions of servers, tens of millions of VMs, petabytes of storage available across the Internet, pay based on usage

Delete Apps → not buy again



Example: Google Drive!  
Amazon





# Cloud Computing Types

- Many types of clouds

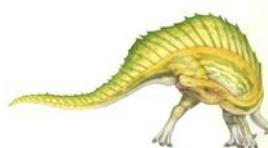
- Public cloud – available via Internet to anyone willing to pay
- Private cloud – run by a company for the company's own use
- Hybrid cloud – includes both public and private cloud components
- Software as a Service (SaaS) – one or more applications available via the Internet (i.e., word processor) → provide software service
- Platform as a Service (PaaS) – software stack ready for application use via the Internet (i.e., a database server) → provide platform
- Infrastructure as a Service (IaaS) – servers or storage available over Internet (i.e., storage available for backup use) → low level, provide
- Increasingly provides other services, such as MaaS or Machine learning as a Service

Internet / Access!

- depend  
on users

build their own machines

ML as a Service





# Free and Open-Source Operating Systems

## Windows

- Operating systems made available in source-code format rather than just binary closed-source and proprietary
  - Microsoft Windows is a well-known example of the closed-source approach.

- Started by Free Software Foundation (FSF), which has “copyleft” GNU Public License (GPL)

*free: can do anything!*

- Free software and open-source software are two different ideas
  - ▶ <http://gnu.org/philosophy/open-source-misses-the-point.html/>
  - Free software not only makes source code available but also is licensed to allow no-cost use, redistribution, and modification. Open-source software does not necessarily offer such licensing

- Popular examples include GNU/Linux, FreeBSD UNIX (including core of Mac OS X - Darwin), and Solaris - OS

- Open-source code is arguably more secure, allowing more programmers to contribute, and is certainly a better learning tool

*Have license → always expertise!  
→ also make yr source code open source!*



# End of Chapter 1

