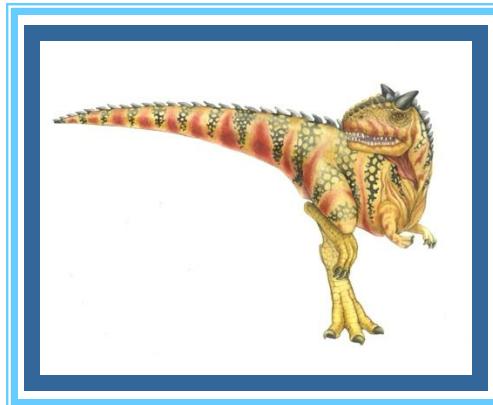


Chapter 3: Processes





Chapter 3: Processes

- Process Concept ← why we need
- Basic Scheduling Concept ← system view
- Operations on Processes – Process Creation and Termination
- Interprocess Communications - IPC
- Communication in Client-Server Systems – Socket

? create, terminate process → system call!



fork() . exec() , wait()

what is done? what is output





Objectives

- Identify the components of a process and illustrate how they are represented and scheduled in operating systems.
- Describe how processes are created and terminated in an operating system, using the appropriate **system calls** that perform these operations.
- Describe and contrast inter-process communications (IPC) using **shared memory** and **message passing** methods

how to use system calls to do processes!





Four Fundamental OS Concepts

Process

Every program is process(es)

(unit for OS to manage running program)

Example:

DATA

control

loop

one control workflow is in one thread

- An instance of an executing program is a process - consisting of an address space and one or more threads of control

show in memory ZI

Control flow, consumed by control control

Thread

more in Chapter 4

- A single unique execution context; fully describes program state – captured by program counter, registers, execution flags, and stack

Address space

- (with address translation) – to be discussed in Chapter 9)

- A program executes in its own address space that is distinct from the address space of another program and different from memory allocated for this process –logical/virtual address discussed in Chapters 9-10)

?

Dual mode operation

– Basic Protection

protected user mode/ kernel mode

- User programs and OS (kernel) programs run in different modes
- Only the operating system is allowed to access certain resources
- The OS and the hardware are protected from user programs
- User programs are protected and isolated from one another





Process Concept

~~multiple processes!~~
→ ~~binary~~ → ~~load! load!~~ → All bring to main memory!

- OS executes a variety of programs, each runs as a process
- **Process** captures a program execution; process executions progress in a **sequential** manner - von Neumann architecture
- The term “process” contains multiple “parts”
 - The program code, also called **text section**
 - Current activities or state - program counter, processor registers, *Fetch first!*
 - **Stack** - temporary data storage when invoking functions such as function parameters, return addresses, and local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time
- A program is a **passive** entity stored on disk (**executable file**), while a process is **active** entity, with a program counter specifying the next instruction to fetch and execute, as well as a set of associated resources including address space
- A process has a **life cycle** – from program execution to termination





? main start from 0 ?!!

Process Concept (Cont.)

ONLY Running program becomes a process!!! logical address

- A program “becomes” a process when an executable file is loaded into main memory and gets ready to run

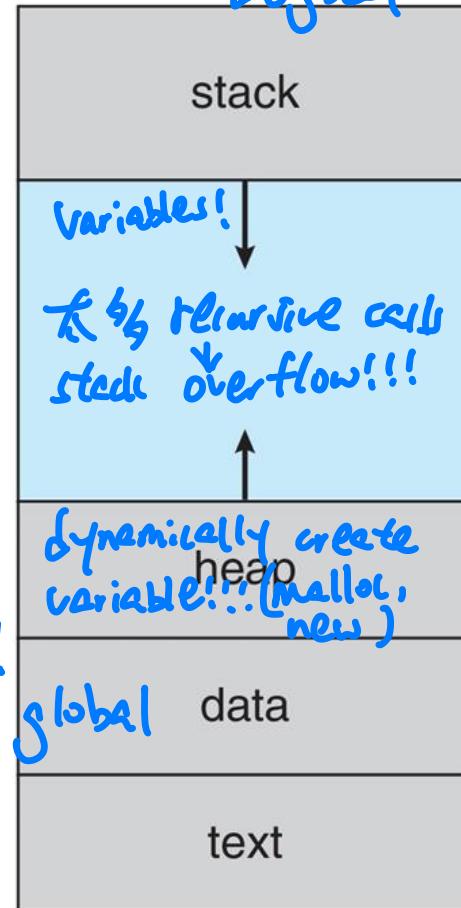
- Double-click a program icon or input a command on *Shell* isolated as!

- Program executes in an **address space** → (right side), **different** from actual main **virtual!!!** memory location – **logical address space** (discussed in Chapters 9-10)

- One program (code) can be executed by multiple processes → **different address!!!!**

- **A process can itself be an execution environment for other programs**, e.g., a Java program is executed within the Java virtual machine (JVM)

max



Different processes cannot access each other, **Address Space Isolated!**
- message passing
- share memory





Loading Program into Memory

Von Neumann Architecture

A program must be loaded inside memory before execution – a program icon is double clicked on a personal computer, or a job is submitted to a remote mainframe

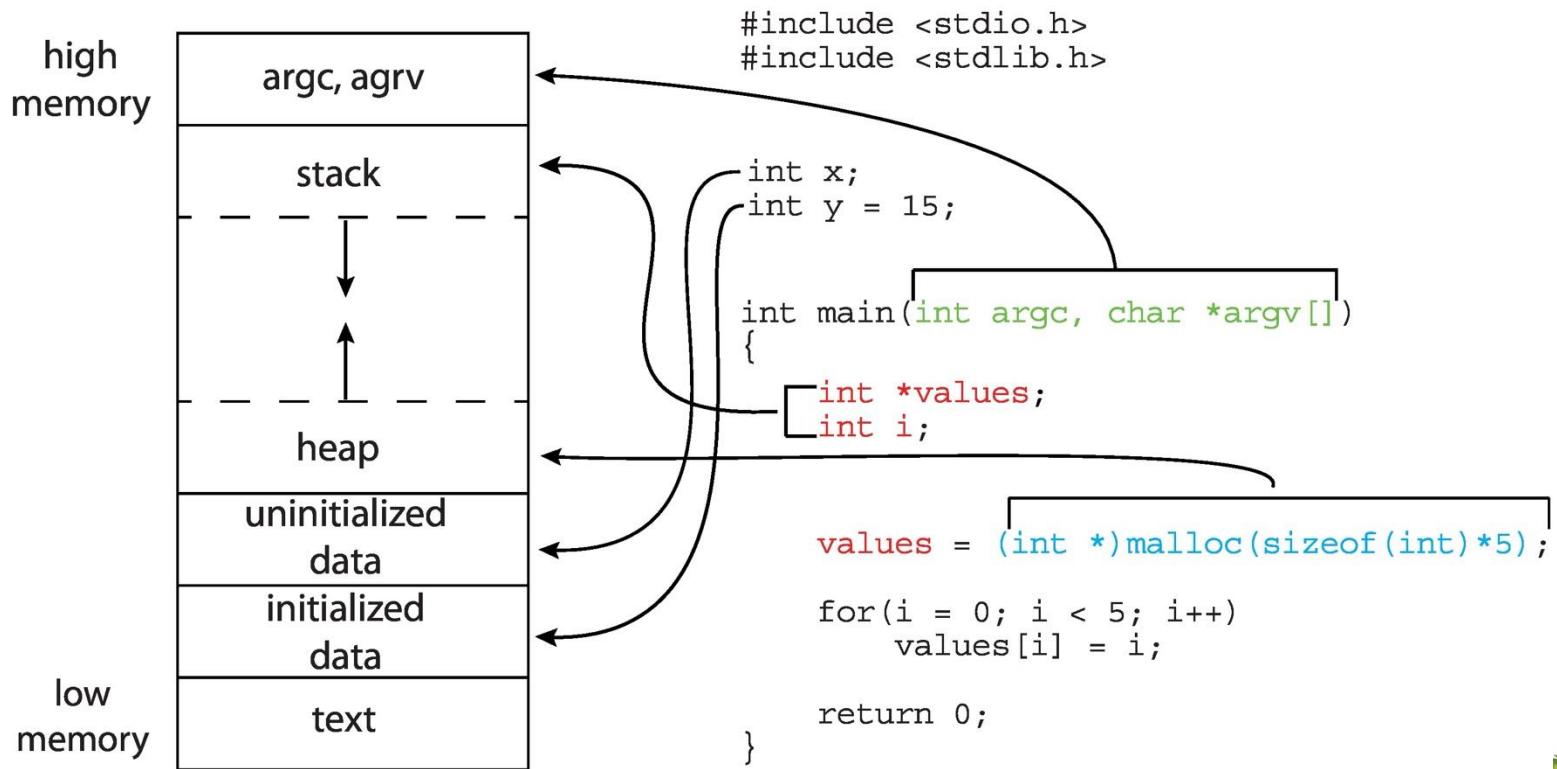
1. Load its code and any static data (e.g., initialized variables) into memory, or into the address space of the process. *higher address!*
2. Some memory must be allocated for the program's **runtime stack** - C programs use the stack for local variables, function parameters, and return addresses. Fill in the parameters to the `main()` function, i.e., `argc` and the `argv` array
3. OS may allocate memory for the program's **heap** - programs request memory space by calling `malloc()` and free it explicitly by calling `free()` *← global data intersektion*
4. OS will also do some other initialization tasks, esp. related to input/output (I/O). For example, in UNIX systems, each process by default has three open file descriptors, for standard **input**, **output**, and **error**.





Memory Layout of a C Program

- The global data section is divided into (1) initialized data and (2) uninitialized data
- A separate section provided for `argc` and `argv` parameters passed to `main()` function

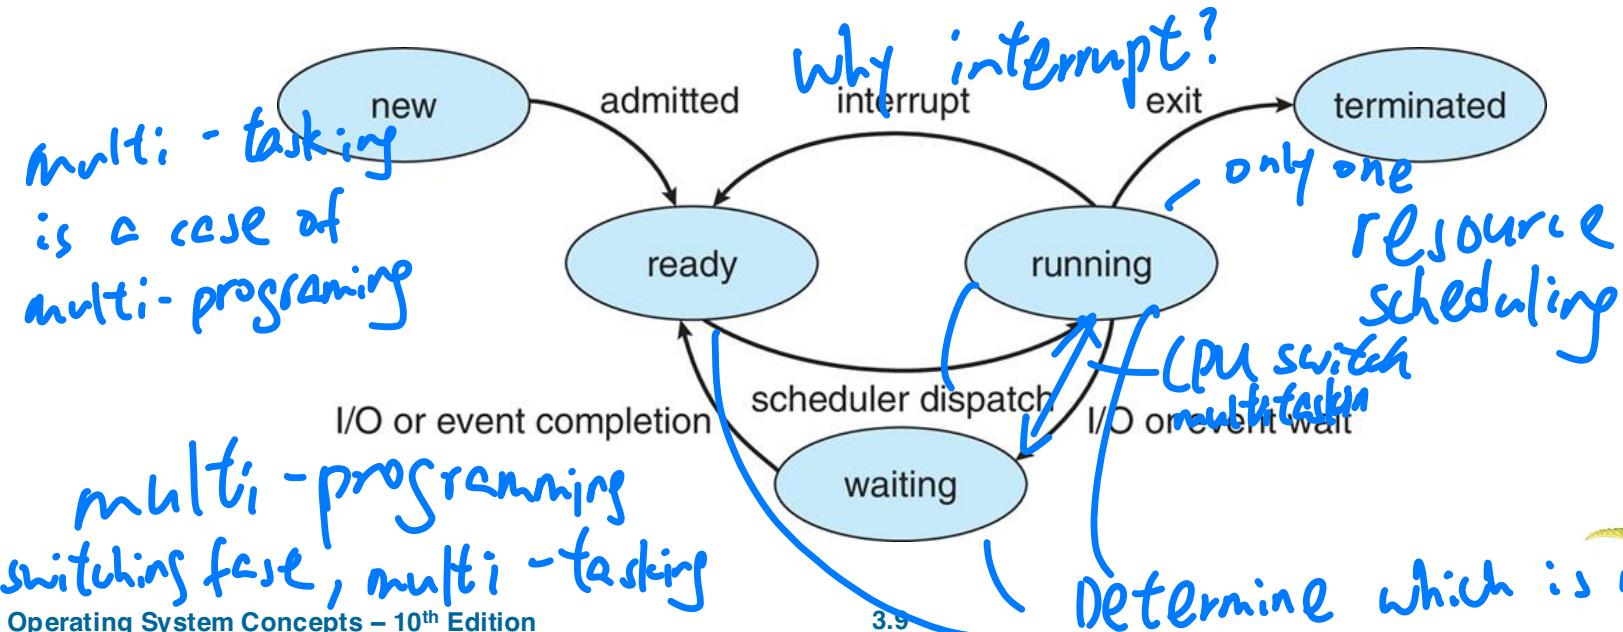


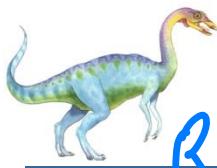


Process State

Nowadays can have multiple states!

- In traditional Unix systems, each process consists of one thread, thus the process state is defined in part by the current activity of that process
- As a process executes, it changes its **state** from one to another
 - **New**: The process is being created – allocating resources
 - **Ready**: The process is waiting to be assigned to a CPU core
 - **Running**: Instructions are being fetched and executed on a CPU
 - **Waiting**: The process is waiting for some event(s) to occur
 - **Terminated**: The process finishes execution – deallocating resources





Kernel things!

Process Control Block (PCB)

Running in kernel mode

I/O = 1 process

Information associated with each process

(also called **task control block**)

AI.S.In different state manage

Process state – running, waiting, etc.
Program counter – location of instruction to be fetched next to execute *+ recording!*

CPU registers – accumulators, index registers, stack pointers, and general-purpose registers, plus any condition-code information.

Value → store in PCB
CPU scheduling information – priorities, pointers to scheduling queue, scheduling parameters

Memory-management information – memory allocated to the process (complicated data structure – paging and segmentation tables) *in chapter 9-10*

Accounting information – CPU used, clock time elapsed since start, time limits

I/O status information – I/O devices allocated to process, for instance, a list of open files

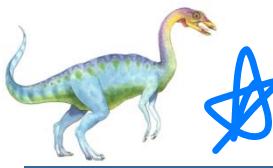
PLB is kernel code, running inside kernel!
3.10

kernel code (C/C++)

process state
process number
program counter
registers
memory limits
list of open files
• • •

data structure
in the Kernel!





Threads

Each core → run one! → more in next lecture

- So far, we assume that a process has a **single** thread of execution
 - It executes a program from the first line of codes to the last line of code, *sequentially*, in a serialized manner. The single execution context fully describes the program state - the current activity of the thread
 - A **thread** is represented by program counter, registers, execution flags, and stack – part of a process
 - A thread is executing on a processor or CPU when it is resident in processor (CPU) registers - hardware registers
- Most modern operating systems allow a process to have multiple threads of execution and thus can perform more than one task at a time, in parallel
 - A web browser (a process) might have one thread displaying images or text while another thread retrieving data from the network
 - On multicore systems, multiple threads of a process can run in parallel
- A process can consist of one or multiple **thread(s)** and they run within the same **address space** (i.e., share code, data and files), each thread has its own unique **thread state**
- Multithreads in a process provide **concurrency**, “active” components

More details to come in Chapter 4...



Thread 不会 share stacks!!!

只会 Share same heap / global / data / text
↑

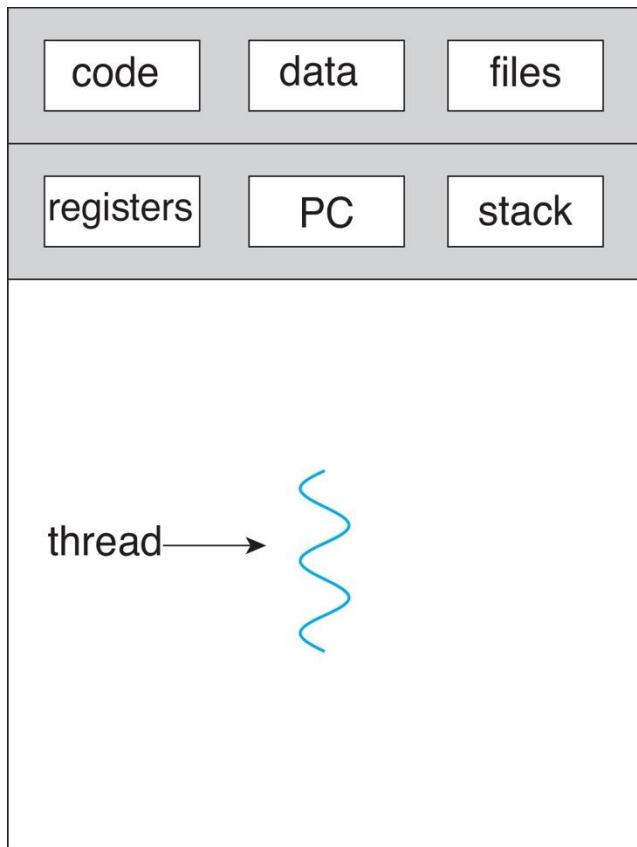
有 address → 可 access heap

因为 不同 functions 可 push 到 stack

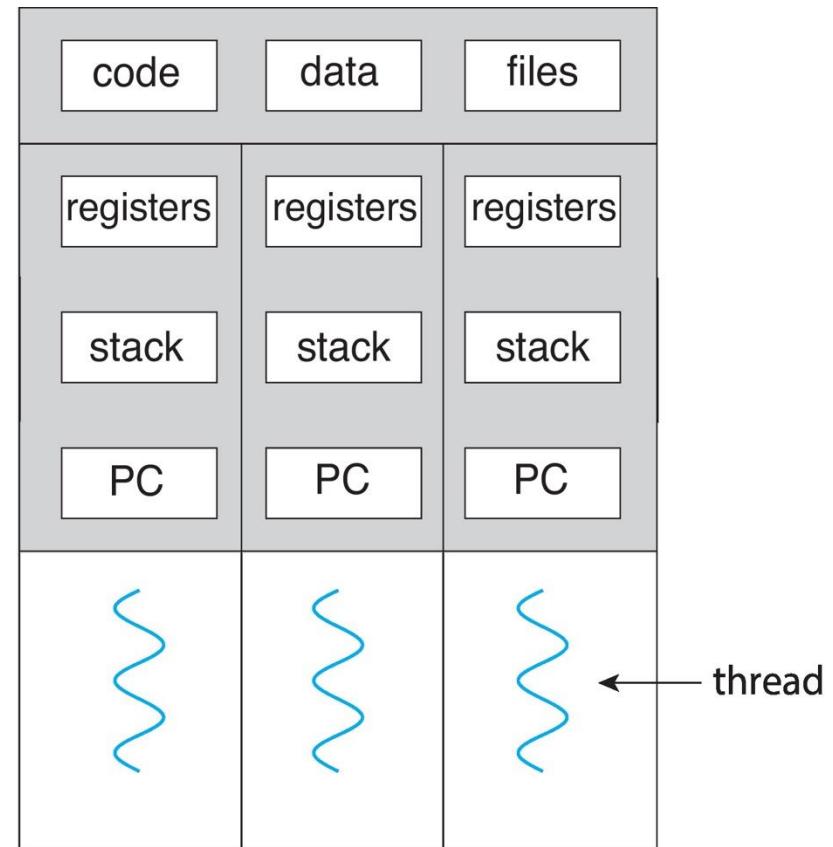
Multithread 适合 Modern Computer



Single and Multi-threaded Processes



single-threaded process



multithreaded process





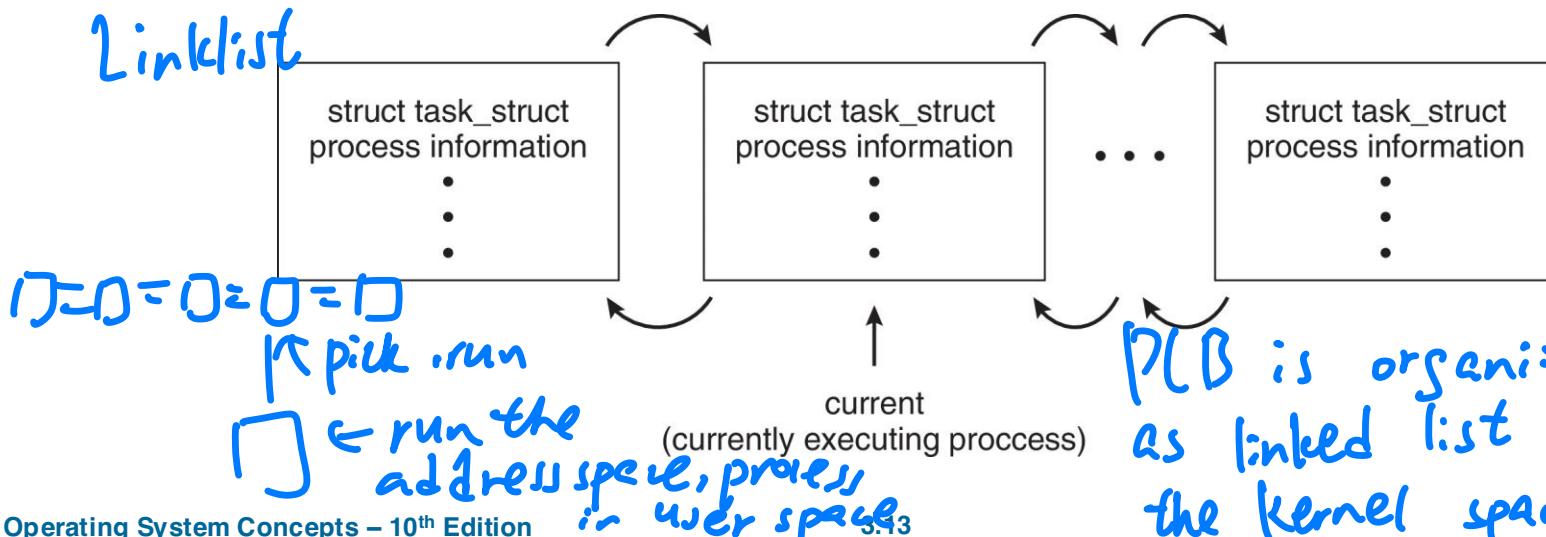
Process Representation in Linux

- The Linux kernel stores the list of processes in a circular doubly linked list called the **task list**. Each element in the task list is a process descriptor (i.e., PCB) of the type struct **task_struct**

```
pid t_pid;                      /* process identifier */  
long state;                     /* state of the process */  
unsigned int time_slice          /* scheduling information */  
struct task_struct *parent; /* this process's parent */  
struct list_head children; /* this process's children */  
struct files_struct *files; /* list of open files */  
struct mm_struct *mm;           /* address space of this process */
```

Good for insert and deletion

Linklist





Process Scheduling

More in

Ch 5

Branch of PGS,

- All modern operating systems are multiprogramming - the primary objective is to try to have some processes always running on the CPU for maximizing CPU utilization – recall CPU usually runs much faster than other devices
- **Process scheduler** – an OS mechanism selecting a process (from available processes inside main memory) for execution on one CPU core
 - Scheduling criteria: CPU utilization, job response time, fairness, real-time guarantee, latency optimization – to be discussed in Chapter 5
- For a system with only a single CPU, there will never be more than one process or thread running on the CPU at a time $(\rightarrow \leq 1)$
- The number of processes currently residing in memory is known as the **degree of multiprogramming** – determining the resource consumption
↑ process in main memory → degree = 3
- Processes (useful for mainframes) can be “roughly” described as either:
 - **I/O-bound process** – spends more of its time doing I/O than it spends doing computations; many short CPU bursts or CPU cycles
 - **CPU-bound process** – in contrast, generates I/O requests infrequently, using more of its time doing computations; few very long CPU bursts

different scheduling policy



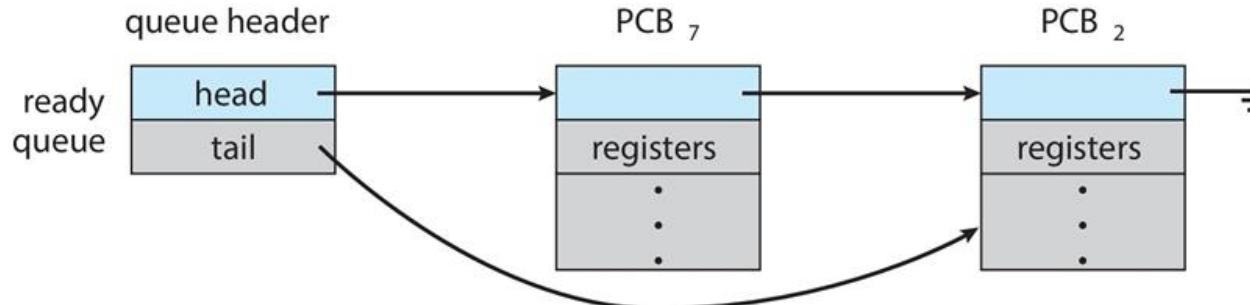


Process Scheduling (Cont.)

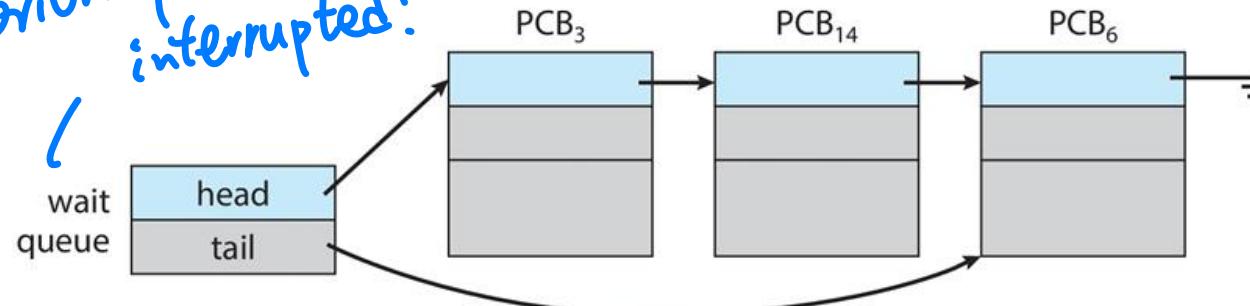
Directly bring to CPU

- The OS maintains different **scheduling queues** of processes
 - Ready queue** – set of processes in main memory, ready and waiting to execute
 - Wait queues** – set of processes waiting for an event such as completion of I/O
 - Processes essentially migrate among various queues during their lifetime

running



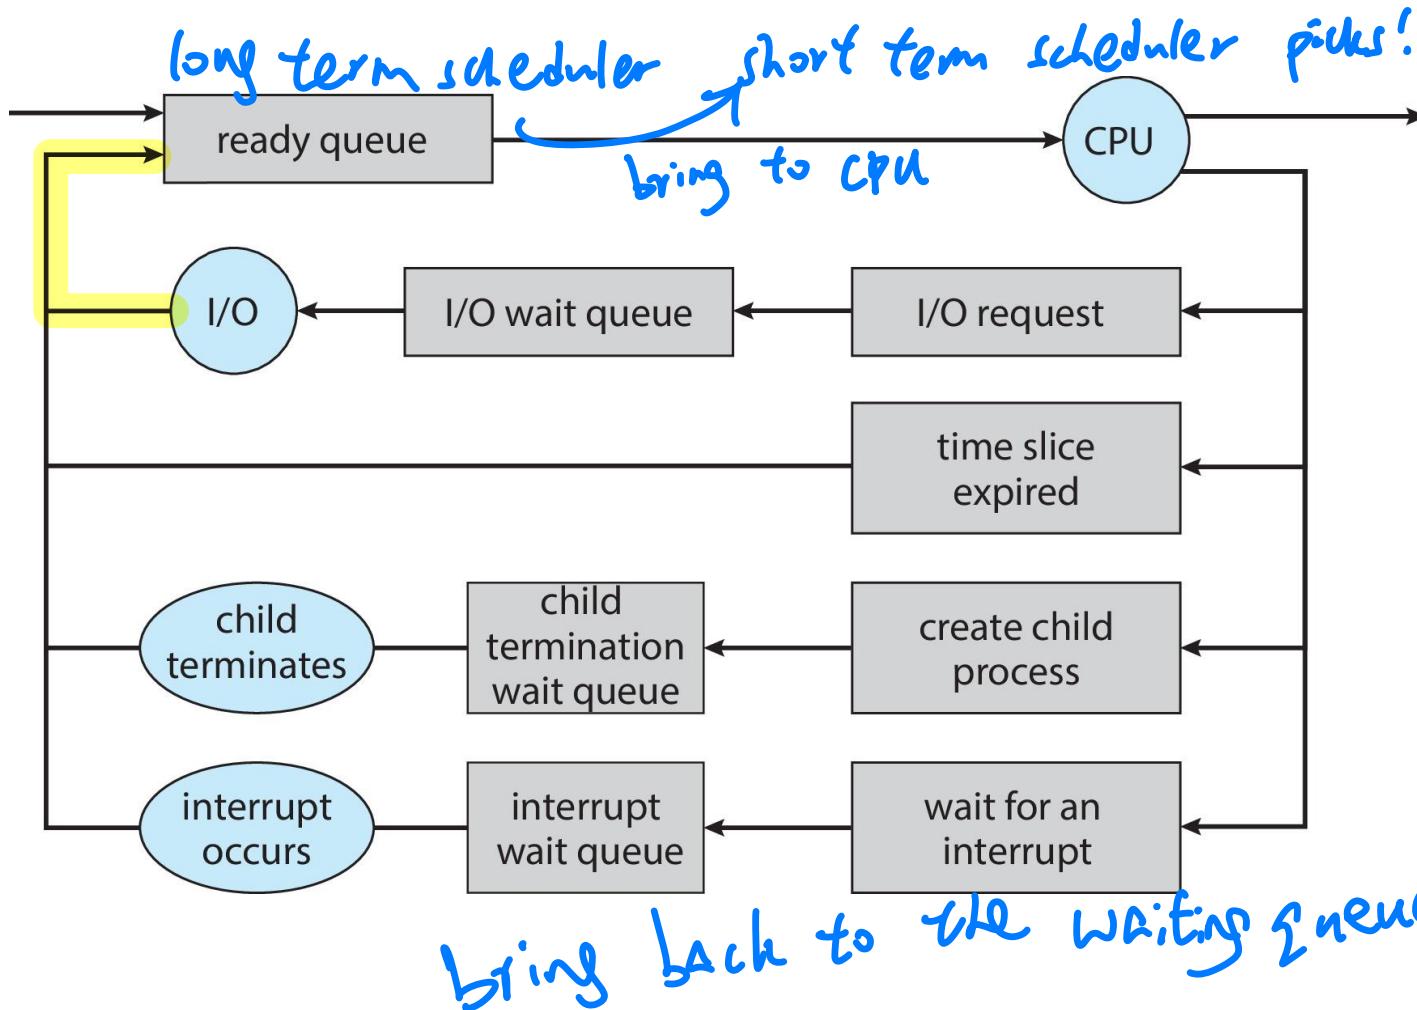
priority to be interrupted!





Representation of Process Scheduling

- A common representation of process scheduling is a queueing diagram. The types of queues present a ready queue and a set of wait queues.





cluster

Schedulers

Ready queue

- Long-term scheduler (or job scheduler) – selects which processes should be brought into memory (after resource allocation such as memory and files) – used by mainframe and minicomputers, not personal computers

Ready queue, it will allocate memory, some other resources and run in CPU

- It consists of a job queue that hosts jobs submitted to mainframe computers (shared by multiple users), yet brought into the memory
- Short-term scheduler (or CPU scheduler) – selects a process from the ready queue to be executed next and allocates a CPU core to run it ~~more frequently~~ means very fast!
- The short-term scheduler is invoked very frequently (milliseconds) ⇒ (must be fast), while long-term scheduler is invoked very infrequently (seconds, or minutes) ⇒ (may be slow)
- The long-term scheduler dictates the degree of multiprogramming – the number of processes that can concurrently be running in a computer.
- Long-term scheduler for a mainframe computer system strives for good process mix – mixed I/O-bound and CPU-bound processes

↓ better !

Decide which process will be fetched to the CPU



Relationship between process, thread, address space

- One thread have multiple threads, \rightarrow different PC
threads will share same of address space!
- Share all, except texts
 - Done by runtime automatically
- Threads from different process not share resources!

```
int main(int argc, char** argv)
```

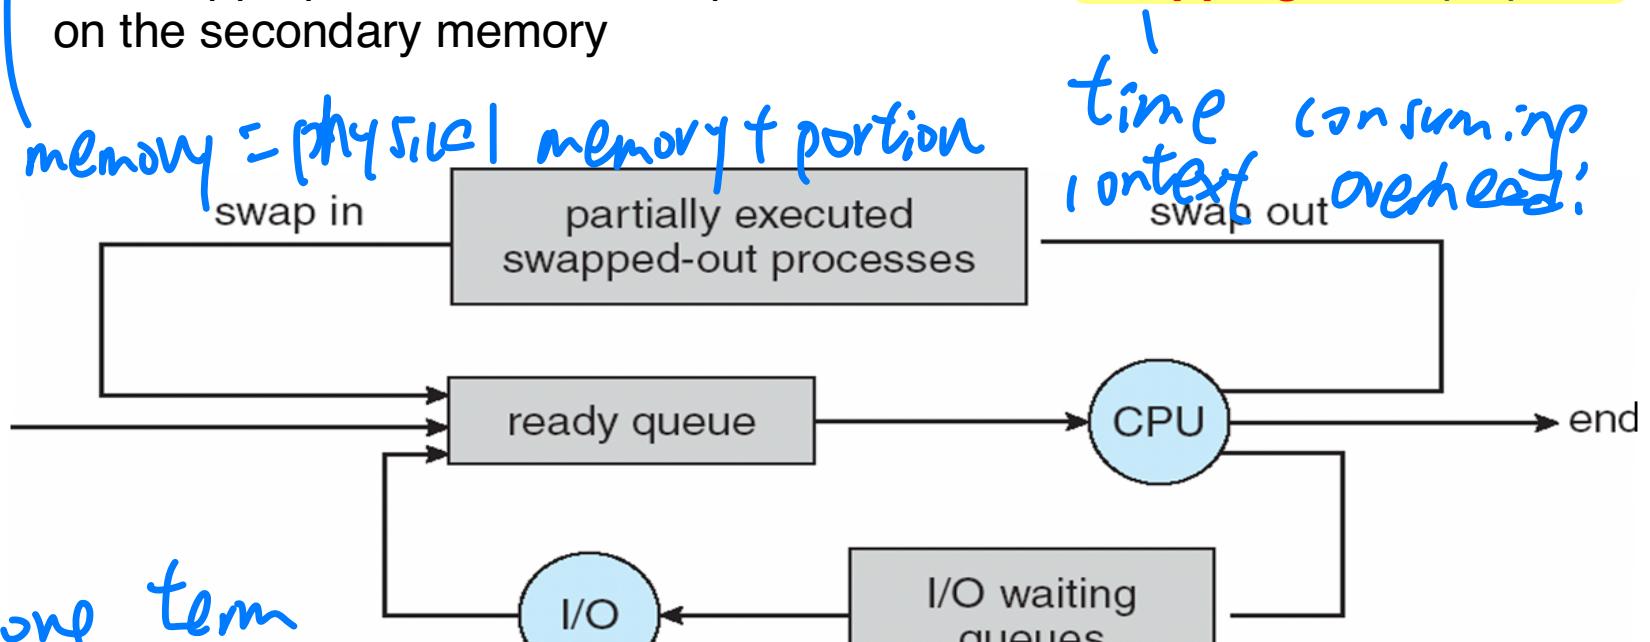
\uparrow \uparrow
by copy by reference
|
different address!



Run out of memory \rightarrow swap some process to secondary storage!

- Medium-term scheduler can be added if degree of multiple programming needs to be reduced - to free up memory space for remaining processes
- It removes a process from memory, store that on disk (secondary storage) temporarily, bring back later from disk to memory for continuing execution when appropriate – this technique is referred as swapping – swap space on the secondary memory

memory = physical memory + portion



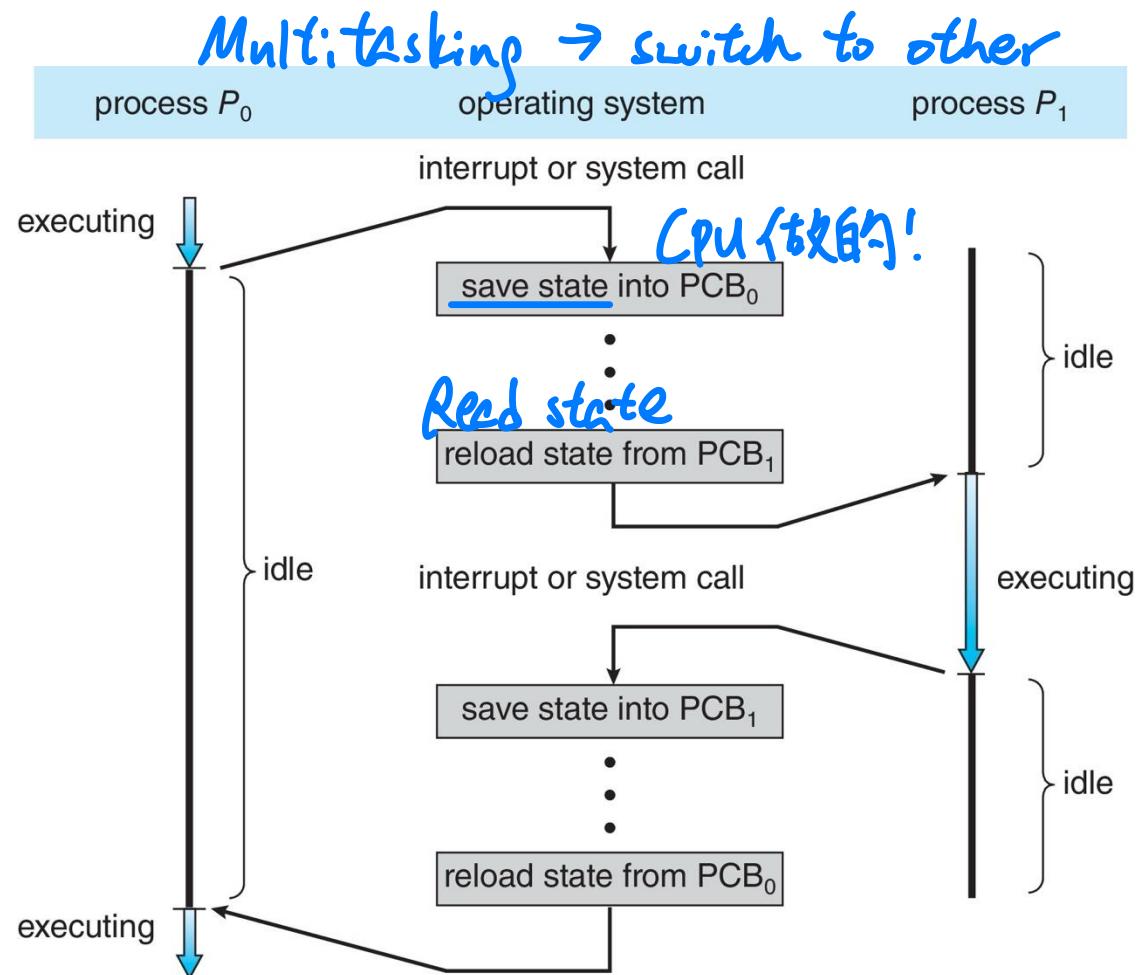
long term
short
Medium





CPU Switch From Process to Process

- A **context switch** occurs when CPU switches from one process to another





Context Switch

No free cost!

Modern computer!

Cause overhead!

- When CPU switches to another process, OS must save the state of the old process and load the saved state from another process via a context switch
 - The context of a process represented in the PCB or the thread
 - Typically, this includes registers, program counter, and stack pointer
 - Context-switch time is an overhead, during which the system does no useful work while switching → waste some resources!
 - The more complex the OS and the PCB, the longer the context switch
- □ Depends! Switching speed varies from machine to machine, depending on the memory speed, number of registers copied, and the existence of special instructions such as a single instruction to load or store all registers
- Context-switch times are highly dependent on the underlying hardware
 - For instance, some processors provide multiple sets of registers (e.g., SUN UltraSPARC). A context switch simply requires changing the pointer to a different register set – in this case, multiple contexts can be loaded at once

Special CPU 2 sets of registers
↓
Context switch → 不用 save!

Depends on hardware
↓



Dual-mode Operation

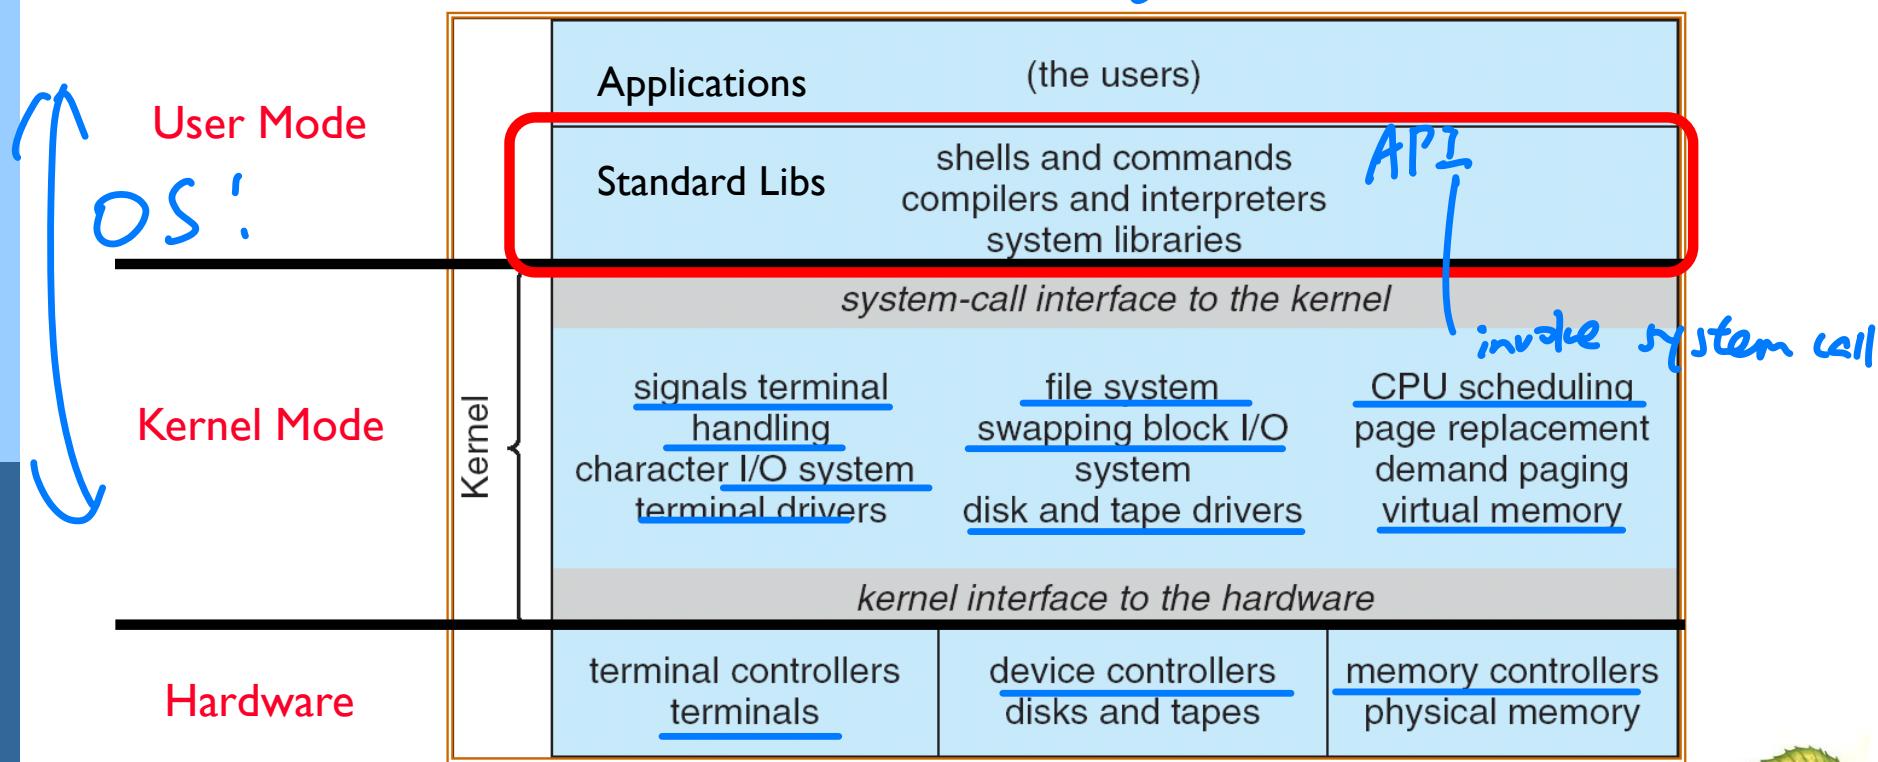
learnt!

Dual-mode operation allows OS to protect itself and other system components

- User mode and kernel mode - to distinguish between the execution of operating-system (kernel) code and user-defined code

!

Protect!





Dual-mode Operation (Cont.)

不同迷路的

- The **dual-mode operation** provides the basic means of protecting the operating system from errant users and errant users from one another
- A **mode bit** provided by hardware - it provides **ability to distinguish when** the system is **running user code** (mode bit=1) or kernel code (mode bit=0)
- Some instructions designated as **privileged**, which can **only be executed in** the **kernel mode**, e.g., I/O control, timer management, and interrupt management
mode = 1 : User mode = 0 : Kernel!
- If an attempt is made to execute a privileged instruction in user mode, the hardware treats it as illegal and **traps** it to the operating system → **switch to**
TRAP!
- The concept of dual modes can be extended beyond two modes.
kernel mode to switch it!
 - For example, Intel processors have **four separate protection rings**, where ring 0 is kernel mode and ring 3 is user mode. Though rings 1 and 2 could be used for various operating-system services, in practice they are rarely used
 - ARMv8 systems have seven modes. CPUs that support virtualization frequently have a separate mode to indicate when the virtual machine manager (VMM) is in control of the system. In this mode, the VMM has more privileges than user processes but fewer than the kernel (to be discussed in Chapter 18)

kernel mode have high privilege!

Mode switch/transfer → switch between kernel/user mode



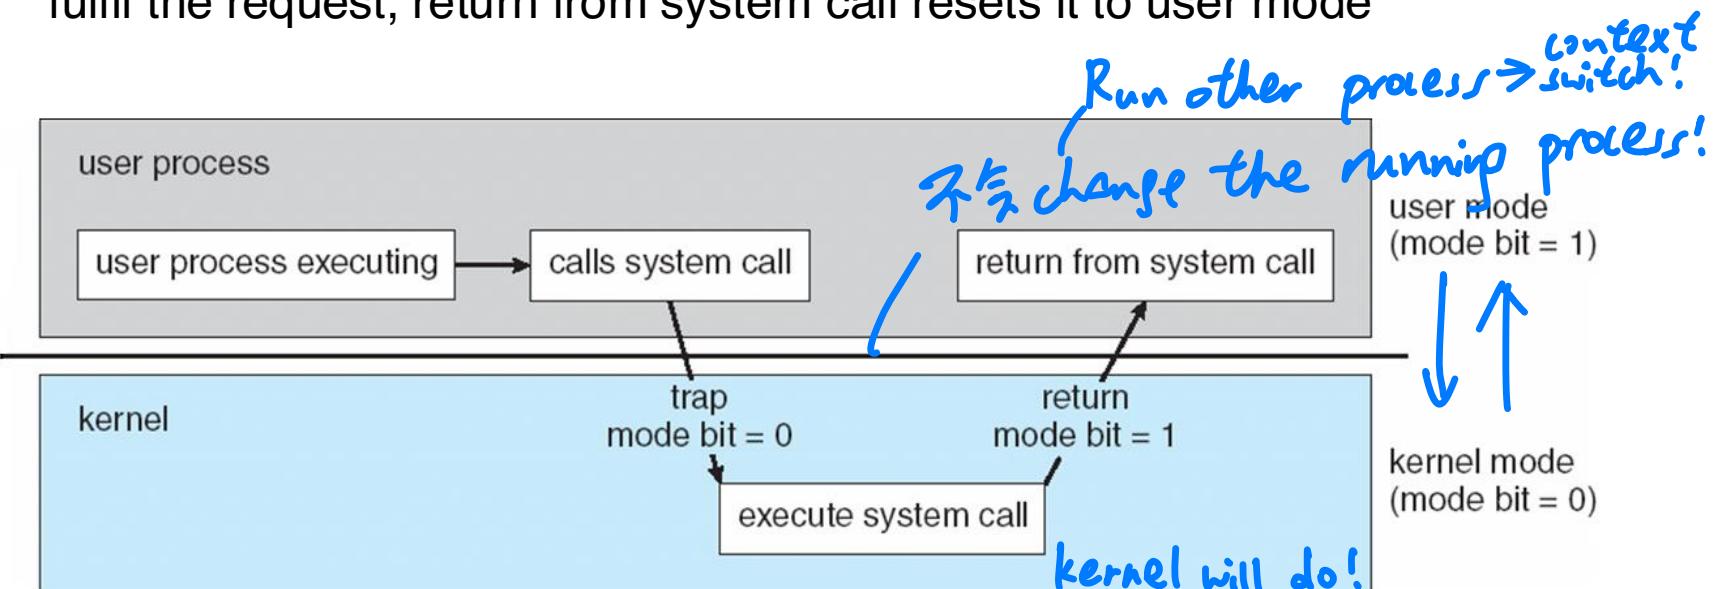


Dual-mode Operation (Cont.)

System calls

kernel provide services

- The initial control resides in the operating system - kernel mode. When control is given to a user application, the mode is set to user mode. Eventually, control is switched back to the operating system via an interrupt, a trap, or a system call
(in kernel)
 - Can we make sure that a user program always returns to OS - **timer**
- When a user program requests a service from the operating system (via a **system call**), the system must transition from user mode to kernel mode to fulfil the request, return from system call resets it to user mode



Do mode switch → sometimes do context switch,
sometimes no!





Three Types of Mode Transfer

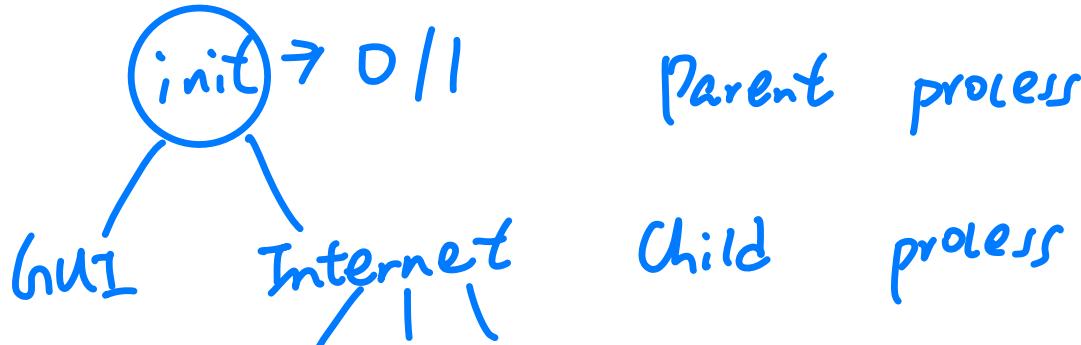
- System call → Kernel mode!
 - A user process requests a service from operating system, e.g., exit(), write(), read(),
 - Like a function call, but “outside” the user process
 - Interrupt 不关 user 事 ! $P_1 \xrightarrow{\quad} P_2$ trigger context switch!
 - External asynchronous events trigger context switch, e.g., timer, I/O
 - Independent of user process *Interrupt is received by kernel space!*
 - Trap or Exception
 - Internal synchronous events in process trigger context switch
 - E.g., protection violation (segmentation fault), algorithmic error (divided by zero), etc.
- Due to error! Go to kernel mode to handle error!





Operations on Processes

- The processes in most computer systems execute concurrently, and they are created and deleted dynamically. Thus, the operating systems must provide mechanisms for:
 - ① Process creation – to run a new program
 - ② Process termination – program execution completes or be terminated
- In general, a **parent** process can create **children** processes, which, in turn, can create other child processes, forming a tree of processes
 - Commonly, the desktop itself can be a parent process, double click an icon (mouse or touch screen), or input a valid command on a Shell, it starts executing a new process, i.e., a child process of the desktop process (i.e., a parent process)





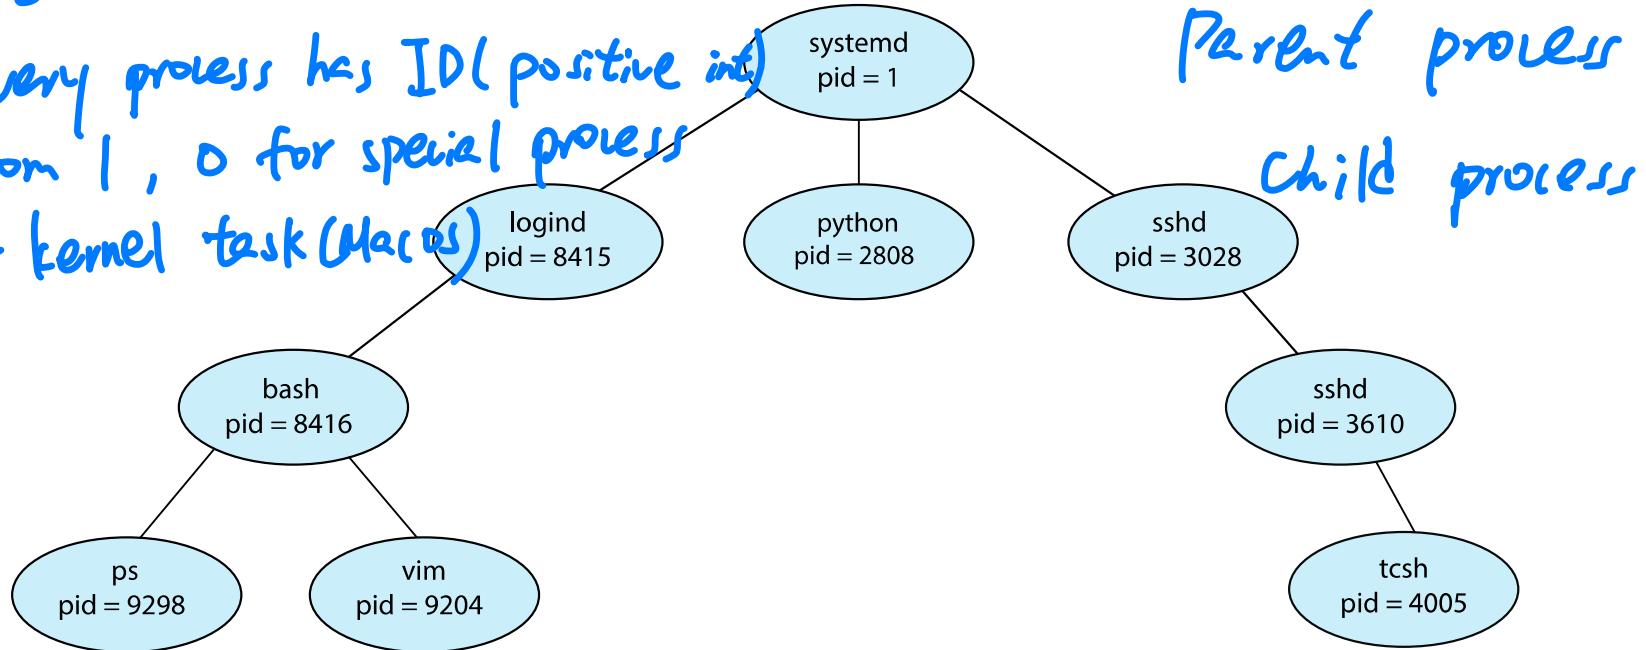
A Tree of Processes on Linux System

- A tree of processes on a typical Linux system

Process \leq Process $\leftarrow \dots$

- Organizing a tree!

- Every process has ID (positive int) from 1, 0 for special process
0 = kernel task (Mac OS)





have different system calls

Process Creation

Parent, child share resources → we fork!

- Most operating systems (e.g., UNIX, Linux, Windows) identify a process according to a unique **process identifier or pid**, typically an integer
- When creating a child process, the child process needs certain resources - CPU time, memory, files, I/O devices to accomplish its task. The parent process may also pass along initialization data (input) to the child process
- During a process creation, there can be **three choices** how resources can be shared (if any) between a parent process and its child process
 - Parent and children almost share all resources - fork() *Everything will be copied!!!*
 - Children share only a subset of parent's resources – clone() (in Chapter 4) *Modern Linux copy will be deleted*
 - Parent and children share no resources
- There are typically **two options** for program execution after process creation
 - The parent continues to execute concurrently with its children *因为很多 overhead!*
 - The parent may wait until some or all of its children have finished execution.
 - For example - Unix shell process waits for a.out process to complete or shell process and a.out& process concurrently execute

p₁ p₂

control how many copied

everything → fork()
part → clone()
copy nothing (meaningless)

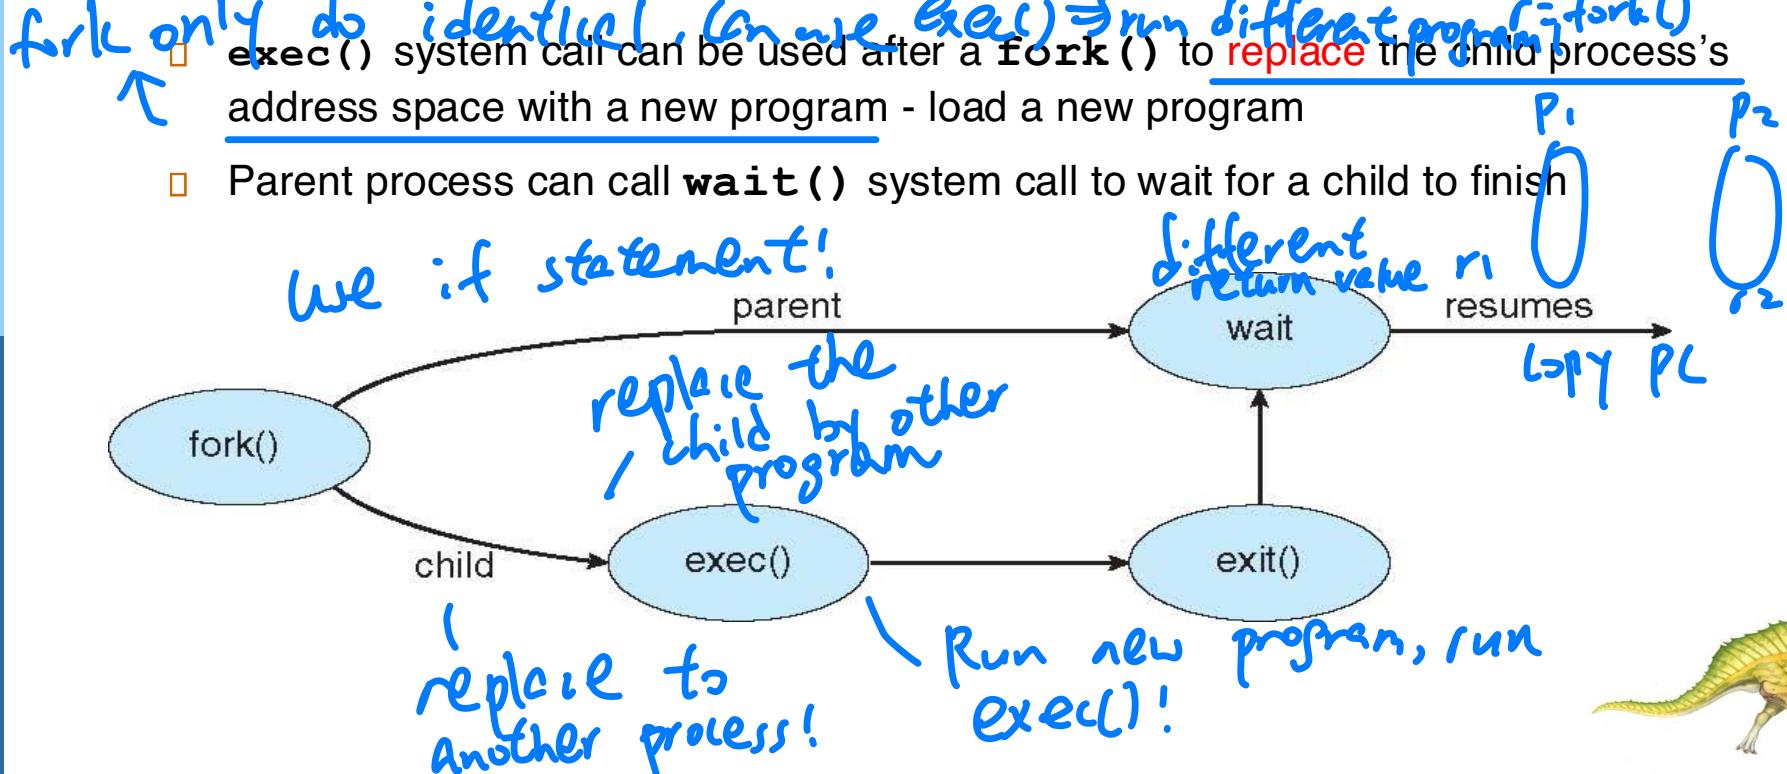


Process Creation (Cont.)

copy = delay = overhead! ??

↓ -fork()
copy everything
but different

- Consider two choices of address space for a new child process
 - A child process duplicates of parent (it has same program and data) – Unix/Linux
 - The child process has a new program loaded into it (different program and data)
- The UNIX example - **fork()**
 - **fork()** creates a new process, which **duplicates** entire address space of the parent. Both processes continue execution at the next instruction after **fork()**
 - **exec()** system call can be used after a **fork()** to **replace** the child process's address space with a new program - load a new program
 - Parent process can call **wait()** system call to wait for a child to finish



`fork()` parent and child

Duplicated

- Address space
- Global & local variables
- Current working directory
- Root directory
- Process resources
- Resource limits
- Program Counter PC

Different

- PID
- Return value from `fork()`
- Running time
- Running state



Process Creation (Cont.)



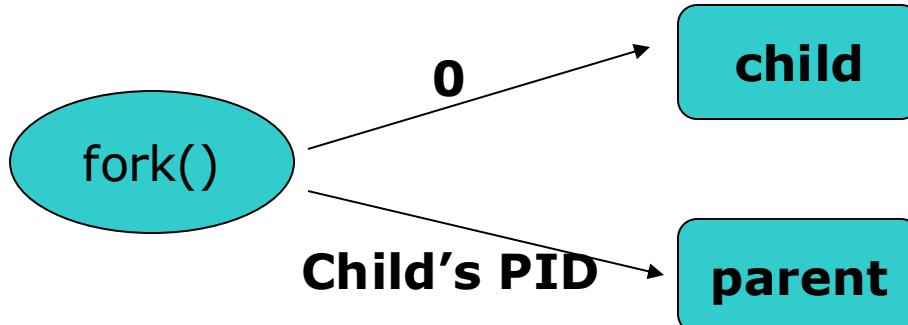
- **fork()** system call creates a **copy** of current process with a new **pid**
- Parent resumes execution after **fork()** and the child also starts execution with **the same program counter**, where the call to **fork()** returns
- The return value from **fork()** – two different integer values, one for parent, one for child; in another word, each process receives exactly one return value from **fork()**, respectively.
- The return value can be:
 - When > 0 : running in (original) parent process ✓
 - The return value (positive integer) is **pid of the newly created child**
 - When $= 0$: running in newly created child process ✓
 - When < 0 : (When = -1) running in original process *will not talk this case!*
 - Error! Must handle somehow – child process not created successfully
- All state of original (parent) process **duplicated** in the child process
 - Memory (program and data), file descriptors, and etc...
- Because the child is a **copy** or **duplication** of the parent, each process has its own copy of the program and all the data – time consuming





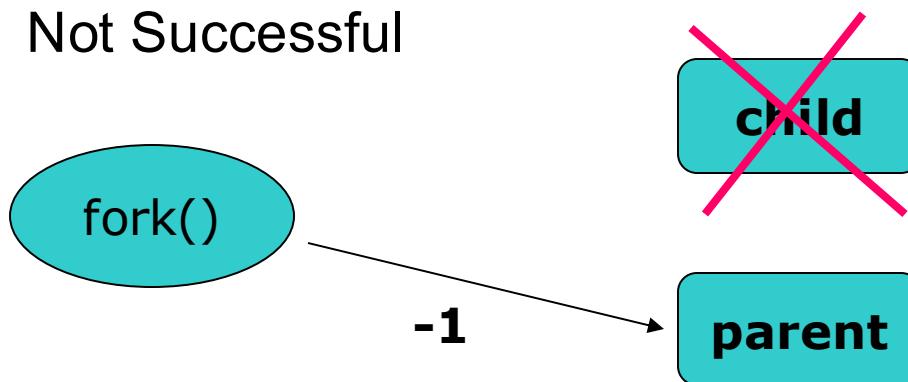
Return values of fork()

Successful



Successful
return value is:
CPID
PID of child process
of parent process!
child process,
return value = 0

Not Successful



P₂ will not
be created

errno is set to indicate error





不重複！

Steps in UNIX fork

Steps to implement UNIX `fork()`

- Create and initialize the process control block (PCB) in the kernel
- Create a new address space – allocating memory
- Initialize the address space with a copy of the entire contents of the address space of the parent – time consuming
- Inherit the execution context of the parent (e.g., any open files)
- Inform the CPU scheduler that the newly created child process is ready to run





Parent vs. Child Processes

process \rightarrow fork() \rightarrow duplicate that process!

After fork() Parent and Child Processes:

Duplicated

- Address space
- Global & local variables
- Current working directory
- Root directory
- Process resources
- Resource limits
- Program Counter - PC
- ...

Different

- PID
 - Return value from fork()
 - Running time
 - Running state
- copy the same PL
exactly same as
that position in PI,
return value is different
- d: child
e: parent
- identity !





fork() example

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

#define BUFSIZE 1024
int main(int argc, char *argv[])
{
    process id.
    char buf[BUFSIZE];
    size_t readlen, writelen, slen;
    pid_t cpid, mypid;
    pid_t pid = getpid(); /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork(); — parent process! Will be run
    if (cpid > 0) { /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) { /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else { // minus value!
        perror("Fork failed");
        exit(1);
    }
    exit(0);
}
```

return child pid

parent process! Will be run

Child process!





System Calls

fork() ↴

□ `exec()`, `execvp()`...

- system call to change the program being run by the current process
- creates a new process image from a regular executable file
- `wait()` – system call to wait for a (child) process to finish or on general
wait for an event CPID > 0 (parent process)
- `exit()` – system call to terminate the current process, and free all its resources
- `signal()` – system call to send a notification to another process
- More details in UNIX man pages and in tutorial

`wait(NULL)` = wait() will block the parent process until any of its child process has finished!

`wait(pid)` = wait until specific child process!





C Program Forking Separate Process

```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
        'wait until this complete'
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL); ←不想 end early!
        printf ("Child Complete");
        exit(0);
    }
}
```





Implementing a Shell

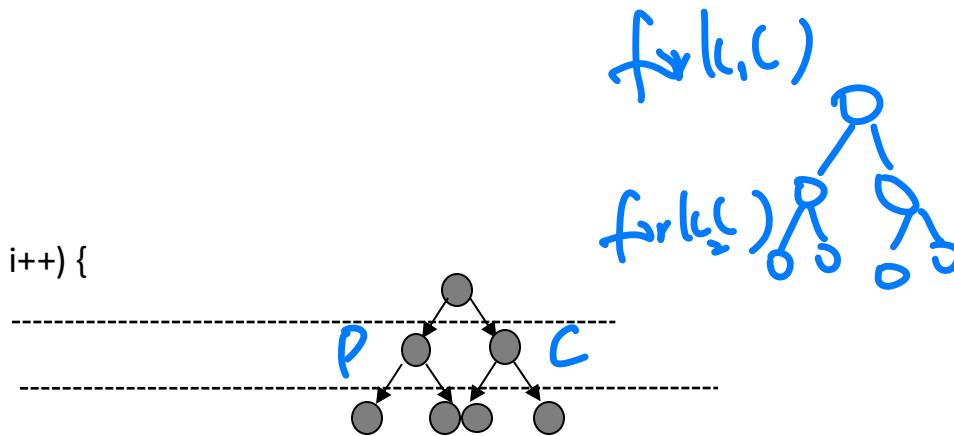
```
char *prog, **args;  
int child_pid;  
  
// Read and parse the input a line at a time  
while (readAndParseCmdLine(&prog, &args)) {  
    child_pid = fork(); // create a child process  
    if (child_pid == 0) {  
        exec(prog, args); // I'm the child process. Run program  
        // NOT REACHED  
    } else {  
        wait(child_pid); // I'm the parent, wait for child  
        return 0;  
    }  
}
```





C Program Forking Separate Process

```
int main()
{
    for (int i = 0; i < 10; i++) {
        fork();
        fork();
    }
    return 0;
}
```



duplicating address space.
also PL

Illustration:

- Totally there are **10** iterations.
- Inside the loop
 - 2 fork() function calls create **4 processes**
- Finally, there are **4^{10}** processes. ✓



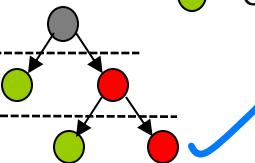


C Program Forking Separate Process

(How many process?)

```
int main()
{
    for (int i = 0; i < 10; i++)
        if (fork())
            fork();
    return 0;
}
```

- - Parent process after fork()
(return value > 0)
- - Child process after fork()
(return value = 0)



only parent will call fork()
since return value ≠ 0

Illustration:

- Totally there are 10 iterations.
- One process creates 2 child processes after each iteration.
- Finally, there are 3^{10} processes.





Final Notes on fork()

- Process creation in Unix is unique. Most operating systems creates a process in a new address space, read in an *executable file*, and begin executing it. Unix takes an unusual approach of separating these steps into two distinct system calls: **fork()** and **exec()**
replace child process / want to create other
- In Linux, **fork()** is implemented via the use of copy-on-write pages. Copy-on-write or COW (to be discussed in Chapter 10) is a technique to delay or prevent copying of data. Rather than copy the address space, the parent and the child can share a single copy, i.e., the child process “points to” the parent process address space without duplication
只改寫 copy
- This delays the copying in the address space until the content is changed or written to. In the case that the pages are never written — for example, if **exec()** is called immediately after **fork()** — they never need to be copied
不會寫的 just we different pointers!
- This greatly simplifies and speeds up the process creation
- Linux also implements **fork()** via the **clone()** system call, which is considered to be more general. **clone()** system call uses a series of flags that allows to specify which set of resources, if any, the parent and child processes should share (to be discussed in Chapter 4)
不修改 → no really copy!





Process Termination

*runtime will automatically
do the t!!!*

- After a process executes the last statement and it asks the operating system to delete it using the **exit()** system call.
 - The process may return a status value (typically an integer) to its waiting parent process (via **wait()** system call from the parent process)
 - Most of the resources including memory (program, data, heap and stack), open files, and I/O buffers – are de-allocated and reclaimed by operating system
- Parent may terminate the execution of children processes using the **abort()** in Unix and Linux, **TerminateProcess()** in Windows
- Notice that a parent needs to know the **identity** or **PID** of its child during termination. Interestingly, when creating a child process, the identity of the newly created child process is passed to the parent with **fork()**
- A variety of reasons for a parent to terminate the execution of a child
 - The child has exceeded allocated resources (parent can inspect the child state)
 - Task assigned to child is no longer required
 - The parent is exiting, and some operating systems do not allow a child to continue if its parent terminates





Process Termination (Cont.)

~~Terminate parent → also terminate child processes~~

- Some OSes do not allow child to exist if its parent has terminated. If a process terminates, all its children must also be terminated. This is referred to as cascading termination, initiated by the operating system
- Under normal termination, **exit()** will be called either directly or indirectly, as the C runtime library includes a call to **exit()** by default
- The parent process may wait for termination of a child process by using the **wait()** system call
- The **wait()** is passed a parameter that allows the parent to obtain the exit status of the child, as well as the process identifier of the terminated child so that the parent can tell which of its children has terminated

```
pid = wait(&status);
```





Process Termination (Cont.)

- When a process terminates, its resources are deallocated by the OS.
- However, there are a few clean up that needs to be done by the parent and OS. For instance, its entry in the process table (a kernel structure that keeps track of all processes in a system) remains until the parent calls `wait()`, as the process table or process list entry contains the process's exit status
PBS will be cleaned!
- A process that has terminated, but whose parent has not yet called `wait()`, is known a **zombie process** – its corresponding entry in the process table, process ID, and PCB still exist. All other resources allocated to the process have been released. So, it is not runnable as it does not have an address space (program and data have all been deleted)
zombie!
- All processes transition into this **zombie state** before termination, but generally they exist as zombies only very briefly
call wait frequently!!!
 - Once the parent calls `wait()`, the process identifier of the zombie process and its corresponding entry in the process table and PCB are released
- Such a design enables the parent, which in turn, informs the OS that a child process has terminated, so to trigger the final cleaning up





Process Termination (Cont.)

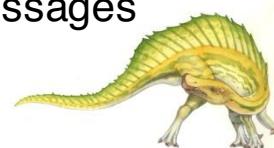
- If a parent terminates without invoking `wait()`, child process becomes an orphan – this process may still be runnable (if cascading is not used) or eventually it will become a zombie process. Some mechanism must exist to reparent such orphan child processes to a new process. Otherwise, parentless terminated processes would forever remain zombies, wasting system memory.
- The `init` or `systemd` is assigned as a parent that periodically invokes `wait()`, allowing exit status of any orphaned process to be collected and releasing orphan's process identifier and the corresponding entry in the process table





As processes are isolated!

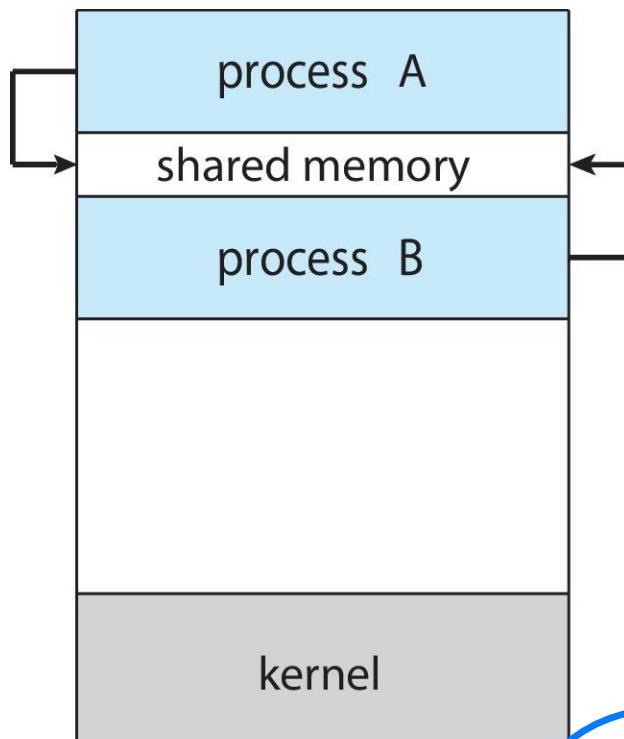
- Processes within a system may be **independent** or **cooperating**
- Cooperating processes can **affect or be affected by other process**, for instance two or more processes share data
- Reasons for processes to cooperate, and also **naturally for a multi-threaded process**:
 1. □ **Information sharing** - several applications may be interested in the same piece of information (for instance, copying and pasting)
 2. □ **Computation speedup** – executing parallel tasks or subtasks, for instance machine learning tasks
 3. □ **Modularity** – functions divided into separate processes or threads
- Cooperating processes require an **interprocess communication (IPC)** - a mechanism that will allow them to exchange data
- There are generally **two models of IPC**
 - **Shared memory** → a region of memory is shared, and cooperating processes can exchange information by reading and writing data to the shared region
 - **Message passing** → explicit communication takes place by means of messages exchanged between cooperating processes





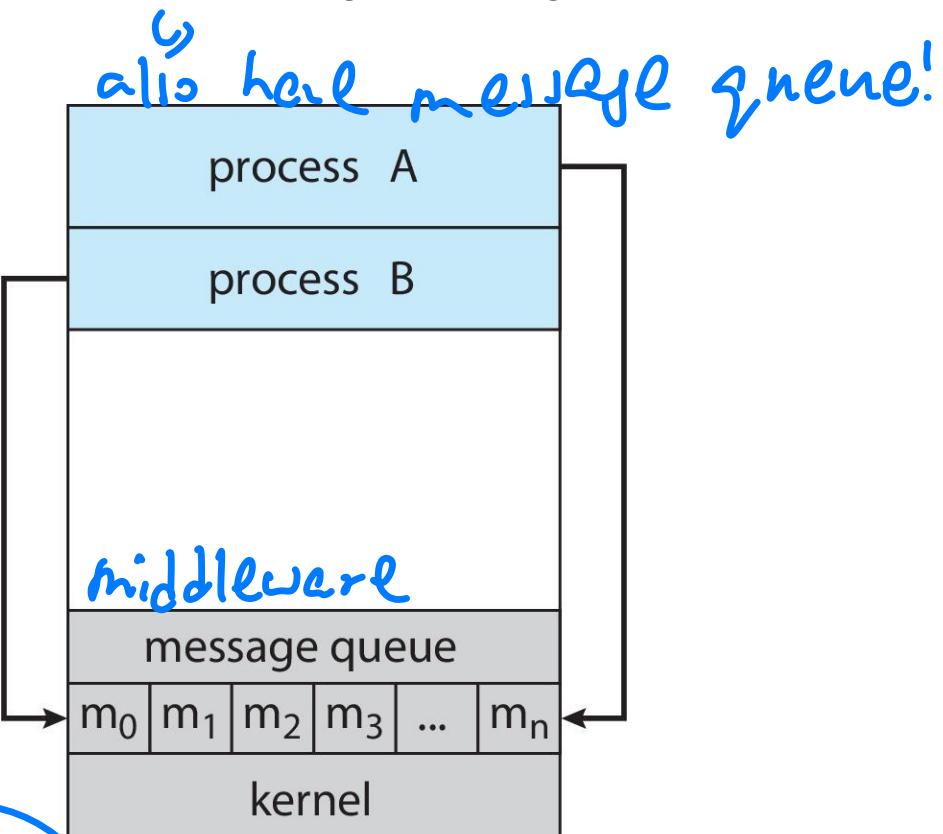
Communication Models

(a) Shared memory



(a)

(b) Message passing



(b)

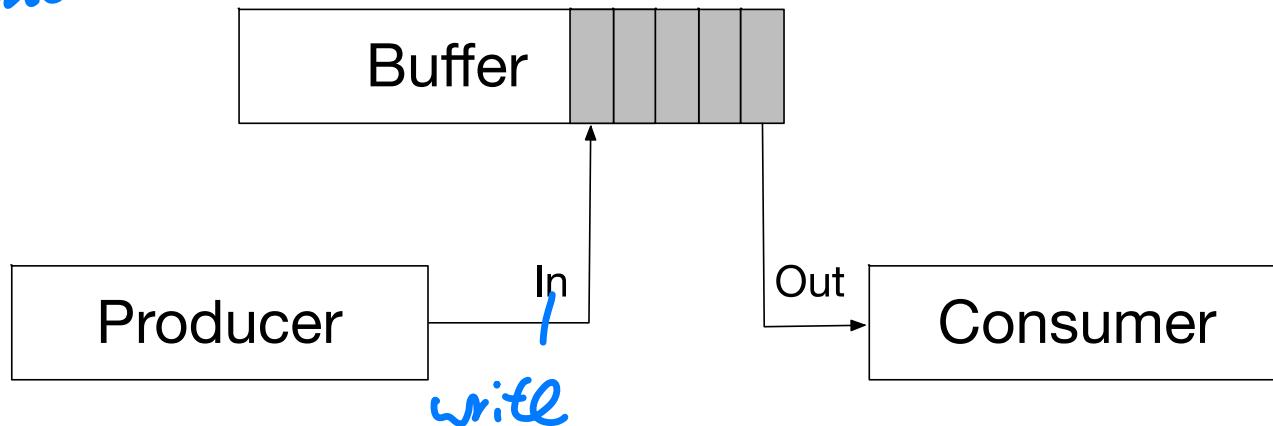




Producer-Consumer Problem

- One paradigm for cooperating processes – **Producer** and **Consumer** represents a common type of cooperation between processes (We will discuss other type of cooperations in Chapters 6 and 7)
- An example: a **producer** process produces information that is consumed by a **consumer** process
 - unbounded-buffer places **no practical limit** on the size of the buffer
 - bounded-buffer assumes that there is a **fixed buffer size**

modern OS, limited space!





Bounded-Buffer – Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10

typedef struct {

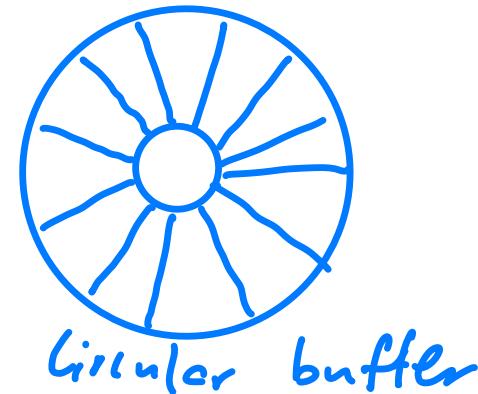
    . . .

} item;

item buffer[BUFFER_SIZE];

int in = 0;

int out = 0;
```



- The above solution is simple and correct, but only allows at most **BUFFER_SIZE-1** items in the buffer at the same time.
- Different implementations can use all **BUFFER_SIZE** items – easily fixed





Producer Process – Shared Memory

Buffer size - 1

```
item next_produced;  
while (true) {  
    /* produce an item in next_produced */  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE; /* move ptr  
forward */  
}
```

0 ... ⑤
10 items!

if out = 0,
infinite loop!
不要用 index!





Consumer Process – Shared Memory

```
item next_consumed;

while (true) {
    while (in == out)

        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}
```

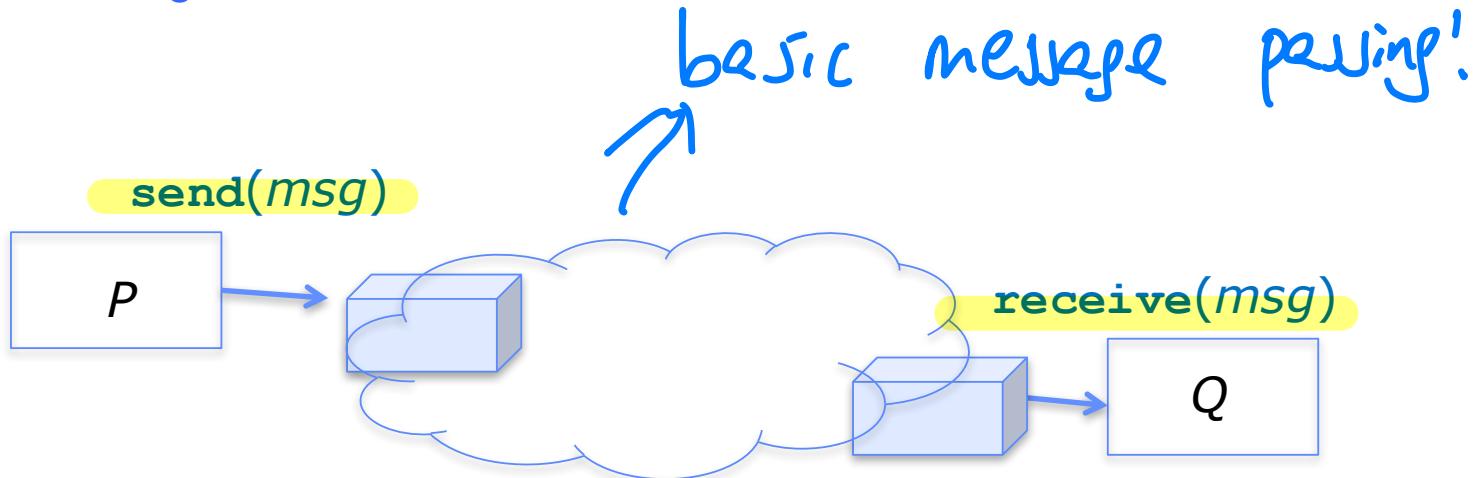
in/out \Rightarrow determine
whether index is full/empty!





Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions without sharing the address space
- It is particularly useful in a distributed environment
- IPC facility provides two operations:
 - `send(message)`
 - `receive(message)`
- The `message` size could be either fixed or variable





Message Passing (Cont.)

FIFO queue?

- If processes P and Q wish to communicate, they need to:
 - Establish a communication link between them
 - Exchange messages via primitives send() and receive()
- There exist different methods in implementation of a link and send() and receive() operations
 - 1. □ Direct or indirect communication
 - 2. □ Synchronous or asynchronous communication
 - 3. □ Automatic or explicit buffering
- **Naming**: processes that want to communicate must have a way to refer to each other. They can use either direct or indirect communication

message passing!

¹
name = pid





Direct Communication

link

- Processes must name each other explicitly:
 - **send** (P , message) – send a message to process P
 - **receive**(Q , message) – receive a message from process Q
- Properties of the communication link
 - A link is established automatically between every pair of processes that want to communicate. The processes need to know **only** each other's identity to communicate.
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
- The **disadvantage** in direct communication is the limited modularity of the resulting process definitions. For instance, changing the identifier of a process may necessitate examining all other cooperating processes – with the new process identifier

one change → all should change the pid!





Indirect Communication

$P_1 \rightarrow P_{11}$

P_2/P_3 exchange the identify
id!

- Messages are sent to and received from mailboxes

- Each mailbox has a unique id

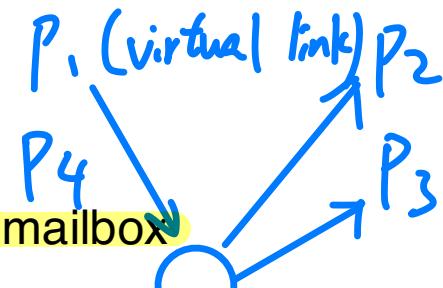
- Processes can communicate only if they share a mailbox

- Properties of communication link

- A link is established between a pair of processes only if both members of the pair have a shared mailbox.

- A link may be associated with more than two processes

- Between each pair of communicating processes, a number of different links may exist, with each link corresponding to one mailbox



mailbox (unique id)
(more stable)

proxy / agent

= someone in the middle

Buy house not directly from owner





Indirect Communication (Cont.)

- The operating system then must provide a mechanism for

- Create a new mailbox
 - Send and receive messages through the mailbox
 - Destroy a mailbox

change the

-function!

≤ 2 processes !

- Primitives are defined as:

send (A, message) – send a message to mailbox A

receive (A, message) – receive a message from mailbox A

- Mailbox sharing

OS can control who receive !

- P_1, P_2 , and P_3 share mailbox A, P_1 , sends; P_2 and P_3 receive?

- Solutions

- Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a **receive ()** operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.





Synchronization

1. Direct / Indirect
2. Syn / Asyn!

- Message passing may be either **blocking** or **non-blocking**, also known as **synchronous** and **non-synchronous** 
- **Blocking** is considered **synchronous** 

xxx
f()
xxx

for heavy task(s)
waiting
functions
- **Blocking send:** The sending process is blocked until the message is received by a receiving process or by a mailbox
- **Blocking receive:** The receiver blocks until a message is available 
(return immediately)
- **Non-blocking** is considered **asynchronous** 

Blocking! → put to another thread
⇒ f() will ensure the whole program to wait!
- **Non-blocking send:** The sending process sends the message and resumes its operation without waiting for the message reception → immediately return
- **Non-blocking receive:** The receiver retrieves either a valid message or null 

need check
whether is empty?
continuous check!

Event driven

f() → call back / interrupt!
after f()
will execute following task:

trigger other threads!

callback() ...



- \bar{w} non-blocking send, blocking receive
- 自己 implement



Synchronization (Cont.)

- Different combinations of `send()` and `receive()` possible. When both `send()` and `receive()` are blocking, we have a **rendezvous** between the sender and the receiver
- Producer-consumer problem solution becomes **trivial** *no queue!*

```
message next_produced;  
while (true) {  
    /* produce an item in next_produced */  
    send(next_produced);  
}  
  
message next_consumed;  
while (true) {  
    receive(next_consumed);  
    /* consume the item in next_consumed */  
}
```

Zero capacity!

Problem to be easy:

*no buffer,
no queue!*

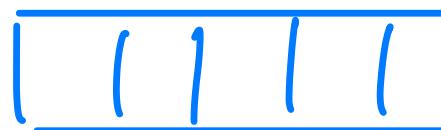
*P1
send()
while loop!
P2
receive*

*Send one by one!
ensure P2 can
receive this message!*





3. how to design! Buffering

- ① Whether communication is **direct** or **indirect**, messages exchanged by communicating processes reside in a temporary queue. Such queues can be implemented in three ways:
 - ② ~~blk/unblk!~~
1. Zero capacity – no messages are queued on a link.
Sender must wait for receiver (rendezvous)
2. Bounded capacity – finite length of n messages
Sender must wait if link is full
3. Unbounded capacity – infinite length
Sender never waits
- printf("Hello world")
→ fork()
return 0;*
- 

Size = 0 // empty

Size = buffer_size // full

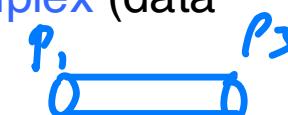




ordinary pipe !

Pipes

message passing system

- **Pipes** - one of the first IPC mechanisms in early UNIX systems. There are four issues to be considered when implementing a `pipe()`
pipe, two ends
- Is communication unidirectional or bidirectional?
only one direction
- In the case of two-way communication (i.e., bidirectional), is it half duplex (data can travel only one way at a time) or full-duplex (data can travel in both directions at the same time)?
half duplex

- Must there exist a relationship (i.e., parent-child) between the communicating processes or between any two processes?
- Can the pipes be used over a network, or only on the same machine?
- **Ordinary Pipes** – cannot be accessed from outside the process that creates it. Typically, a parent process creates a pipe and uses it to communicate with a child process that it creates via `fork()`.
- **Named Pipes** – can be accessed without a parent-child relationship.



Ordinary Pipes

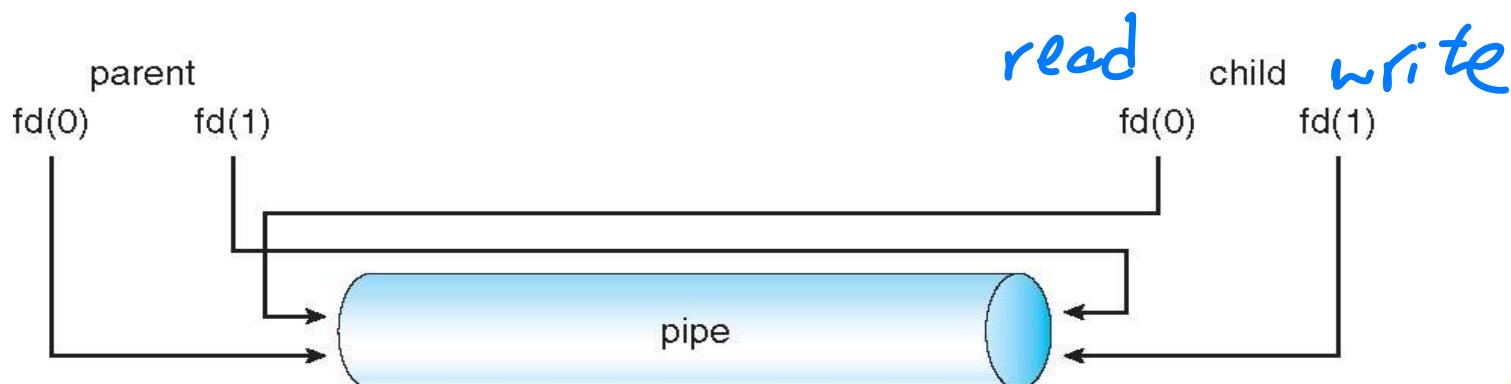
- only one direction
- only can exist in parent-child process!



Unix / Linux / Anymous ↑ Ordinary Pipes

vector of two elements!

- Ordinary pipes communicate in a standard *Producer-Consumer* fashion - **pipe(int fd[])** in UNIX – unidirectional, allowing one-way communication
 - Producer writes to one end (the **write-end** of the pipe) **fd[1]**
 - Consumer reads from the other end (the **read-end** of the pipe) **fd[0]**
- This requires a parent-child relationship between two communicating processes running on the same machine
- An ordinary pipe cannot be accessed from outside the process that creates it
- An ordinary pipe ceases to exist after processes have finished communicating and terminated



two way pipe → create two of that!





Ordinary Pipes (Cont.)

↗ Everything!

- UNIX treats a pipe as a special type of file, standard **read()** and **write()**, and the child **inherits** the pipe from its parent process
- In UNIX, the parent process uses **pipe()** creates a pipe and then uses **fork()** to create a child process. The child process inherits everything from the parent including the ordinary pipe (a special file). In this case, the parent and the child can communicate with each other using this pipe
- Noticing, if the parent needs to write to the pipe and the child reads from it, the parent must explicitly close the read-end of the pipe, and the child has to explicitly close the write-end of the pipe. **One directional!**
- Ordinary pipes on Windows systems termed **anonymous pipes**, and behave similarly, access through **ReadFile()** and **WriteFile()** functions

Linux → everything in file!

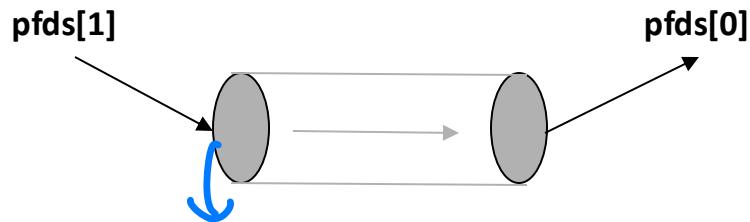




Pipe Example (Parent => Child)

```
int main()
{
    int pfds[2];
    char buf[30];
    pipe(pfds); /* Create a message pipe*/
    pid_t pid = fork(); /* 0 (child), non-zero (parent) */
    if ( pid != 0 ) {
        printf("PARENT: writing to pipe\n");
        close(pfds[0]); -close read site
        write(pfds[1], "test", 5);
        wait(NULL); /* Wait until the child returns*/
        printf("PARENT: exiting\n");
    } else {
        printf("CHILD: reading from pipe\n");
        close(pfds[1]);
        read(pfds[0], buf, 5);
        printf("CHILD: read \"%s\"\n", buf);
    }
    return 0;
}
```

only can be used in parent - child relationship



*只能读写！
不能多写！*





Named Pipes

- Named pipes are more powerful than ordinary pipes - communication is bidirectional, no parent-child relationship is required, and several processes can use it for communication, typically with several writers
- Name pipes continue to exist after the communicating processes have finished, until they are explicitly deleted
- *first in first out queue*
Named pipes are referred to as FIFOs in UNIX systems, created using `mkfifo()` system call. They appear as typical files in the file system
 - FIFOs only supports half-duplex transmission. If data must travel in both directions simultaneously, two FIFOs are typically used. Additionally, the communicating processes must reside on the same machine
- Named pipes on Windows provide a richer communication mechanism
 - Full-duplex communication is allowed, and the communicating processes may reside on either the same or different machines
 - Named pipes are created with the `CreateNamedPipe()` function





Communications in Client-Server Systems

X in scope in OS.

two process,

Also can be done in networking! two machines!
rich functions / kernel/user spaces

- **Sockets** Also can be done in networking! two machines!
A socket is defined as an endpoint for communication. A pair of very processes communicating over a network employs a pair of socket strong!!!
 - A socket is uniquely identified by an IP address and a port number specifying the process or application on either client or server side
 - The server (always running) waits for incoming client requests by listening to a specified port (well knowns). Once a request is received, the server can choose to accept the request from the client to establish a socket connection for communication

- **Remote Procedure Calls** or **RPC** sender/receiver have different IP address
 - One of the most common forms of remote services - the procedure-call mechanism for use between systems with network connections function call through network
 - A message-based communication scheme to provide remote services
 - The RPC scheme is commonly used for implementing distributed file systems by providing a set of RPC daemons and clients

- Use IP address to find yr prcs





Client-Server Communication - Sockets

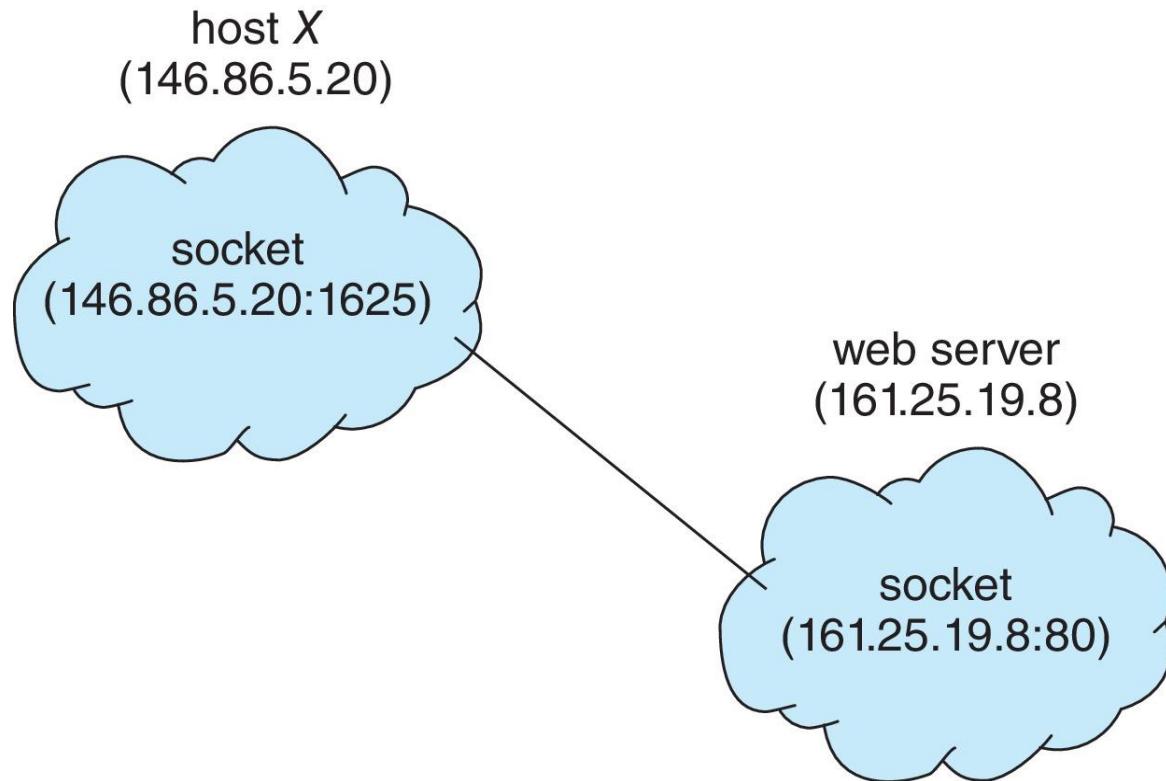
- Servers (such as SSH, FTP, and HTTP) listen to well-known ports (an SSH server listens to port 22; an FTP server listens to port 21; and a web, or HTTP, server listens to port 80). All ports below 1024 are well known and are reserved for standard services - server port numbers
- Server host address (IP address) or name (DNS translates that into an IP address) and server port number are well known. Servers are always ready waiting clients to contact
- Only clients can initiate the communication or send request to a server
- When a client initiates a request for a connection, it knows its own IP address, selects a client port number (greater than 1024), not current in use
- IP address – globally unique ↴ random
 - 128.32.244.172 (IPv4 32-bit)
 - Fe80::4ad7:5ff:fecf:2607 (IPv6 128-bit)
- The 4-tuple (source IP address, source port number, destination IP address, destination port number) uniquely determines a pair of communication

between 2 processes
can be different/same machine
Port number → identify the process!





Socket Communication



Communication between a pair of sockets



\$ - Message passing

- Shared memory

- fork()

- PCB in kernel space!

End of Chapter 3

