

? servl threads \rightarrow fork() will do?

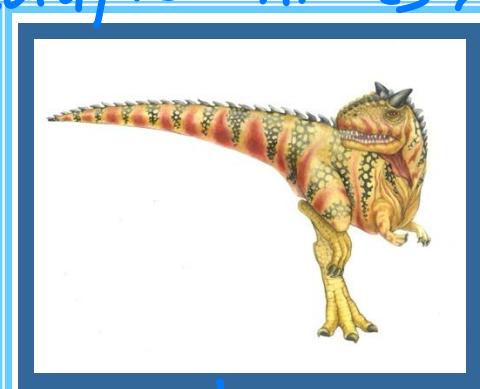
worker: - 1 thread

→uz

Chapter 4: Threads & Concurrency

- process, isolated
- have multiple threads, will share some resources

- fork()
 \downarrow
execve()



\downarrow
run a new process, replace the address space
wait() \rightarrow parent wait for child process finish
- blocking \rightarrow x design explicitly buffer!



Chapter 4: Threads

- Overview
- Multicore Programming
- Multithreading Models
- Threading Issues
- Operating System Examples

how to run program much more fast

one thread \Rightarrow fork()

other thread \Rightarrow ?





Objectives

- Identify the basic components of a thread, and contrast thread and processes.
- Describe the major benefits and challenges in designing multithreaded processes.
- Describe how Windows and Linux operating systems represent threads

*move fast
particular!*

kernel > 100 threads (run parallel)





Motivation

- Many applications or programs are **multithreaded**
 - A web browser might have one thread displaying images or text while another thread retrieving data from the network
one process → data → multiple threads (parallel execution)
 - A word processor may have a thread for displaying graphics, another thread for responding to keystrokes from user, and a third thread for performing spelling and grammar checking in the background
- Most operating system **kernels** are multithreaded
 - For example, during system boot time on Linux, several kernel threads are created. Each thread performs a specific task, such as managing devices, memory management, or interrupt handling
- It can take advantage of **processing capabilities** on multicore systems
 - Parallel programming widely used in applications such as data mining, graphics, and artificial intelligence
involve copying the address space run together!
- **Process creation** is time consuming and resource intensive
 - Process creation is **heavy-weight** while thread creation is **light-weight**, in which different threads belonging to the same process share code, data and others
share the code, local variables, heap





Examples of Multithreaded Programs

不重要！

- Embedded systems
 - Elevators, planes, medical systems, wristwatches
 - Single program, concurrent operations
- Most modern OS kernels
 - Internally concurrent to deal with concurrent requests by multiple users
 - But no protection needed within kernel
- Database Servers
 - Access to shared data by many concurrent users
 - Also background utility processing must be done
- Network Servers *← Middleware!*
 - Concurrent requests from network
 - Again, single program, multiple concurrent operations ✓
 - File server, Web server, and airline reservation systems
- Parallel Programming */ Run faster!*
 - Split a program and data into multiple threads for parallelism

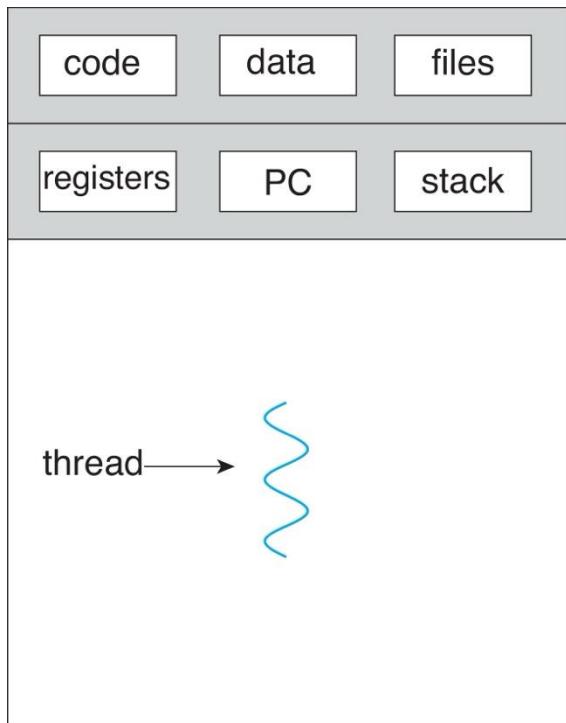
Every program is multithreaded nowadays



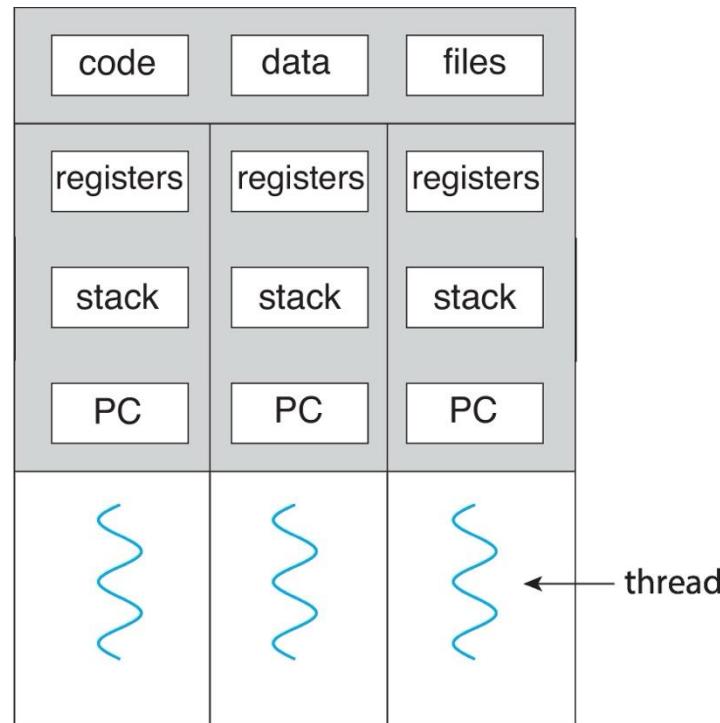


Single and Multithreaded Processes

- Recall that a **thread** is a basic unit of CPU utilization – **independently scheduled and run** (referred as an **instance of execution**) represented by a **thread ID**, a **program counter (PC)**, a **register set**, and a **stack**. It shares with other threads of the same process its code, data, and other OS resources, such as open files and signals

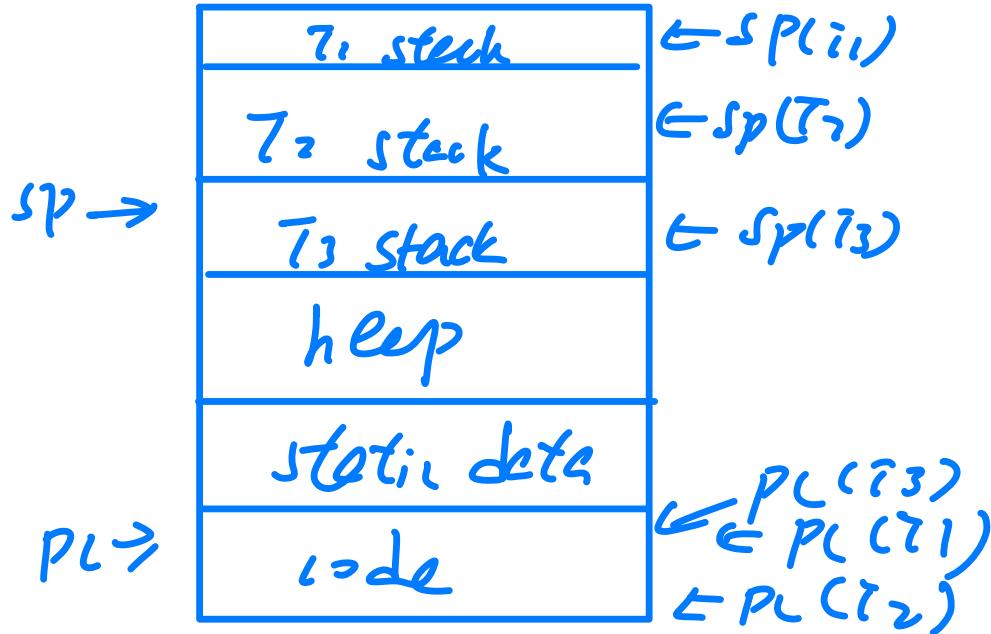


single-threaded process



multithreaded process





- Different PC in different threads
- Share the same code (address space)
- have own thread!
- 不同 Sp



Benefits

one process → more responsible

- ① **Responsiveness** – may allow continued execution if part of process (e.g., one thread of a process) is blocked, this could be especially important for interactive applications having user interfaces
- ② **Resource Sharing** – threads ^{2 or 3} share resources of a process by default, easier than shared memory or message passing between processes – as they essentially run within the same address space of the process
- ③ **Economy** – ^{fork()} thread creation is much “cheaper” than process creation (consuming less time and memory), context switching is typically faster between threads of a process than between processes
- ④ **Scalability** – processes can take advantage of multicore architecture

Multiple core → run together → fits

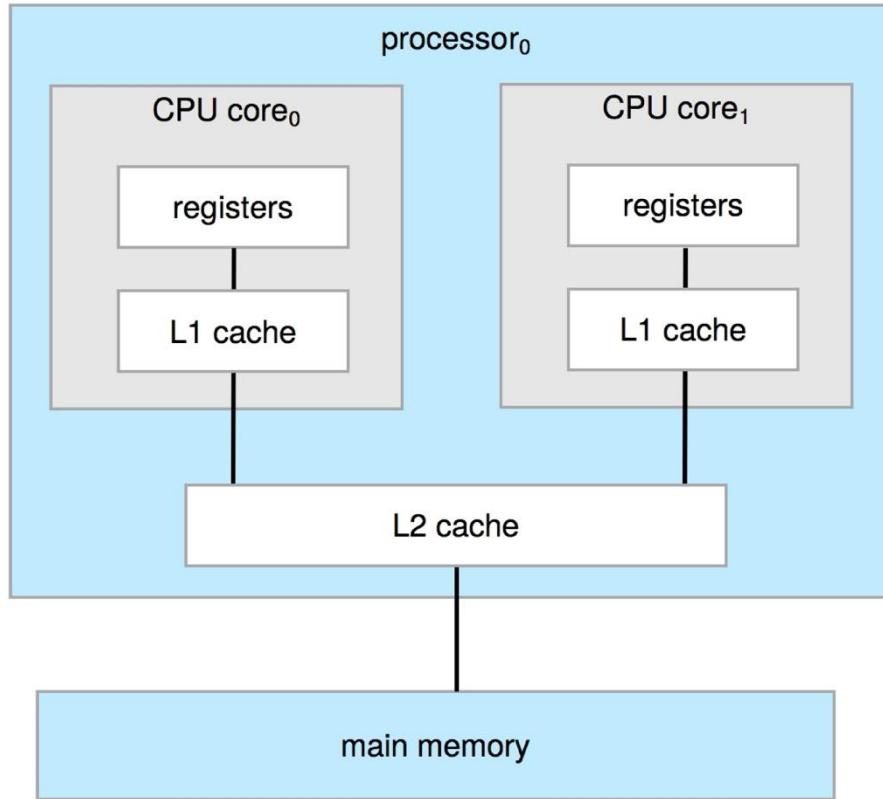




A Multi-Core Design

- The multithreaded programming provides a mechanism for more efficient use of multicores and improves concurrency in **multicore** systems.

distribute 2 threads! Perfect fits!





Multicore Programming

follow principles!

- In multicore or multiprocessor systems, there are significant and new challenges in programming design include:
 - Dividing tasks – how to divide into separate, concurrent tasks
 - Balance – each task perform “equal” amount of work
 - Data splitting - data used by tasks must be divided to run on separate cores
 - Data dependency – if there is a data dependency, synchronization is required
 - Testing and debugging – different path of executions makes debugging difficult
- How to speed up:
 - There is a clear distinction between parallelism and concurrency
 - Parallelism → run multiple part of yr program simultaneously → same time
 - Refer to the state → parallel all progress
 - Concurrency supports more than one task for making progress
 - Multiplexed over time, single processor core, scheduler provides concurrency
- The advent of multicore systems require an entirely new approach in designing software systems with particular emphasis on parallel programming.

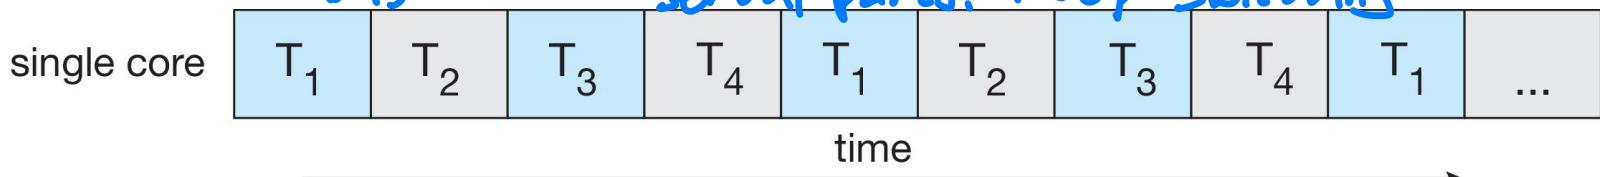




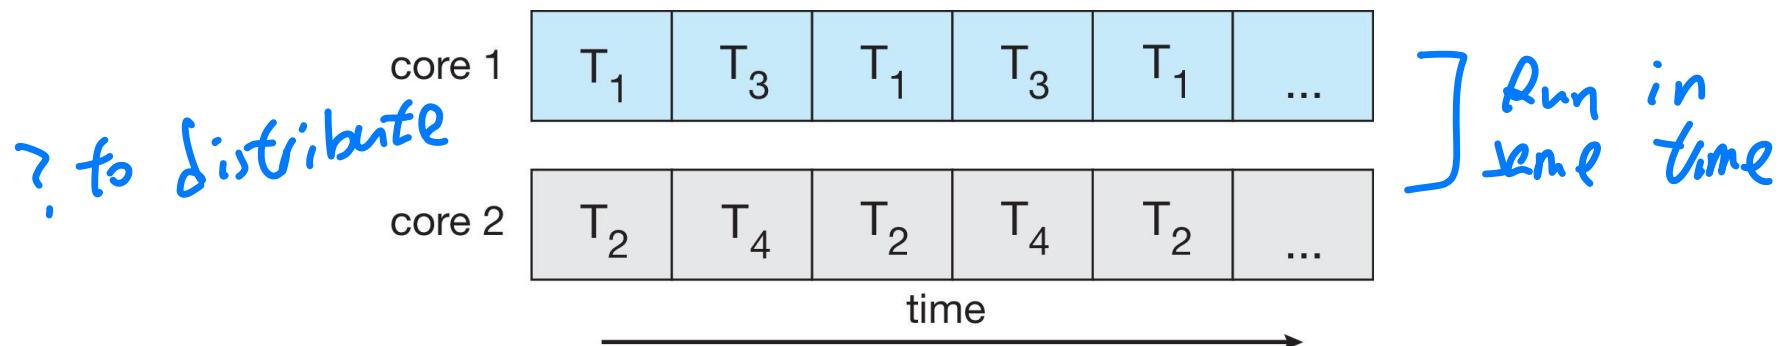
Concurrency vs. Parallelism

- Concurrent execution on single-core system – multiplexing over time

Distribute several parts! Keep switching



- Parallel execution on a multi-core system:



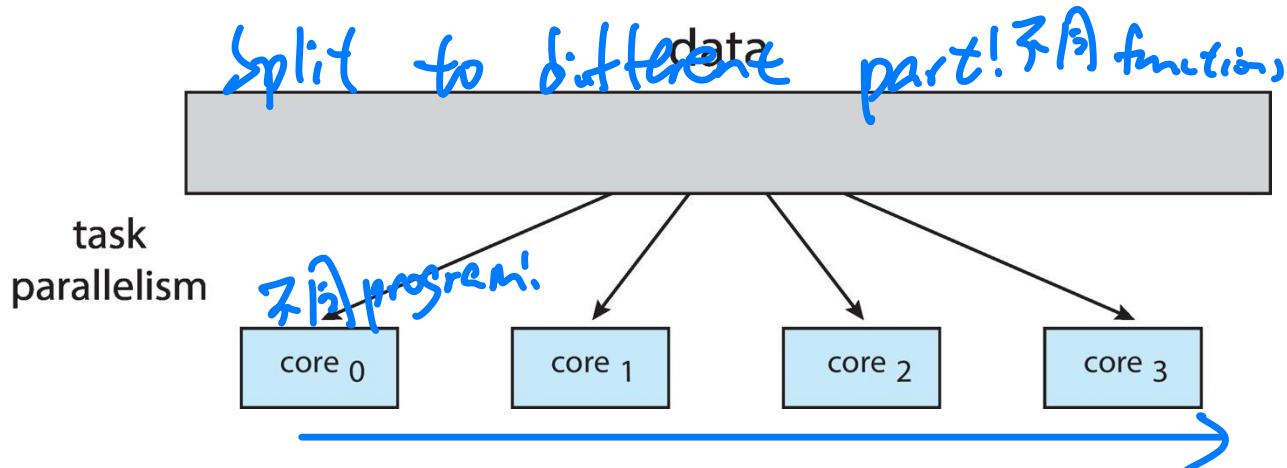
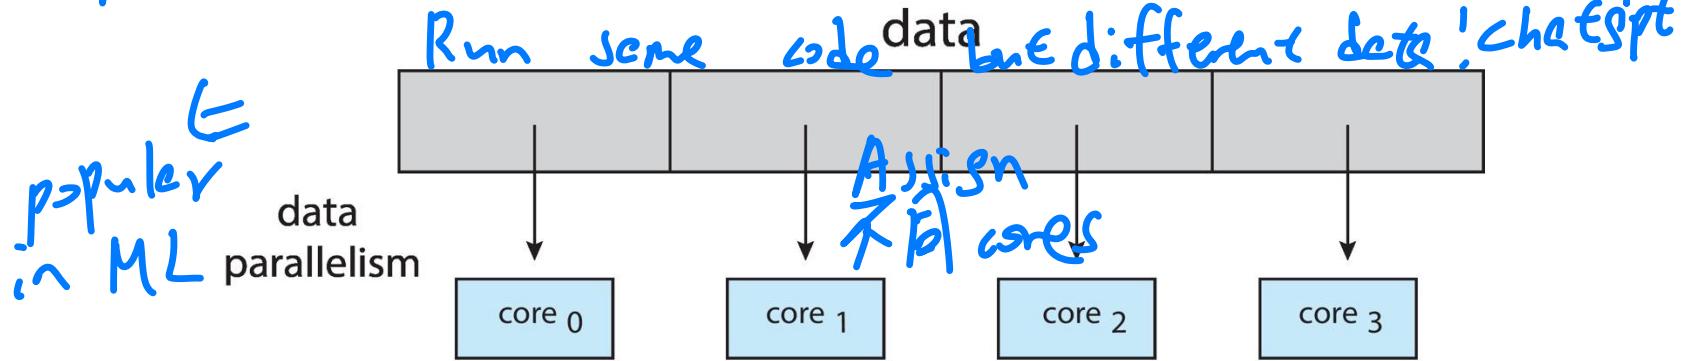


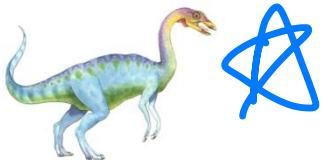
Data and Task Parallelism

very common Same code → but different data!

- **Data parallelism** – distributes subsets of the data across multiple cores, same operation on each core (common in distributed machine learning tasks) *distribute data to GPUs, GPU run the same code!*
- **Task parallelism** – distributing threads across cores, each thread performing unique operation *divide program to different tasks, run one by one*
- Data and task parallelism are not mutually exclusive, an application may use both - hybrid *run different function!*

Example:





Amdahl's Law

*1/20 → 10 groups!
decrease parallelism*

- It identifies performance gains (theoretical speedup in latency of the execution of a task at fixed workload) from adding additional cores to an application that has both serial and parallel components in programs
- S is **serial portion** and $1-S$ is **parallel portion**
- N processing cores

; ideal, have overhead!

$$\text{speedup} \leq \frac{1}{S + \frac{(1-S)}{N}}$$

$S \rightarrow 0 \text{ to } 1$

$$\frac{T}{ST + \frac{(1-S)T}{N}} = \text{speedup}$$

- That is, if application is 75% parallel and 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times

- As N approaches infinity, speedup approaches $1/S$

(infinite core!)

$$\frac{1}{0.25 + \frac{0.75}{2}} \approx 1.6$$

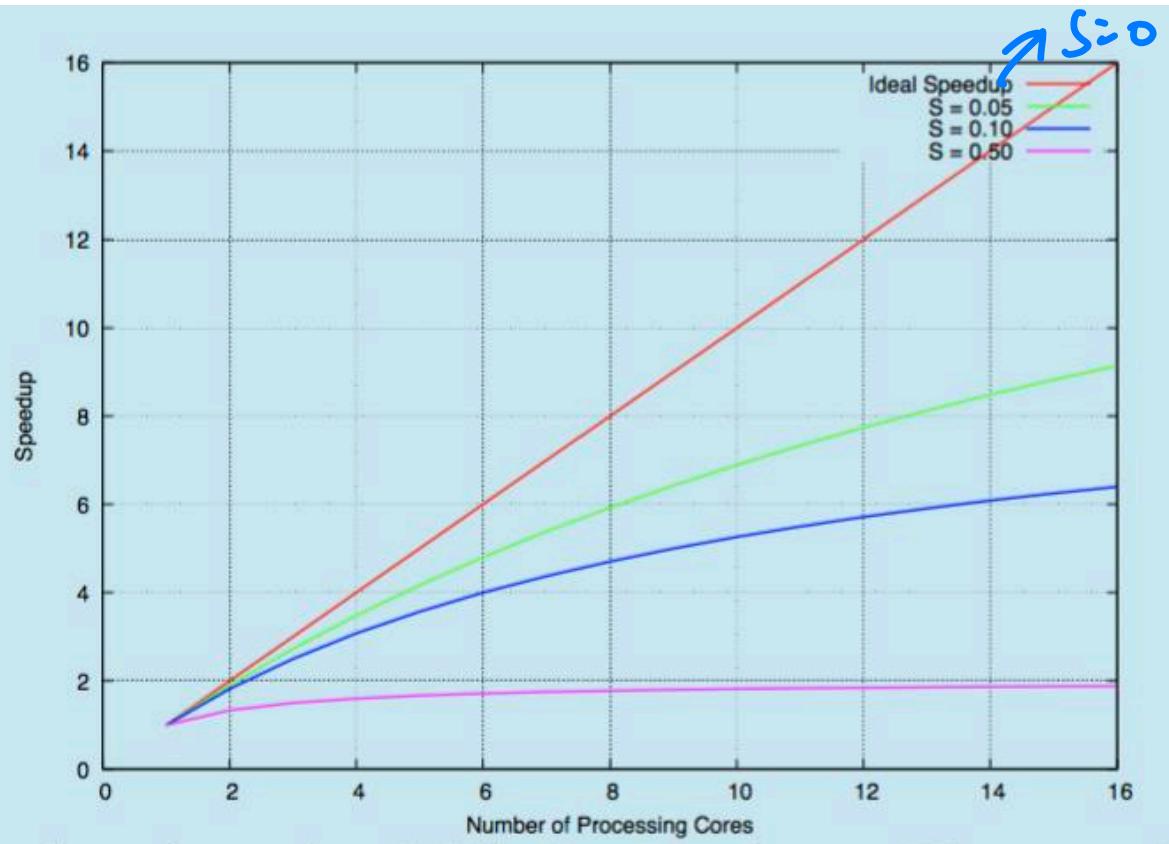
Serial portion of an application has a disproportionate effect on performance gained by adding additional cores

↓↓ → Good speed up!





Amdahl's Law



Large code \rightarrow x parallel \rightarrow very slow \Rightarrow .
? design the code!

Subj as small as possible!





- why not only one process with lots of thread?

Why not multi-process rather than multithread

Multithreaded Process

used for computation

process:

Do isolation

- A multi-threaded process has more than one instance of execution - each of which with a program counter is being fetched and executed
- A thread is similar to a process, except they share the same address space and thus can access the same set of data with other thread(s) in a process
- The state of a single thread is also similar to that of a process - It has a program counter (PC) that tracks where the program is fetching instructions from. Each thread has its own private set of registers for execution
- If two threads run on a single processor, when switching from running one (T1) to running the other (T2), a context switch must also take place
 - The register state of T1 must be saved and the register state of T2 restored from T2's stack before running T2
 - The address space remains the same, the context switch overhead is much smaller
 - When switching to a thread belonging to a different process, overhead is more
- This provides parallelism in a process (multithreaded) execution and can enable overlap of I/O with other activities within a single program
 - One thread is running on CPU, and another thread of the process is doing I/O

不同线程 不同线程 task → task context switch → CPU scheduling



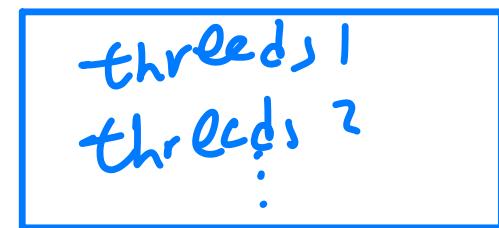


Linux → distinguish between thread Thread

- **Thread**: The single unique execution context - **lightweight process**
 - Program counter, registers, execution flags, stack
 - A thread is executing on a processor when it is resident in the registers.
↳ by control unit!
 - **PC register** holds the address of executing instruction in the thread
 - Registers hold the root state of the thread (other state in memory)
- Each thread has a **Thread Control Block** (TCB) → *like PCB in kernel*
 - **Execution state**: CPU registers, program counter, pointer to stack
 - **Scheduling info**: state (more later), priority, CPU time
 - **Accounting Info** → *Remember who run threads* ↗ *Unique!*
 - **Various Pointers** (for implementing scheduling queues)
↳ for PCB
 - **Pointer to enclosing process**: PCB, which process it belongs to *↳ for corresponding process!*

process →

↳ : solution



→ CPU





sketching ↗ next chapter!

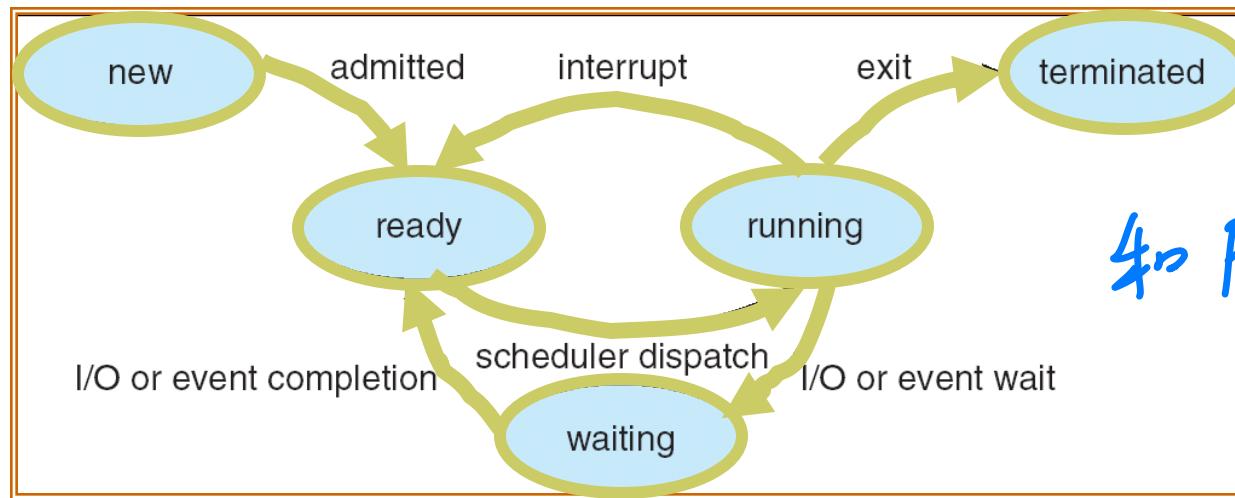
Thread State

- Threads in a way encapsulate concurrency, which is the “active” component of a process
↳ only one process!
- Address space encapsulates protection: which can be considered as the “passive” part of a process
 - One program address space is different from that of another program so to keep buggy program from thrashing the entire system
- State shared by all threads in process/address space
 - Contents of memory (global variables, heap) *↳ 2⁶⁴ states*
 - I/O state (file descriptors, network connections, etc.)
- State “private” to each thread
 - Kept in TCB = Thread Control Block
↳ unique
 - CPU registers (including, program counter)
 - Execution stack (parameters, temporary variables, PC saved)





Lifecycle of a Thread



- As a thread executes, it changes **state**:
 - **new**: The thread is being created
 - **ready**: The thread is waiting to run → by CPU scheduling!
 - **running**: Instructions are being executed
 - **waiting**: Thread waiting for some event to occur
 - **terminated**: The thread has finished execution
- “Active” threads are represented by their TCBs
 - TCBs organized into queues based on their states



Some thread is waiting
→ How to know whether process is running or not?

Process is running

- ≥ 1 thread is running → process is running

Process in ready

- ≥ 1 thread in ready → process is ready
- X threads are running

Process in waiting state

- All threads in process is waiting (to

process in terminated state

All of its threads have finished and are in terminated state

★ context switch of process produces
much more overhead!

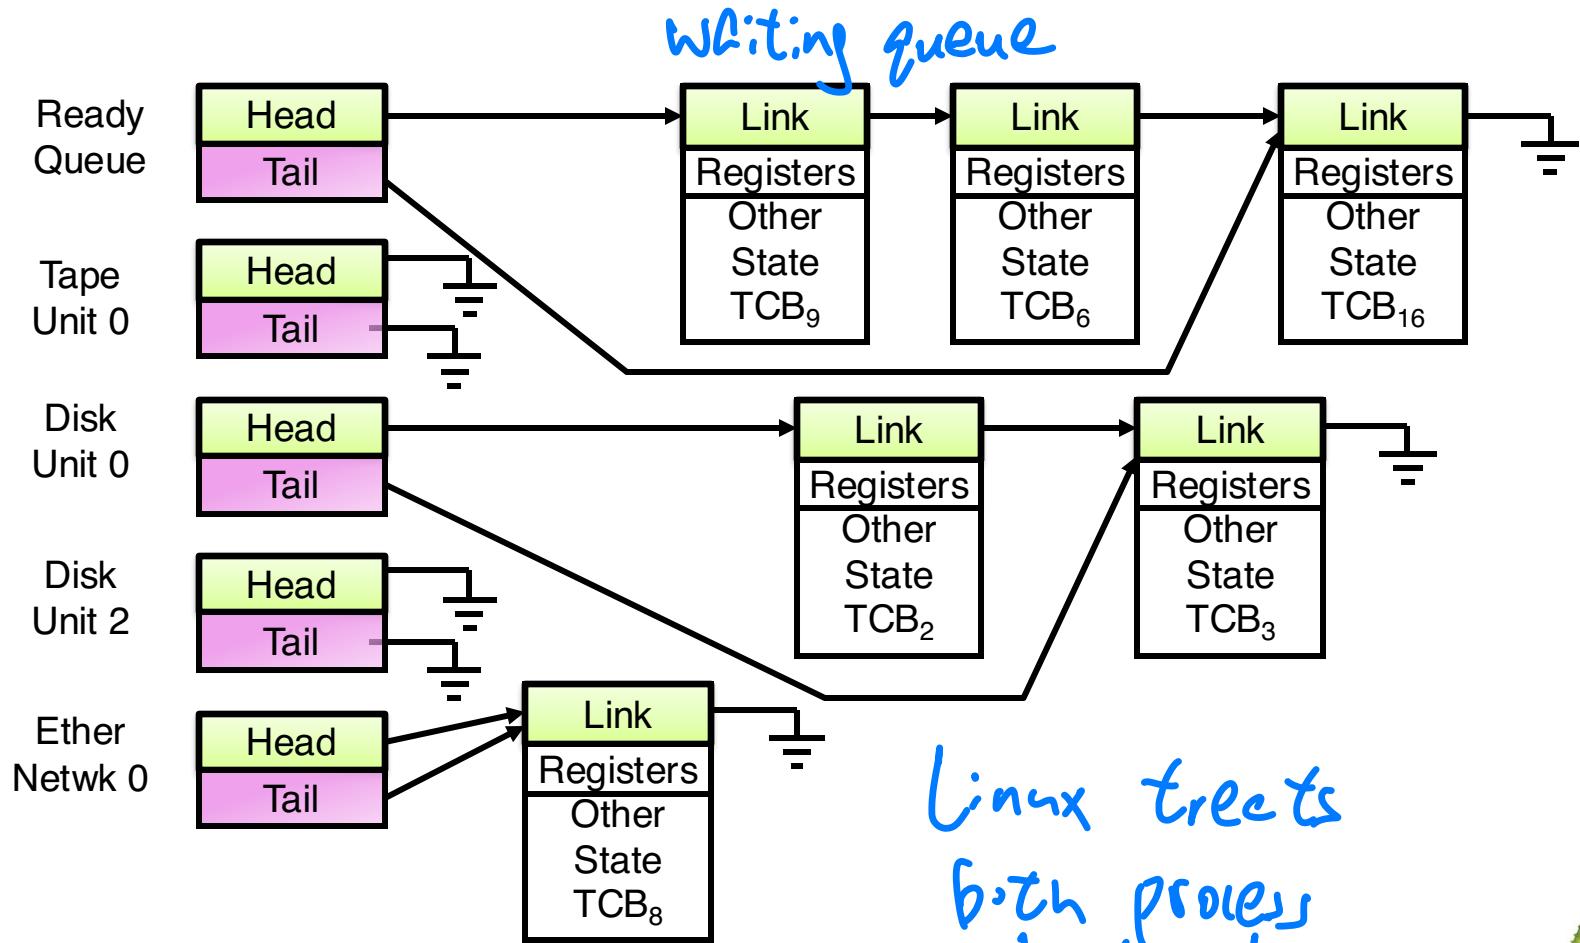
Single thread, only one-to-one mapping
CPU does Scheduling



Similar to PLB! Ready Queue And Various I/O Device Queues

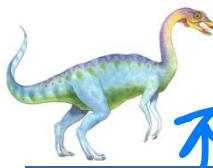
Thread not running \Rightarrow TCB is in some other queue

- Queue for each device/signal/condition, each with its own scheduling policy



Linux treats
both process
and threads as
same thing, called tasks





User Threads and Kernel Threads

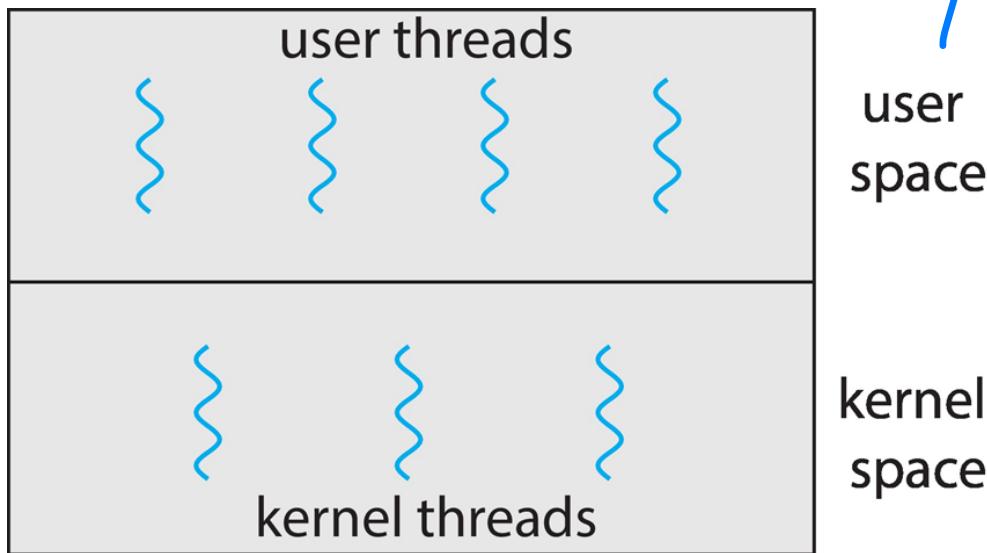
↑ Complicated

OS has own scheduling!

- **User threads** – independently executable entity within a program, created and managed by user-level threads library
 ↳ *Defined by you*
- **Kernel threads** – can be scheduled and execute on a CPU, supported and managed by the operating system
 → *help create threads*
- Examples – virtually all general-purpose operating systems, including:
 - Windows, Linux, Mac OS X iOS, Android

(controlled by OS!)

?.





Multithreading Models

- **User-level threads** are visible only to programmers and unknown to the kernel. Consequently, they cannot be scheduled to run on a CPU. In another word, OS only manages and schedules kernel threads
- **Threads libraries** provide APIs for creating and managing user threads in programs. There are three primary thread libraries - POSIX Pthreads, Windows threads and Java threads, corresponding to the three most common APIs, e.g., Win32 API, POSIX API and Java API *write in program*
- In a way, user-level threads provide the concurrent and modular entities in a program that the OS can potentially take advantage of. Or if there is no user threads defined or specified in a program, there will be no kernel thread(s) that can be scheduled for the program execution

logically, cannot directly execute in CPU

```
pthread_t tid;
```

→ in labs!

```
/* create the thread */
```

```
pthread_create(&tid, 0, worker, NULL);
```

OS only maps user threads !





Multithreading Models (Cont.)

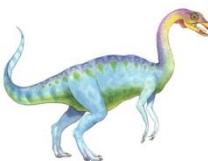
- Ultimately, a mapping must exist between user threads and kernel threads
There are three common ways of establishing such a relationship:
 - **Many-to-One**: many user-level threads mapped to one kernel thread
 - **One-to-One**: each user-level thread mapped to a kernel thread – most common in modern operating system
 - **Many-to-Many**: many user-level threads mapped to many kernel (usually smaller number of) threads



Hw / Exam!

many core,
choose one
Block all
others except
only one!



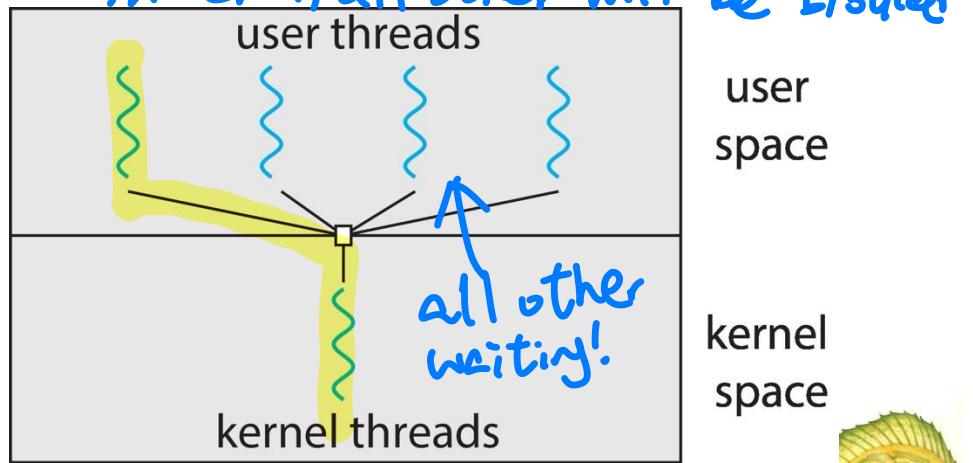


Many-to-One

- Many user-level threads mapped to single kernel thread
- Which user thread is currently mapped to the kernel thread is a **scheduling problem**, known as **process-contention scope** or PCS (to be discussed in Chapter 5)
- **One thread blocking** (i.e., the kernel thread) causes all threads mapped to this thread to block, i.e., the entire process is blocked
⇒ X speed up!
- Multiple user threads **can not run in parallel** on multicore system because only one kernel thread can be active at a time
only one thread can be scheduled in CPU, all other will be blocked
- **Few systems** currently use this model
- Examples:
 - Solaris Green Threads
 - GNU Portable Threads

*No meaningful!
Take multi-threading!*

Separate to other core.





One-to-One

Windows,

use middleware / software

Most straightforward,
easy!

- Each user-level thread maps to one kernel thread
- It provides the **maximum concurrency**, and it also allows multiple threads to run in parallel on multiprocessor or multicore systems
- This implies that creating a user-level thread mandates creating a corresponding kernel thread – the number of threads per process sometimes can be restricted due to overhead (kernel threads consume resources such as memory, I/O, etc.)

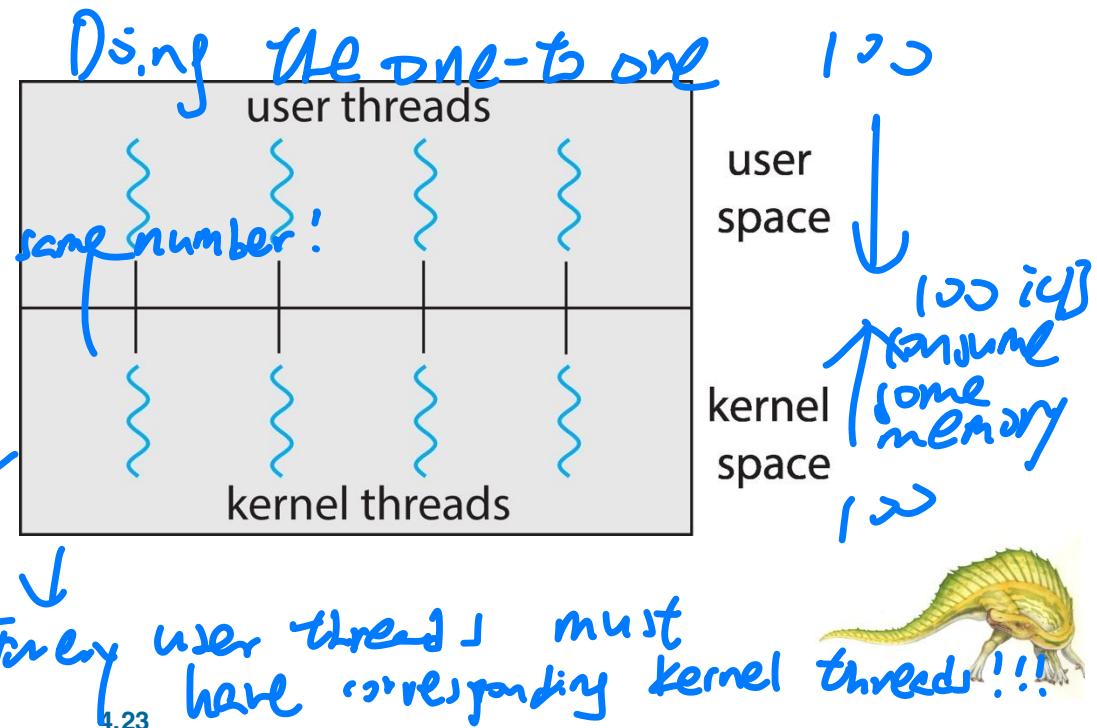
Examples

- Windows
- Linux



Some resource
is wasted!

Schedule
in CPU!





Complicated! Few systems Many-to-Many Model

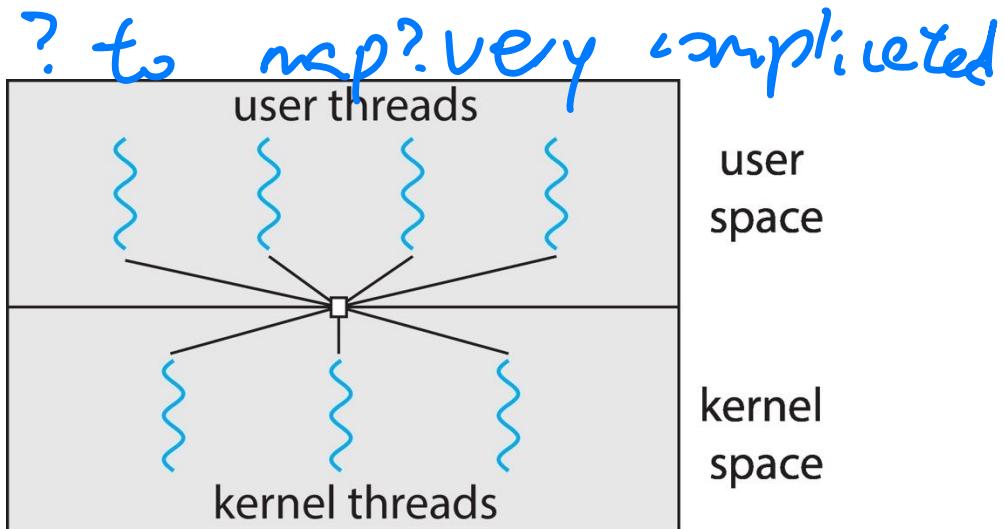
by fiber! → do this

- Multiplexe many user-level threads to a smaller or equal number of kernel threads, specific to a particular application or machine
- This provides certain level of concurrency for a process execution – also **process-contention scope** or **PCS** is involved for scheduling
- This allows the operating system to create a sufficient number of kernel threads in advance – **thread pool**
- Windows with the ThreadFiber package
- Otherwise not very common

Benefit from
multicore structure

X allocate extra
space → reduce
memory!

More flexible!

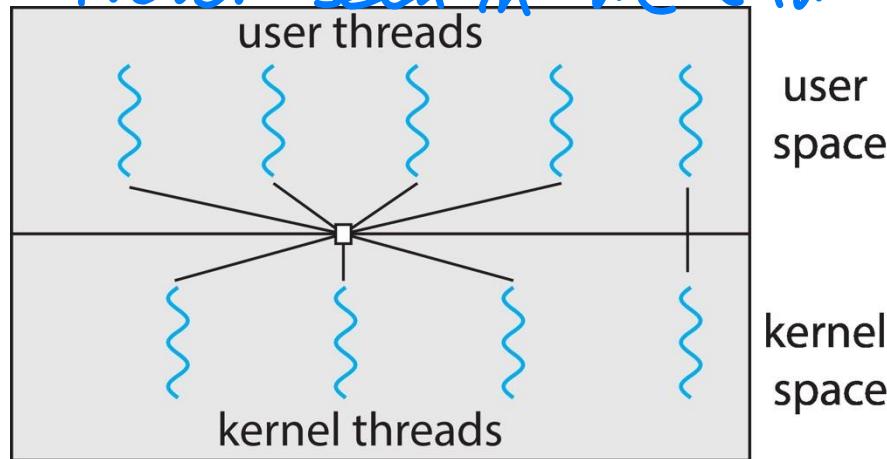




Two-level Model

- Similar to the Many-to-Many model, except that it allows a user thread to be **bound** to a kernel thread

*combine one to one to many to many!
Never seen in the environment*





Multithreading Mapping Models

- The many-to-many model is most flexible, but can be difficult to implement in practice *Few OS uses this!*
 - For instance, an application may be allocated more kernel threads on a system with eight processing cores than on a system with four cores
- With an increasing number of processing cores in modern computer systems, limiting the number of kernel threads has become less important.
- Most operating systems now use the one-to-one model

*resource waste
or be ignored*





Threading Issues

duplicate!

- Semantics of `fork()` and `exec()` system calls
- Signal handling
 - Synchronous and asynchronous
- Thread cancellation of target thread → *also manage!*
 - Asynchronous or deferred cancellation





Semantics of fork() and exec()

- Duplicate all the threads*
- Does `fork()` called by a thread duplicate only the calling thread or all threads of a process?
 - Some UNIX systems have chosen to have two versions of `fork()` - If `exec()` is called immediately after forking, then duplicating all threads is unnecessary, as the program specified in `exec()` will replace the entire process. In this instance, duplicating *only* the calling thread is appropriate.
do not copy everything!
 - `exec()` usually works as normal – replace the entire process including all threads
Only copy the calling thread!
Save some overhead
boot the performance





Signal Handling

low level → cell signal → high level → sent to the process!
? / to trap → generate an interrupt!

- Signals are used in UNIX systems to notify a process that a particular event has occurred. A signal may be received either synchronously or asynchronously, depending on the source and the reason.
- Synchronous signals include illegal memory access and division by 0
 - delivered to the same process that performed the operation that caused the signal (that is why they are considered synchronous)
- When a signal is generated by an event external to a running process – asynchronous signals. Examples include terminating a process with specific keystrokes (such as <control><C>) and a timer expires → terminating the program
- Signals are handled in different ways *又寫, 有 default handler!*
 - Some signals may be ignored, while others (for example, an illegal memory access) are handled by terminating the program

can overwrite
this default!





Signal Handling (Cont.)

multithreaded may be different

- A **signal handler** is used to process signals, following the pattern below
 - Signal is generated by a particular event (e.g., process termination) ✓
 - Signal is delivered to a process ✓
 - Signal is handled by one of two signal handlers, **default** or **user-defined** ✓
- Every signal has **default handler** that kernel uses to handle that signal
 - **User-defined signal handler** can override default handler
- !
□ For a single-threaded process, signals are delivered to that process or thread
Question?
- Where should signal be delivered for a multi-threaded process?
 - Deliver the signal to the thread to which the signal applies
 - Deliver the signal to every thread in the process
 - Deliver the signal to certain threads in the process
 - Assign a specific thread to receive all signals for the process





Signal Handling (Cont.)

16/1 > <

- The method for delivering a signal depends on the type of signal generated
 - For example, **synchronous signals** need to be delivered to the thread causing the signal and not to other threads in the process
 - Some **asynchronous signals** — such as a signal that terminates a process (<control><c>, for example)—should be sent to all threads
- UNIX function for delivering signal `kill(pid t pid, int signal)`, which specifies the process (**pid**) to which a particular signal (**signal**) is to be delivered
- POSIX Pthreads function allows a signal to be delivered to a specified thread (**tid**): `pthread_kill(pthread t tid, int signal)`

*not kill
the process!* *Signal should be directly send
to this thread!*

+ send to particular thread!

Terminate all process → sent to all threads!





Thread Cancellation

how to terminate the thread!

- Terminating a thread before it has finished, and thread to be canceled is referred to as the target thread.
- The cancellation occurs in two different scenarios
 - Asynchronous cancellation terminates the target thread immediately *right now*
 - Deferred cancellation The target thread periodically checks whether it should terminate, allowing it an opportunity to terminate itself in an orderly fashion
- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state and type *Tallocate some memory
tell thread → terminate
delete the unnecessary resources!*
- In Pthreads, if thread has cancellation disabled, cancellation remains pending until thread enables it – the default type is deferred

write done in flag!

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

→ Different Strategy





Thread Cancellation (Cont.)

- ❑ Pthread code to create and cancel a thread:

- ❑ Pthreads: POSIX standard for thread programming
 - ❑ Need to #include <pthread.h>

```
pthread_t tid;
```

```
/* create the thread */
```

```
pthread_create(&tid, 0, worker, NULL);
```

```
. . .
```

Start routine

```
/* cancel the thread */
```

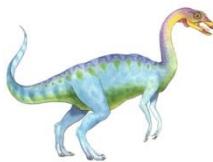
```
pthread_cancel(tid);
```

↑ not immediately
→ make cancellation

```
/* wait for the thread to terminate */
```

```
pthread_join(tid, NULL);
```





Operating System Examples

- Windows Threads
- Linux Threads





Windows Threads

- A Windows application runs as a separate process, and each process may contain one or more threads
- Windows API – primary API for Windows applications
- Windows uses the one-to-one mapping, where each user-level thread maps to an associated kernel thread *(failed!)*
- The general components of a Window thread include:
 - A thread ID identifying the thread
 - Register set representing the state of the processor
 - A program counter (PC)
 - Separate user and kernel stacks when thread runs in user mode or kernel mode
 - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- The register set, stacks, and private storage area are known as the context of the thread





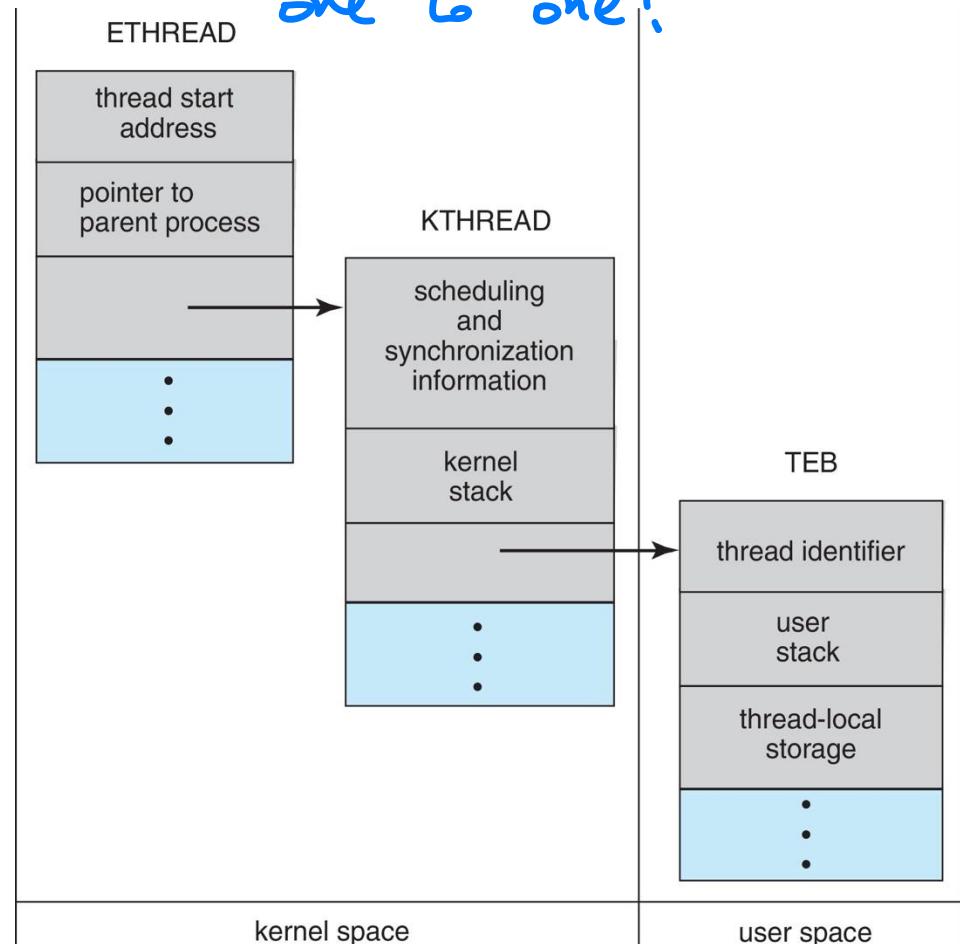
Windows Threads (Cont.)

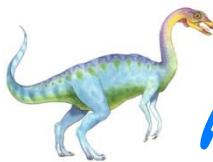
Memory windows, thread/process are different

one to one!

The primary data structures of a thread include:

- ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD,
- KTHREAD (kernel thread block) scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
- TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space





More important! Linux Threads

一样!!!

- Linux does not differentiate between threads and processes. To Linux, a thread is just a special kind of process, all referred to as a **task** structure!
- A thread in Linux is a process that may share certain resources with other threads. Each thread has a unique **task_struct** (illustrated in Chapter 3) and appears as a normal process - threads in Linux just happen to share resources, such as an address space, with other processes
- This approach to threads contrasts greatly with operating systems such as Microsoft Windows or Sun Solaris (Unix), which have explicit kernel support for threads (and sometimes call threads *lightweight processes*)
- For example, if a process consists of two threads

On Microsoft Windows, one PCB exists describing the shared resources such as address space or open files, and that, in turn, points to the two different threads. Each thread TCB describes the resources it alone possesses

- In Linux, there are simply two tasks with normal **task_struct** structures.

The two processes are set up to share resources

process → leading task → link list!
share resources
 $PLD = \text{thread ID} - \text{task ID}$ → `clone()`
whether two tasks to share something



Linux Threads

- Threads are created in the same way as normal tasks, with the exception that the **clone()** system call is used to pass flags corresponding to the specific resources to be shared:

```
clone(CLONE_VM | CLONE_FS | CLONE_FILES | CLONE_SIGHAND, 0);
```

- clone()** allows a child task to share the part of the execution context with a parent task, such as address space, file descriptors, signal handler
- Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

fork()

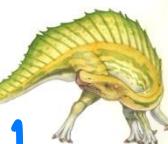


*clone system
do not share
everything*

- Instead of copying all data structures like **fork()**, new task points to data structures of parent task, depending on set of flags passed to clone()
- A normal **fork()** can be implemented as **clone(SIGCHLD, 0)**;

*(control two
execution instance!)*

*create between
process and thread*



End of Chapter 4

