

very important!!!

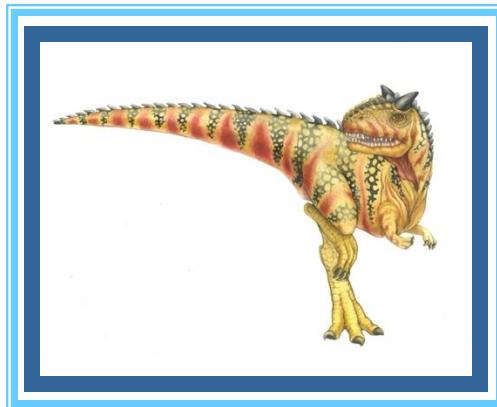
Short / Medium / Long term scheduling!

- Affect the performance

⇒ AW2

Chapter 5: CPU Scheduling

Want the utilization percentage $\rightarrow 1$





Chapter 5: CPU Scheduling

- Basic Concepts
- Scheduling Criteria
- Scheduling Algorithms
- Thread Scheduling
- Multi-Processor Scheduling
- Real-Time Scheduling
- Algorithm Evaluation

Draw diagrams!
FIFO/Ringing/Multi-level

— Simulation!





Objectives

Advantages / Disadvantages

- Describe various CPU scheduling algorithms
- Evaluate CPU scheduling algorithms based on scheduling criteria
- Explain the issues related to multiprocessor and multicore scheduling
- Describe real-time scheduling algorithms
- Apply modeling and simulations to evaluate CPU scheduling algorithms





Basic Concepts

holes ↗

DMA →

SC; trf I/O

Make
as busy
as possible
•
run it

- The objective of multiprogramming is to try to have a process running at all times - **maximize CPU utilization**

- Process execution consists of **cycles** of **CPU execution** and **I/O wait** – referred as **CPU burst** and **I/O burst** (when not running on CPU)

- Whenever CPU is idle, OS tries to select one of processes on the ready queue to execute unless the ready queue is empty, i.e., no process is in ready state

- The selection of process is done by the **CPU scheduler** or also called **process scheduler, short-term** → in ch 3 scheduler

load store
add store
read from file

wait for I/O

store increment
index
write to file

wait for I/O

load store
add store
read from file

wait for I/O

CPU burst

I/O burst

CPU burst

I/O burst

CPU burst

I/O burst

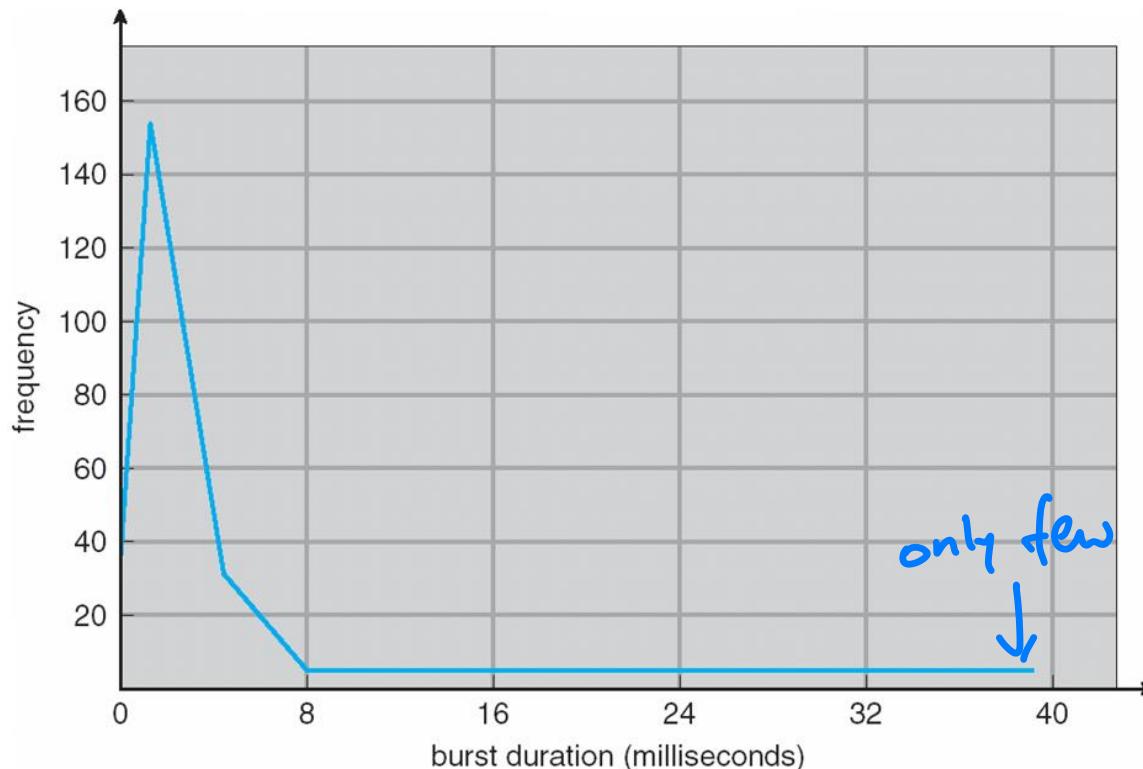




Histogram of CPU-burst Times

long tail distribution (short burst ↑↑)

- The durations of CPU bursts have been measured extensively over the years. The frequency curve is similar to that shown below
- There are many short CPU bursts and a small number of long CPU bursts (long-tail distribution). An I/O-bound program typically has many short CPU bursts, while a CPU-bound program might have a few long CPU bursts.
- This distribution is important for designing a CPU-scheduling algorithm



? Naive algorithm
X works?

Heavy tasks!





In first part of this chapter!

CPU Scheduler

Process = Process with one thread thread: unit of execution

- The CPU scheduler selects one process or thread from processes on ready queue, and allocates the CPU core to the selected process
 - Queue may be ordered in various ways, single or multiple queues
- CPU scheduling may take place in the **following four conditions**:
 1. Switches from running to waiting state, e.g., I/O request, or wait()
 2. Process terminates
 3. Switches from running to ready state, e.g., interrupt
 4. Switches from waiting to ready, e.g., completion of I/O or from new to ready, a new process arrives on the ready queue with a higher priority
- Scheduling under 1 and 2 is **non-preemptive**, in which a process gives up CPU voluntarily by itself
 - ↑ Actively do something → can pick
- Scheduling under 3 and 4 is **preemptive** *someone kick u out* *完Running!*
 - Consider access to shared data (discussed in Chapter 6)
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities



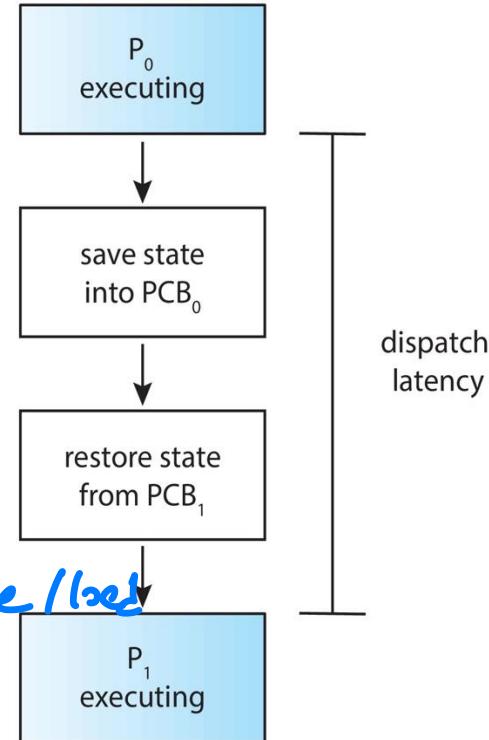


Dispatcher

Context switch!

module in CPU

- **Dispatcher** module allocates the CPU to the process selected by CPU scheduler; this involves:
 - switching context from one process to another
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- **Dispatch latency** – the time it takes for the dispatcher to stop one process and start another running – an overhead
 - time to store / load the status
- The number of context switches can be obtained by using `vmstat` on Linux, typically hundreds of context switches per second





? What is yr requirement?

Scheduling Criteria

keep high Always

- **CPU utilization** – fraction of the time that CPU is busy → want 100%
every useful, always have process!
- **Throughput** – # of processes or jobs completed per time unit
 - Efficient use of the resources (CPU, memory, disk, etc.), favor short processes
- **Waiting time** – amount of time a process waiting on the ready queue
 $T_1 = \text{submitted time}$, $T_2 = \text{job finished time}$ $\rightarrow T_2 - T_1$
Easy to calculate, preemptive by others → order changes → first
- **Turnaround time** – the amount of time to execute a particular process, measured by CPU burst time, I/O burst time and waiting time → finish the whole job!
 - With a single CPU burst, the turnaround time = waiting time + CPU burst time
 - Waiting time = turnaround time – CPU burst time = the completion of the CPU burst time – the arrival time (to the ready queue) – CPU burst time
- **Response time** – the amount of time it takes from when a request is submitted until the first response is produced – RR scheduling
More important in interactive → core whole program just first response.
 - Time to echo a keystroke in editor, or games
 - This is more relevant to interactive programs (typically using RR scheduling)
 - Considering single CPU burst, this is the time between the completion of first CPU time minus the time that this process joins the ready queue
- **Fairness** – adequate allocation for different jobs or/and processes

↑ don't care! Maxing (very complicated)





Scheduling Criteria (Cont.)

depends on requirement!

- It is desirable to maximize CPU utilization and throughput and to minimize turnaround time, waiting time, and response time. But these can be **conflicting** set of criteria, there are different considerations in practice
- In most cases, we optimize an **average measure**, e.g., the average waiting time or average turn-around time. However, under special circumstances, we prefer to optimize the minimum or maximum values rather than the average
 - Considering all users, we may want to minimize the **maximum response time**
- For interactive systems (such as a desktop or laptop), it might be more important to minimize the variance in the response time than to minimize the **average response time**
 - A system with reasonable and predictable response time may be considered more desirable than a system that is faster on the average but is highly variable
- Different CPU-scheduling algorithms have different properties. we next describe several scheduling algorithms in the context of **only one CPU core** – the system is capable of only running one process or thread at a time





Algorithms!

First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

Just Q queue

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The Gantt Chart for the schedule is:

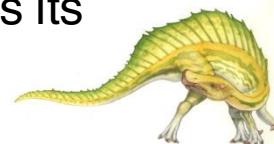


- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$
- Average turn-around time: $(24+27+30)/3 = 27$
- In earlier systems, FCFS means that one program is scheduled to run until completion including all I/O
- In multiprogramming systems, this usually means a process finishes its current CPU burst time

But depends on process!
You never know
what's the burst time!



- Just know the arrival order → can use!





FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for P₁ = 6; P₂ = 0, P₃ = 3
- Average waiting time: $(6 + 0 + 3)/3 = 3$ (much shorter!)
- Average turn-around time: $(30+3+6)/3 = 13$
- Convoys effect?** A short process stuck behind a long process, bad for short jobs, depending purely on the arrival order
↑↑ time ?.
↳ long time process block others!
- Consider one CPU-bound and many I/O-bound processes, FCFS results also in low I/O device utilization → long time process block others!
- Waiting in banks: depositing a check, stuck behind new account opening
 - ↑ waiting time / turnaround time



only know the arrival order



FCFS

long tail
distribution

→ Convoy Effect

Makes it worse

Reduce time of long process

kick them off if too long (Allow short)

Round Robin



Round Robin (RR)

- The FCFS scheduling algorithm is non-preemptive. Once the CPU core allocated to a process, the process keeps the CPU until it releases CPU
Want more responsible OS!
- The FCFS is thus particularly troublesome for interactive systems, where it is important that each process get a share of the CPU at regular intervals.
- The **round-robin (RR)** scheduling algorithm is similar to FCFS scheduling, but preemption is added to enable system to switch between processes
difficult to design
- Each process gets a small unit of CPU time (time quantum q), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue (exactly like FCFS). *加入就绪队列*
- Given n processes, each process gets $1/n$ of the CPU time in chunks of at most q time units at once – context switching time is ignored
 - No process waits more than $(n-1)q$ time units.
 - A **timer** interrupts every quantum to schedule next process, or the process blocks upon completing its current CPU burst time when its CPU burst time or remaining CPU burst time $< q$





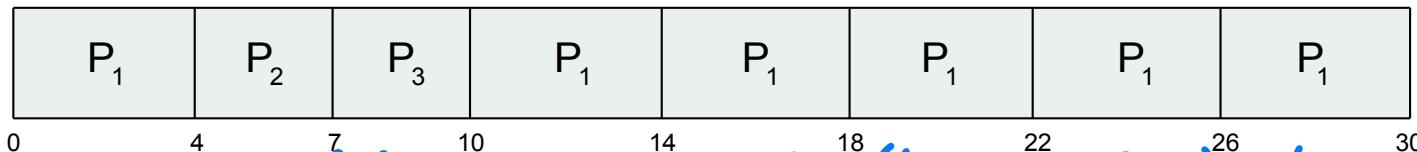
Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

$Q = 4 \text{ times}$.

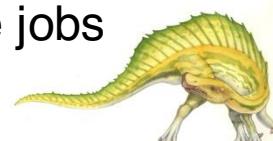
- The Gantt chart is:

(Calculate ALL waiting time !!!)



never response more than $(3-1) \times 4 \text{ time}$.
Waiting time for $P_1=6, P_2=4, P_3=7$

- Average waiting time $(6+4+7)/3 = 5.67$ (better!) \rightarrow ~~bad~~ \downarrow
- Average turn-around time: $(30+7+10)/3 = 15.67$ *First CPU burst - Survival time!!!* \uparrow very good reduce
- Response time** for $P_1=4, P_2=7, P_3=10$, average = $(4+7+10)/3 = 7$
- The average waiting time under RR policy can be long, but is inherently more "fair" (FIFO order), usually perform better for short jobs than FCFS, and offers better average response time – important for interactive jobs

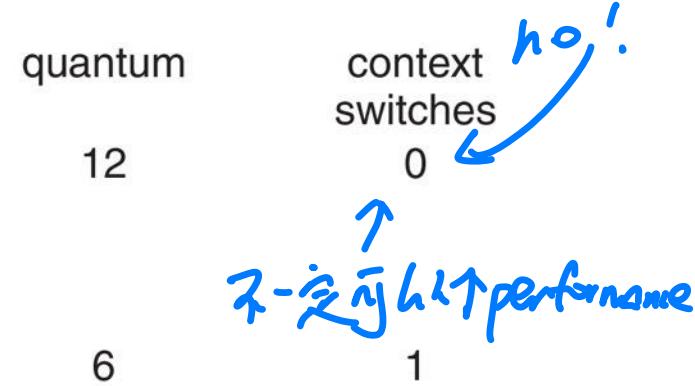
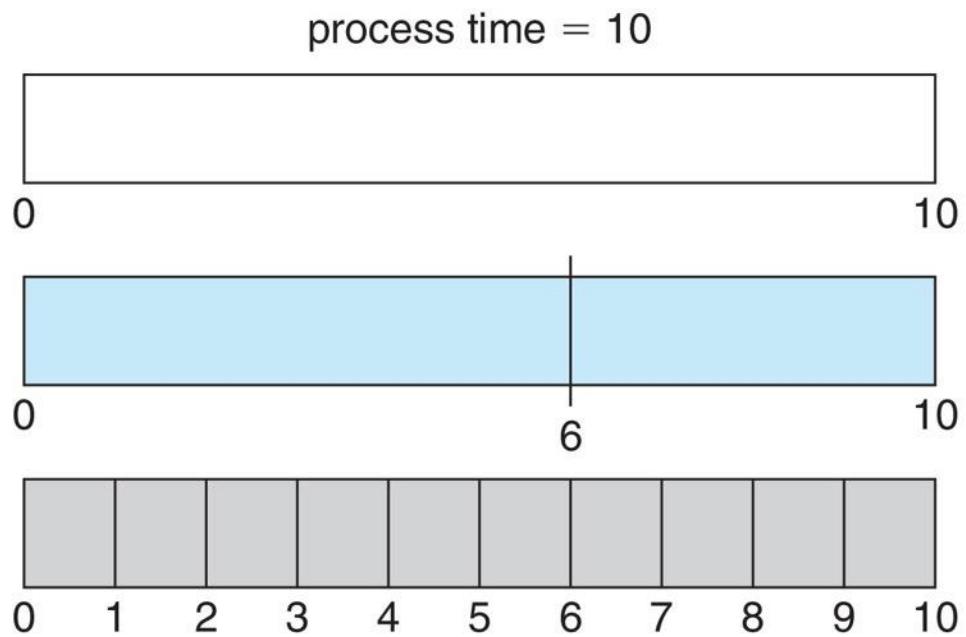




Time Quantum and Context Switch Time

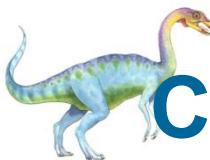
How to set τ ?

- The performance of the RR algorithm depends heavily on the size of the time quantum
 - τ large \Rightarrow FCFS \rightarrow meaningless!
 - τ small \Rightarrow interleaved, but τ must be large with respect to context switch time (usually $< 10 \mu\text{sec}$), otherwise overhead is too high



Ignore the context switch overhead





Comparisons between FCFS and RR

Assuming zero-cost context-switching time, is RR always better than FCFS?

- An example: 10 jobs starting at the same time, each taking 100s of CPU time; RR scheduler quantum of 1s;

Job #	FIFO	RR
1	100	991
2	200	992
...
9	900	999
10	1000	1000

不是 long tail distribution!

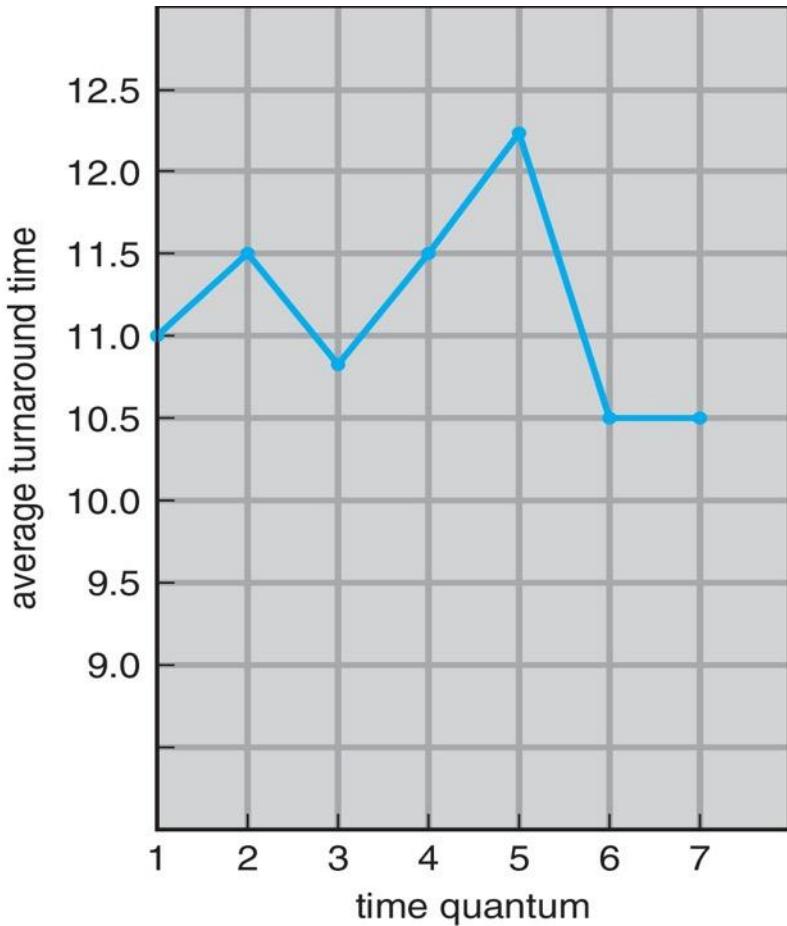
- The average job turn-around time is much worse under RR!

- ▶ Bad when all jobs have the same length *No-one can finish earlier!*
- ▶ The average response time is much better





Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

long-tail distribution

Want short job finish first!

- The average turnaround time does not necessarily improve when the quantum increases
- In general, the average turnaround time can be improved if most processes finish their current CPU bursts within a single quantum
- The time quantum can not be too big, in which RR degenerates to an FCFS policy
- **A rule of thumb:** 80% CPU bursts should be shorter than the time quantum q

RR \rightarrow good response time!

Rule!





Shortest-Job-First (SJF) Scheduling

If we know the order and the burst time

- Noticing that in FCFS and RR, we do not need to know the next CPU burst time of each process during scheduling, and scheduling is done solely based on the arrival order to the ready queue
- What if we knew the future – **the next CPU burst time of each process**
only know the arrival time
- Associate with each process the length of its next CPU burst
 - To schedule the process with the shortest next CPU burst
- The **Shortest Job First** or **SJF** scheduling algorithm is optimal – it produces the minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request *tut1* ?
 - The basic idea is to get the short jobs out of the system sooner
 - Big effect on short jobs, relatively small effect on long jobs
 - This can be applied to an entire program, or the current CPU burst
 - Perhaps a more precise term should be the **shortest-next-CPU-burst** algorithm, but shortest job first or SJF is commonly used.



SJF
↓ preemptive
SRJF

know the order and the burst time!

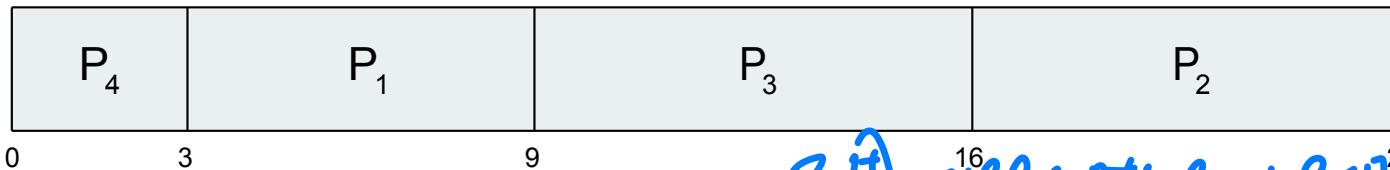


Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- SJF scheduling chart

we know $P_4 < P_1 < P_3 < P_2$



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$ Optimal scheduling time!
- The “best” FCFS perform the same if arrival order happens to be the same
- **Sketch Proof:** Moving a short process before a long one decreases the waiting time of the short process more than it increases the waiting time of the long process. Consequently, the average waiting time decreases





Determining Length of Next CPU Burst

- How to estimate the length based on the past behavior
 - Then pick the process with shortest predicted next CPU burst
 - *Predict burst time!*

但 practical
75%

- Can be done by using the lengths of previous CPU bursts and exponential averaging algorithm

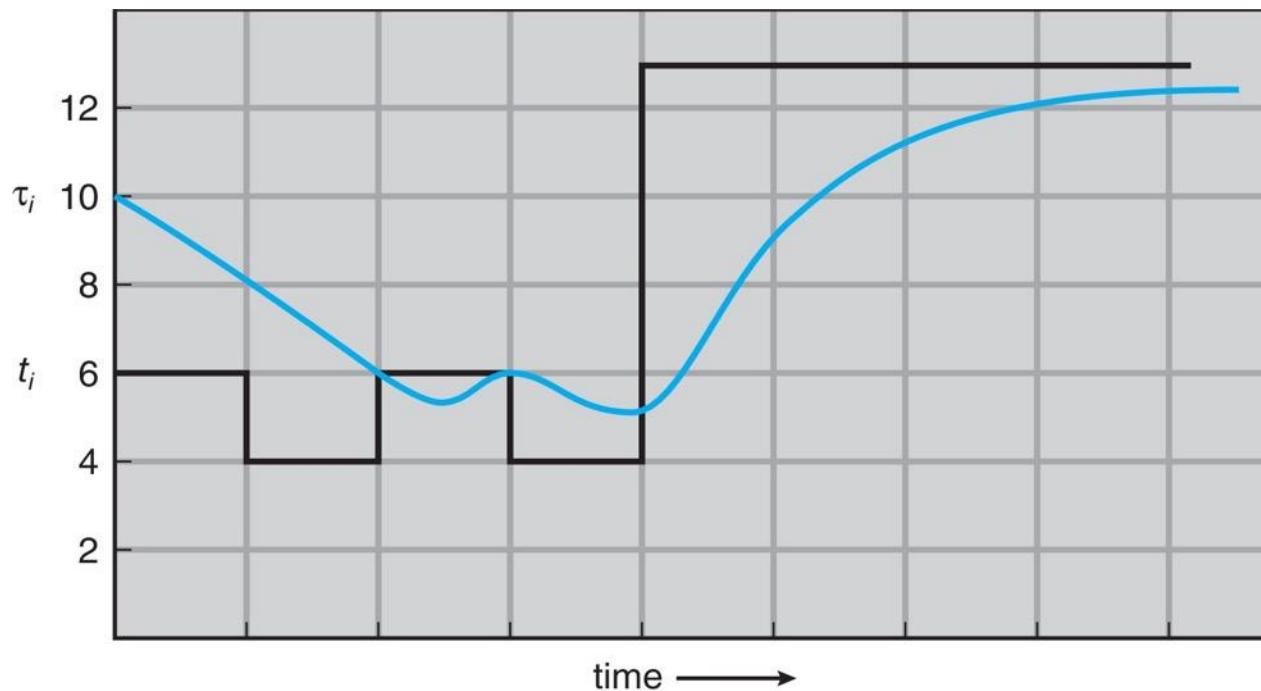
1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.

- Commonly, α set to $1/2$ - the relative weight of recent and past history in the prediction
- Preemptive version called **shortest-remaining-time-first (SRTF)**

use multi-level feedback queue



Prediction of the Length of the Next CPU Burst



CPU burst (t_i)	6	4	6	6	4	13	13	13	...
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

This shows an exponential average with $\alpha = 1/2$ and $\tau_0 = 10$.





Exponential Averaging Algorithm

- This is commonly used to estimate a future value based on historical data – the recent historical value(s) weight more than the past historical value(s)

1. t_n = actual length of nth CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. α , $0 \leq \alpha \leq 1$
4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$

If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} &= \alpha t_n + (1 - \alpha)\tau_n \quad \text{Downward arrow} \\ &= \alpha t_n + (1 - \alpha) \cdot (\alpha t_{n-1} + (1 - \alpha)\tau_{n-1}) \\ &= \alpha t_n + (1 - \alpha)\alpha t_{n-1} + (1 - \alpha)^2\tau_{n-1} \\ &= \alpha t_n + (1 - \alpha)\alpha t_{n-1} + (1 - \alpha)^2 \cdot (\alpha t_{n-2} + (1 - \alpha)\tau_{n-2}) \\ &= \alpha t_n + (1 - \alpha)\alpha t_{n-1} + (1 - \alpha)^2\alpha t_{n-2} + (1 - \alpha)^3\tau_{n-2} \\ &= \dots \\ &= \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots + (1 - \alpha)^j\alpha t_{n-j} + \dots + (1 - \alpha)^{n+1}\tau_0\end{aligned}$$

exp. term
Downward arrow





Examples of Exponential Averaging

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count

- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts

- If we expand the formula, we get: *it contains ALL historical data!*

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$

More recent value → weight more

- All past values (i.e., the length of the CPU bursts) matter, since they ($t_n, t_{n-1}, t_{n-2}, \dots$) are included in the formula
- But recent values matter more. The past value effect is diminished **exponentially fast**, with each time multiple by a fraction $(1 - \alpha)$.
- For example, if $\alpha = 0.5$, the effect after 9 rounds, $(1 - \alpha)^{9+1} = 1/1024$, which can be ignored.

→ it's accurately predict value!



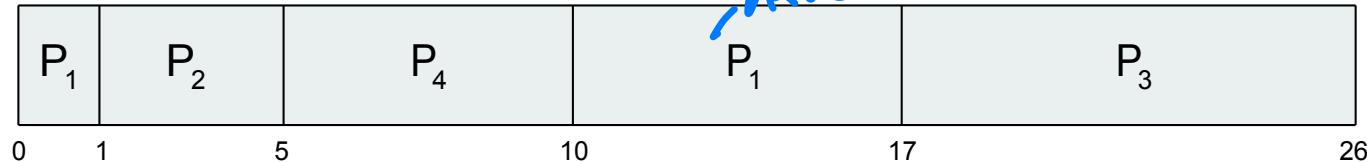


Example of Shortest-Remaining-Time-First

- The SJF algorithm can be either **preemptive** or **nonpreemptive**. The choice arises when a new process arrives at the ready queue while another process is still executing

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>	<i>SJF → optimal!</i>
P_1	0	8	
P_2	1	4	
P_3	2	9	
P_4	3	5	

- Preemptive SJF Gantt Chart



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$
- Now scheduling needs to be invoked when there is an arrival to the ready queue (scheduling condition 4)





Comparison of SJF/SRTF and FCFS

- SJF/SRTF are the best we can do towards minimizing the average waiting time. or minimizing the average turnaround time

- Provably optimal (SJF among non-preemptive, SRTF among preemptive)

better

- SRTF is always at least as good as SJF

- SJF/SRTF performs the same as FCFS if all processes have the same CPU burst times

starvation → long term must be in the last!

- SJF/SRTF can possibly lead to starvation for long process if there is always shorter process joining the ready queue

- “fairness” can not be enforced

X interactive OS





Priority Scheduling

- A priority number (e.g., integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer = highest priority), it can be
 - Preemptive (upon new arrival of a higher priority process) ✓
 - Nonpreemptive ✓
- Equal-priority processes are scheduled in FCFS order
- SJF is a special case of the general priority-scheduling algorithm, where priority is the inverse of predicted next CPU burst time
↳ may be long process!
- Problem = **Starvation** – low priority processes may never execute → *unfair*
- Solution = **Aging** – as time progresses increase the priority of the process





Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>	<u>Difficult to set priority in real world!</u>
P_1	10	3	
P_2	1	1	
P_3	2	4	
P_4	1	5	
P_5	5	2	

- Priority scheduling Gantt Chart



- Average waiting time = 8.2

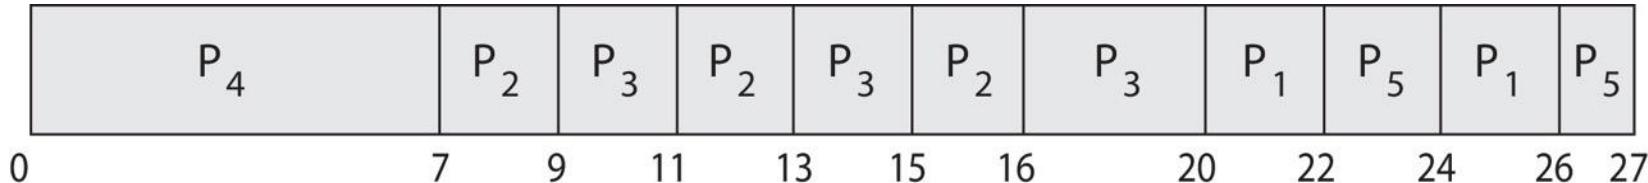




Priority Scheduling w/ Round-Robin

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>	
P_1	4	3	
P_2	5	2	Priority Scheduling!!!
P_3	8	2	RR only for some priority!
P_4	7	1	
P_5	3	3	

- Run the process with the highest priority. Processes with the same priority run round-robin
RR → make more fair!
- Gantt Chart with 2 ms time quantum



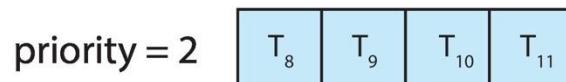
First consider Priority → then RR !!!





Multilevel Queue

- Multilevel queue scheduling can still be a priority scheduling combined with round-robin
- A priority is assigned statically to each process, and a process remains in the same queue for the duration of its runtime



•
•
•



↑ Priority, ↑ queues!!!

↓ CPU scheduling first

Each queue can
use different
scheduling algorithm!!!
- usually RR



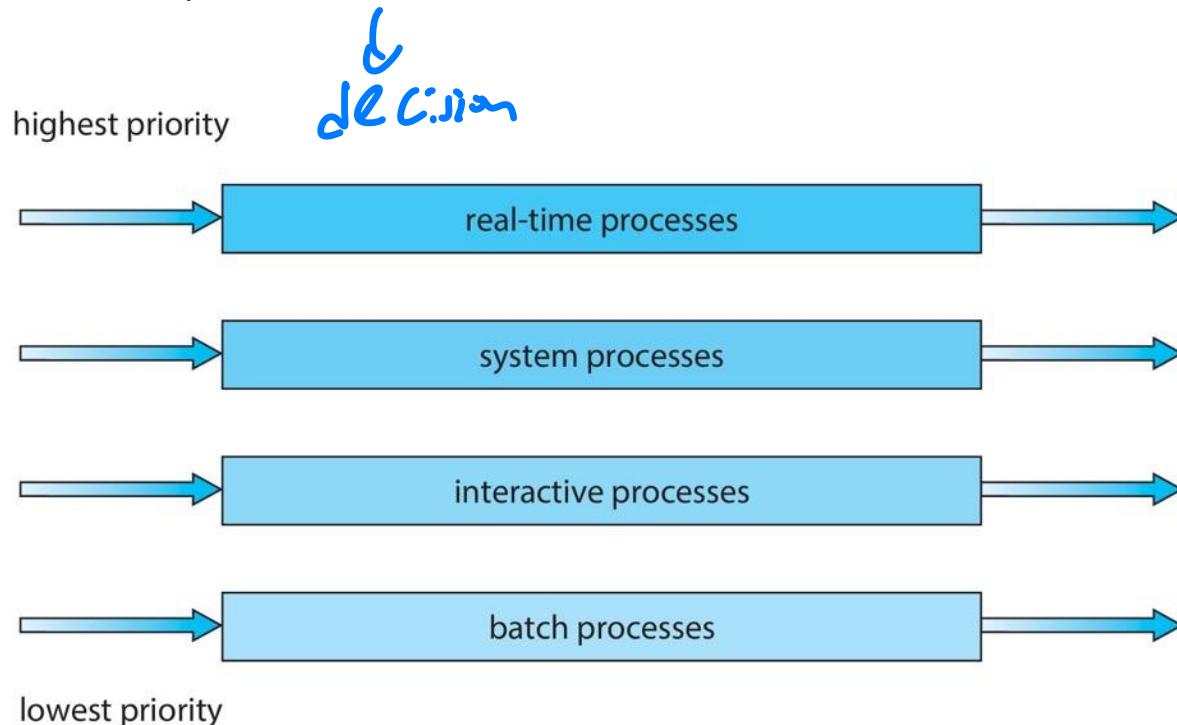


Not important!

Multilevel Queue (Cont.)

Partition processes into different queues based on process type

- Each queue can have its own scheduling algorithm based on the needs
- The scheduling among the queues, is commonly implemented as **fixed-priority preemptive scheduling** or each queue gets certain amount of CPU time – time-slice (for instance 60%, 20%, 10%, 10%)





Multilevel Feedback Queue (MLFQ)

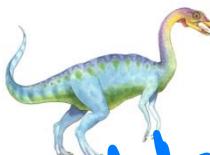
- A process can move between the various queues; aging can be implemented this way. This provides the flexibility
- Multilevel-feedback-queue or MLFQ scheduler defined by the following parameters:
 - The number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service

↑ 用太多 → aging problem
↑ timer!

→ are !

process can move between various queue
solve for aging problem, move to 1 priority queue





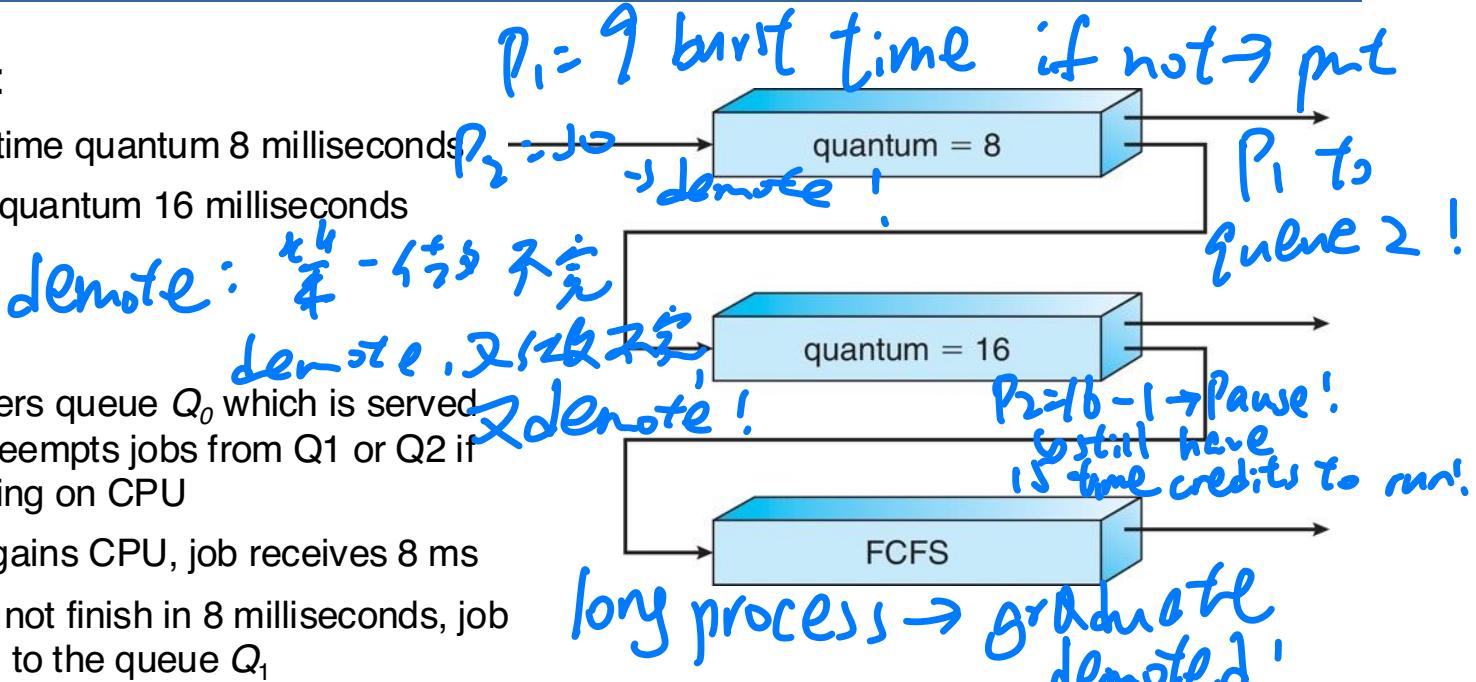
Example of Multilevel Feedback Queue

~~AKA~~ Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS

Scheduling

- A new job enters queue Q_0 , which is served FCFS, also preempts jobs from Q_1 or Q_2 if currently running on CPU
 - When it gains CPU, job receives 8 ms
 - If it does not finish in 8 milliseconds, job is moved to the queue Q_1
- At Q_1 , job is again served FCFS and receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2
- If a job from Q_1 or Q_2 is preempted by a new job from Q_0 , it joins the head of the queue Q_1 or Q_2 , respectively



□ This approximates SRTF:

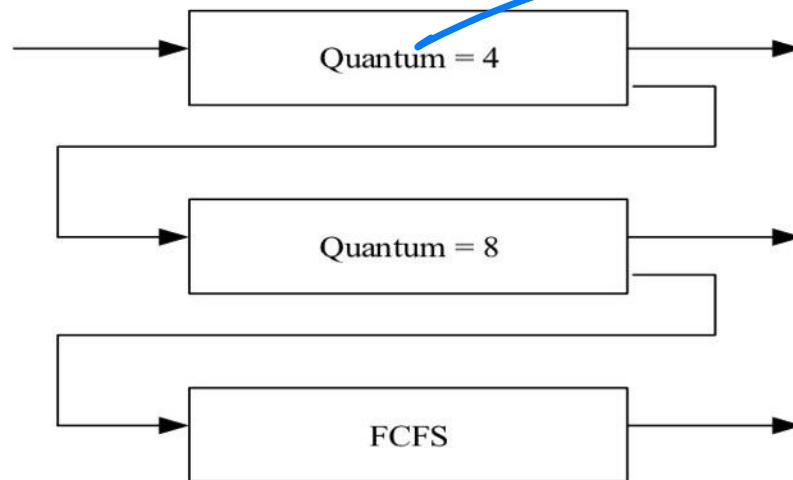
- CPU bound jobs drop like a rock
- Short-running I/O bound jobs stay near top



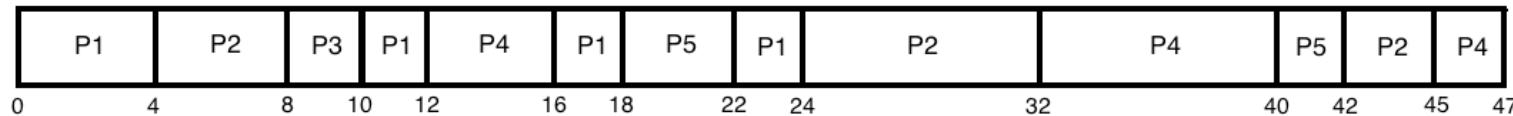


MLFQ Example

time credits



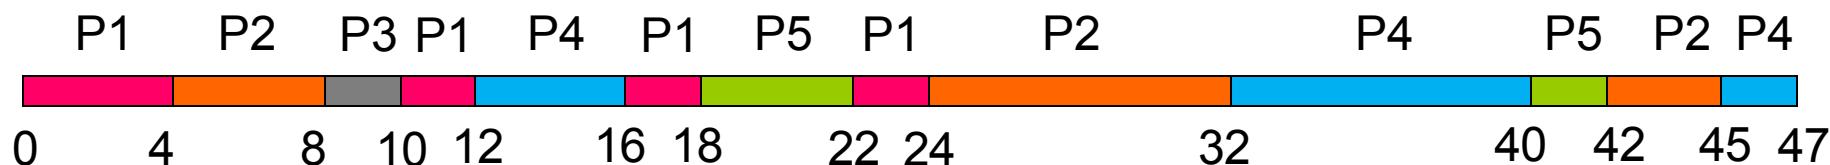
Process	Arrival Time (ms)	Burst Time (ms)
P1	0	10
P2	2	15
P3	5	2
P4	12	14
P5	18	6



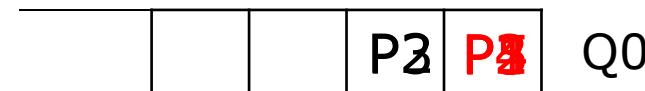


MLFQ Scheduling: Example

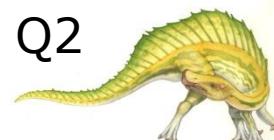
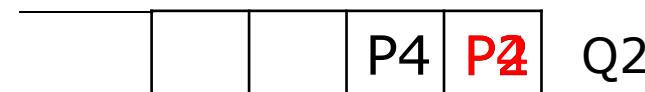
Process	Burst Time	Arrival Time	Remaining Time
P1	10	0	10
P2	15	2	15
P3	2	5	0
P4	14	12	14
P5	6	18	6



0 4 8 10 12 16 18 22 24 32 40 42 45 47



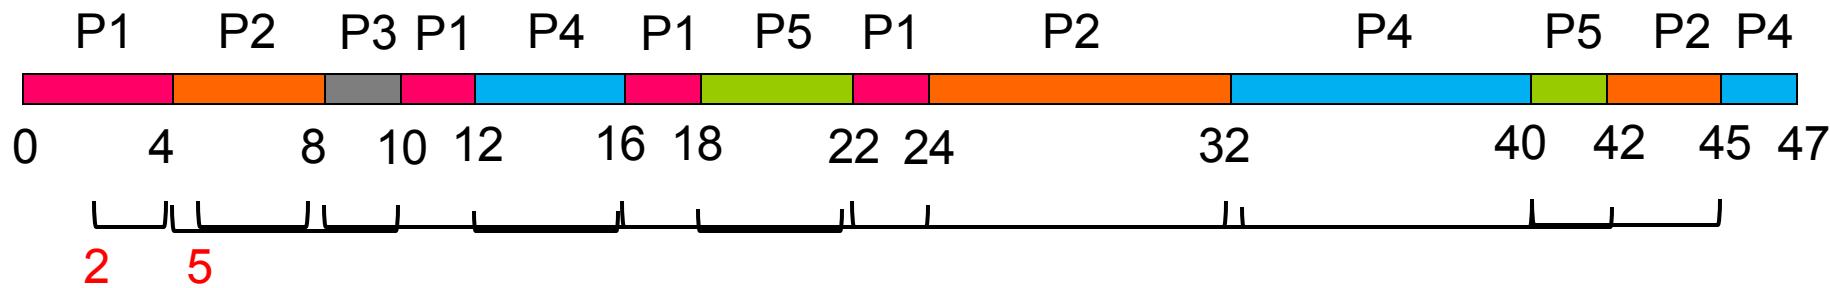
- at time 24 P3 gets service, P4 waits
and waits in Q1; P1 gets service in Q0





MLFQ Scheduling: Example

Process	Burst Time	Arrival Time	Remaining Time
P1	10	0	0
P2	15	2	0
P3	2	5	0
P4	14	12	0
P5	6	18	0



- Waiting time for P1=64, P2=28, P3=5, P4=0, P5=18
- Average waiting time: $(14+28+3+21+18)/5=16.8$





Multilevel Feedback Queue (MLFQ)

- MLFQ is commonly used in many systems such as BSD Unix, Solaris, Window NT and subsequent Window operating systems
- MLFQ has several distinctive advantages: *Approximate!*
 - It does not need prior knowledge of the next CPU burst time *long process*
 - It handles interactive jobs well by delivering better performance in term of response time than RR *gradually demoted!*
 - It produces similar performance as that of SJF or SRTF
 - It is also “fair” by making *demoted gradually* progress on CPU-bound jobs
- The possible starvation problem can be handled by reshuffling the jobs to different queues periodically
 - E.g., after some period, move all jobs to the top queue



- Targets
 - ① ↑ CPU utilization
 - ② ↓ average waiting time / response time / time

SJF → non-preemptive

SRJT → preemptive

Not practical → hard to know next CPU burst time

- exponential averaging

- MIQ , window / Unix , popular in modern OS.

Demoted / Upgraded

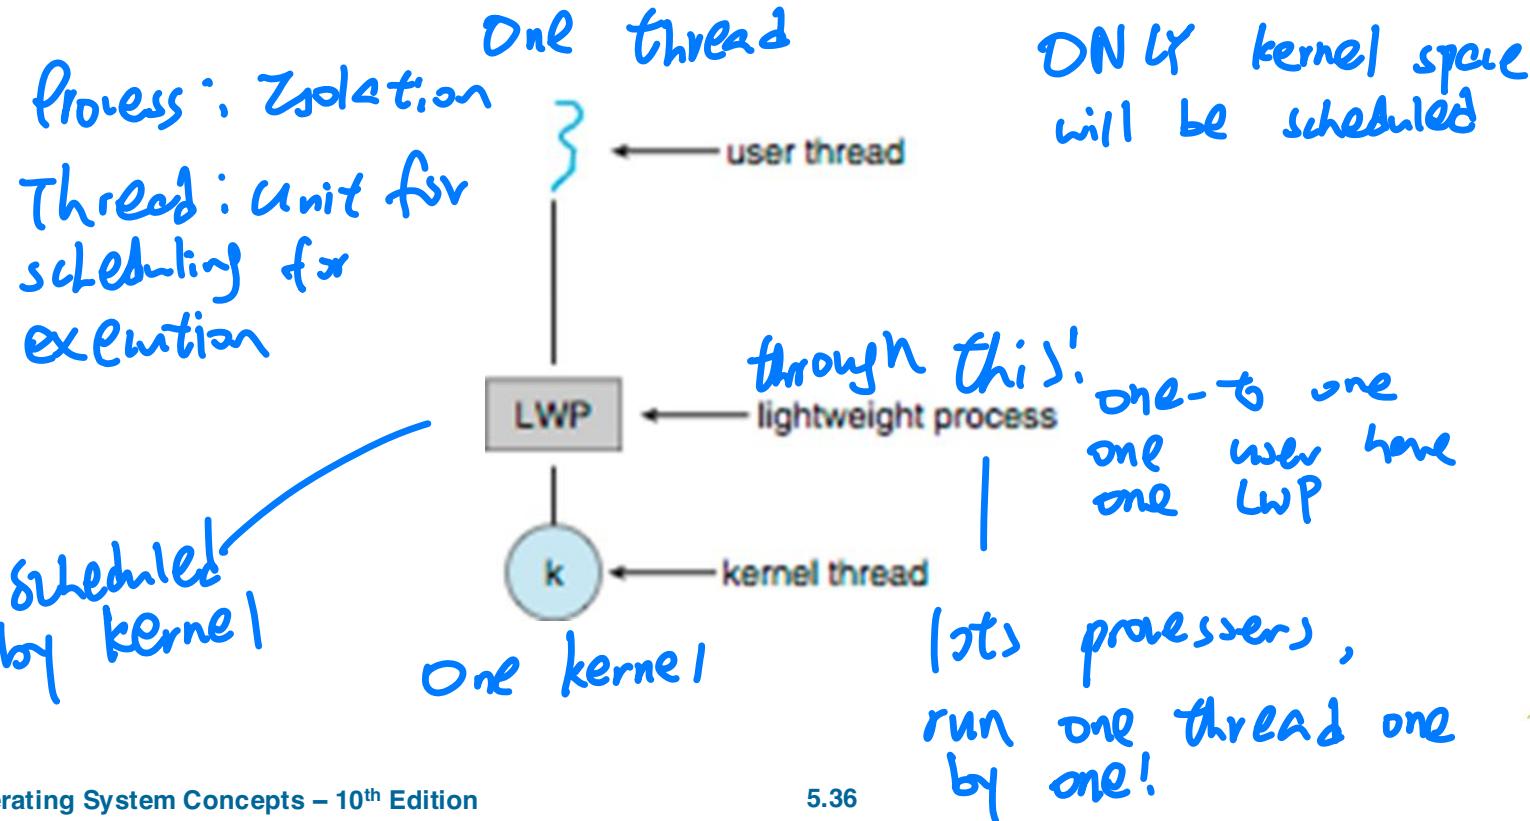
Assumptions :

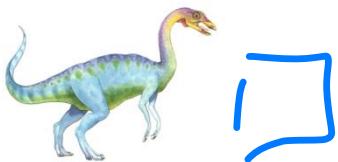
- Only one thread / one process with one thread
- only one CPU
- But modern OS, have different cores!
 - ? Mapping to user space and kernel space



Thread Scheduling

- On modern OS, kernel-level threads are the ones being scheduled, user-level threads are managed by a thread library instead
- The OS uses an intermediate data structure between user threads and kernel threads, a lightweight process (LWP) *May > than yr core*
 - Appears as a virtual processor on which user threads are scheduled to "run"
 - Each LWP attached to kernel thread (one-to-one)





Thread Scheduling

- Under *many-to-one* and *many-to-many* models, thread library “schedules” user-level threads to run on LWP. This is known as **process-contention scope (PCS)** – determine which user thread is mapped to which kernel thread
 - in a 1:1 X one to one, only need to many to one/many!*
- Since scheduling competition takes place among the threads belonging to the same process, typically done via **priority** set by programmers
 - Thread library usually can not adjust the priority *Linux X care PCS*
 - PCS will typically **preempt** the thread currently running in favour of a higher-priority user-level thread *Modern OS → one to one*
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** – competition among all threads in the system – **CPU scheduling**
- Systems using one-to-one mapping model, such as Windows, Linux, and Solaris, schedule threads using only SCS

*kernel use the algorithms!
- Just schedule those threads*

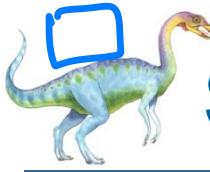




Multiple-Processor Scheduling

- Maximize the utilization* More complicated
- CPU scheduling far more complex with multiple CPUs – load sharing
 - Traditionally, the term multiprocessor referred to systems that provide multiple physical processors, where each physical processor chip contains one single-core CPU *↳ multiple physical CPU*
 - The definition of multiprocessor has evolved significantly, and in modern computing systems, multiprocessor now applies to multicore CPUs, multithreaded cores, NUMA systems, and heterogeneous multiprocessing
 - There are generally two types of multiprocessing systems, *↳ own memory!* Asymmetric multiprocessing and symmetric multiprocessing *↳ only one can be asymmetric, can't be symmetric*
 - Asymmetric multiprocessing – only one processor can access kernel data structures, alleviating the need for data sharing. The other processors execute only user codes
 - All scheduling decisions, I/O processing, and other system activities handled by a single processor — the master server
 - The master server can become a potential bottleneck
- X in modern day! ↑ Everything goes that!





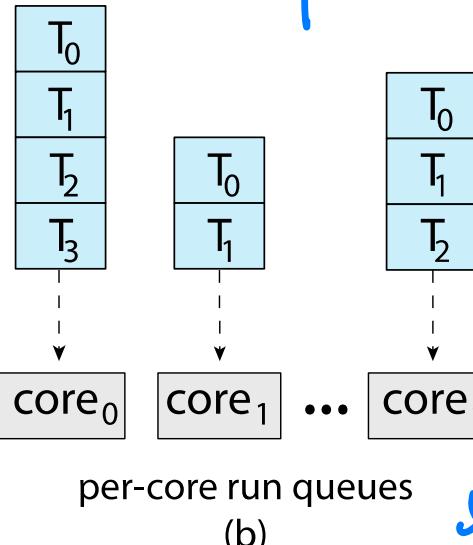
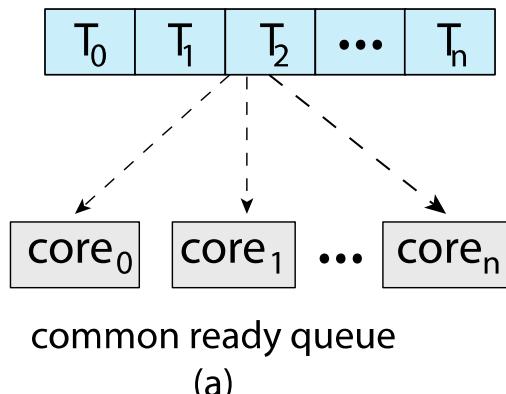
Symmetric multiprocessing (SMP)

execute any thread!

- Symmetric multiprocessing (SMP): Each processor is self-scheduling, with one common ready queue, or each having its own private ready queue
 - Scheduling proceeds by having the scheduler of each processor examines the ready queue and select a thread to run
 - To ensure two separate processors do not choose to schedule the same thread with a common ready queue - possible race condition (discuss in Chapter 6)
 - All modern OS support SMP, including Window, Linux, Mac OS X, as well as mobile systems including Android and iOS

Can Schedule at more as possible, keep all busy!

Ready queue!



common ready queue

(a)

per-core run queues

(b)

Easy for load and share!

how to distribute yr workload evenly, fairly!

Schedule in the same order, implement in affinity!



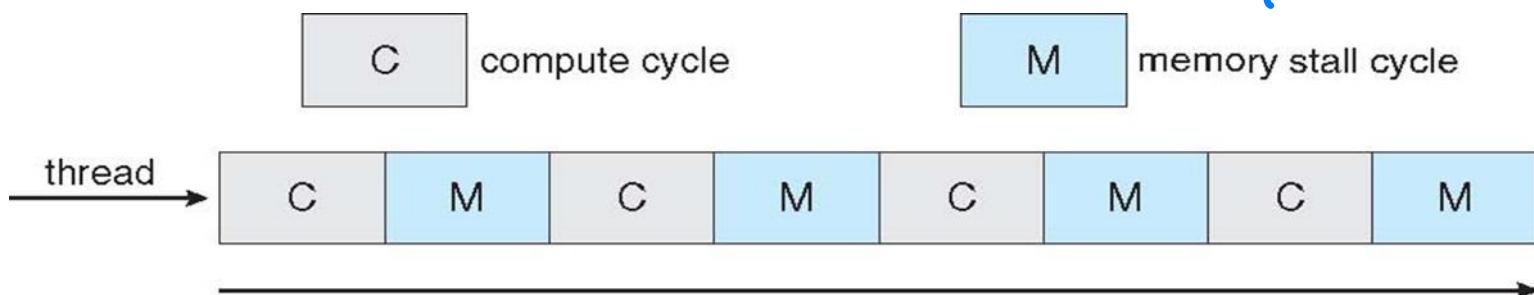
Multicore Processors

- 1 chip, 534 processor!

- Recent trend to place multiple processor cores on same physical chip
 - Faster and consumes less power, but complicate the scheduling design
- Memory stall:** when a processor accesses memory or cache, it spends a significant amount of time waiting for data to become available, primarily because modern processors operate at much faster speeds than memory, esp. when there is a cache miss
- Multiple hardware threads per core - each hardware thread has its own state, program counter (PC), register set appearing as a logical CPU to run a software thread. This is known as chip multithreading (CMT)

Done by hardware!

Consider memory access time!



overlap two threads
time
computation — two threads concurrently
memory computation

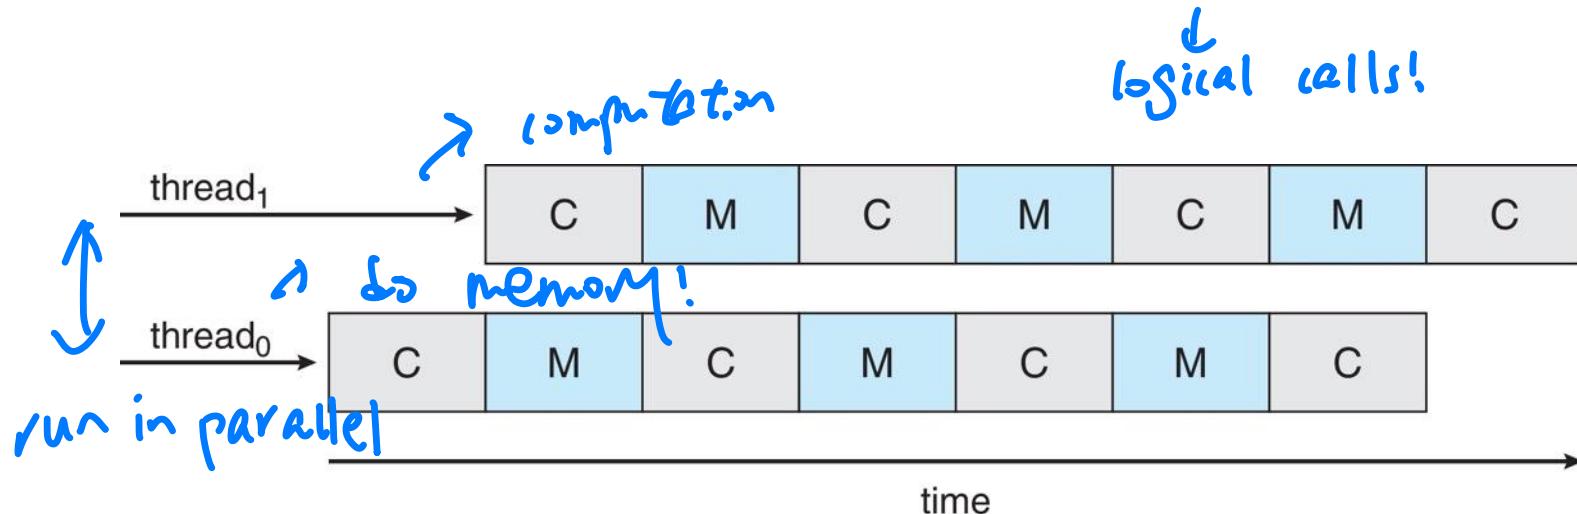


Multithreaded Multicore System

The scheduling can take advantage of **memory stall** to make progress on another hardware thread while memory retrieve happens for one hardware thread

one physical core!

- If one thread stalls while waiting for memory, the core can switch to another thread. This becomes a **dual-thread processor core**, or resembles two logical processors
- A **dual-threaded, dual-core system** presents **four logical processors** to the operating system
- UltraSPARC T3 CPU has 16 cores per chip and 8 hardware threads per core, from operating system perspective, this appear to be **128 logical processors**

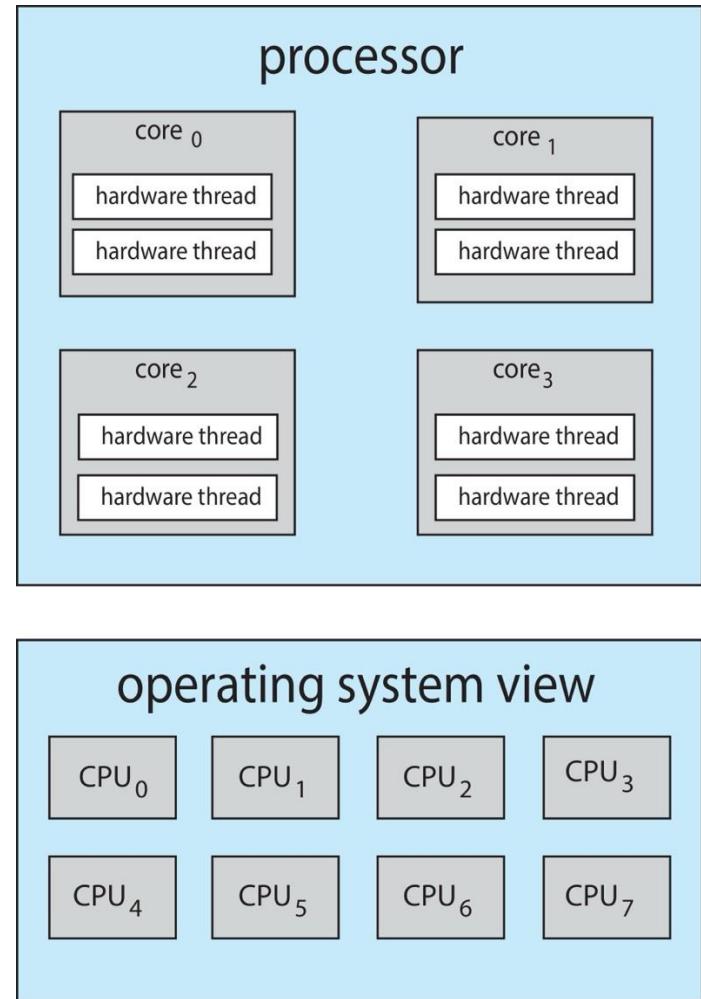




Multithreaded Multicore System

- From OS perspective, each hardware thread maintains its architectural state, such as instruction pointer and register set, and thus appears as a logical CPU that is available to run a software thread, i.e., a kernel thread
- Chip-multithreading (CMT) assigns each core multiple hardware threads. (Intel refers to this as hyperthreading)
- On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.

↓
usually done by RR!





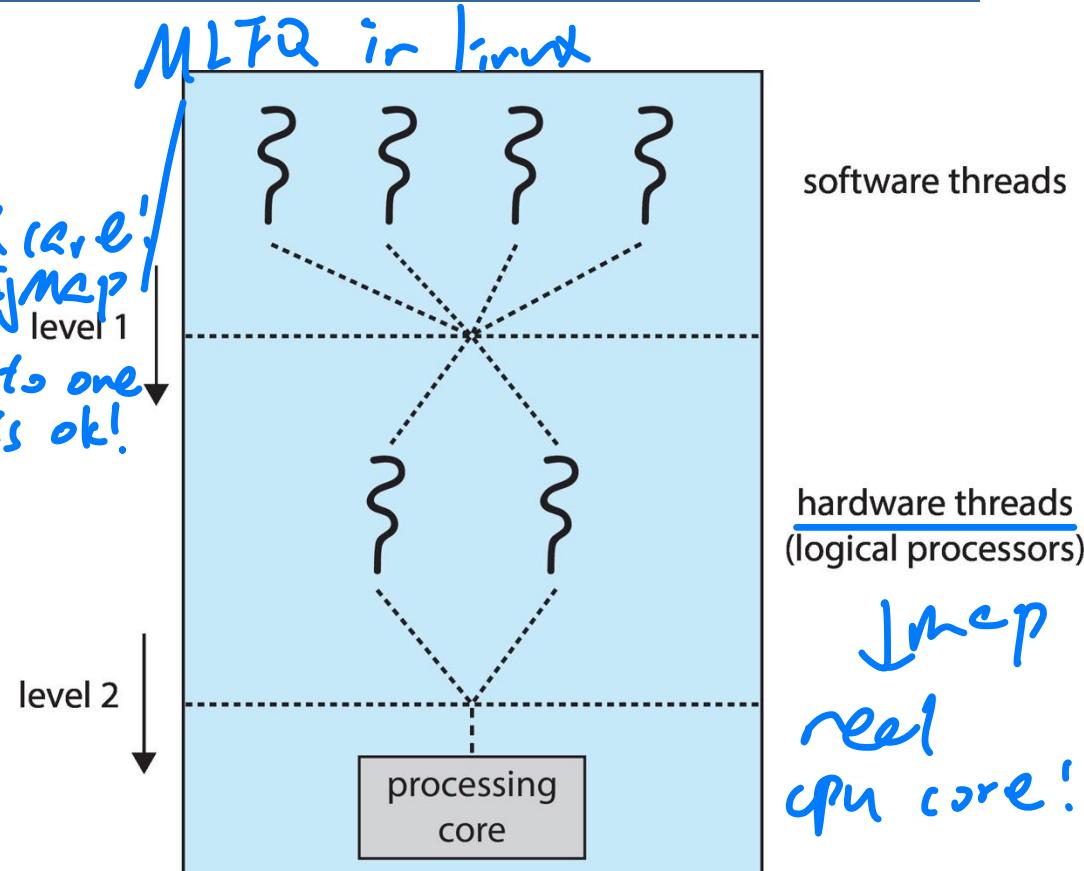
Multithreaded Multicore System

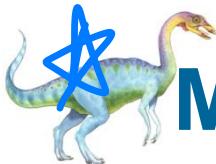
CP

- Two levels of scheduling:

1. The operating system deciding which software thread (kernel thread) to run on a logical CPU – CPU is ok! scheduling that we cover in this Chapter

↳ one!





Multithreaded Multicore Scheduling

How to schedule to run in the actual CPU

- A user-level thread is scheduled to a LWP – a kernel-level thread
 - Under many-to-one and many-to-many models, thread library schedules user-level threads known as process-contention scope (PCS), typically based on the priority set by programmers

Done by OS!

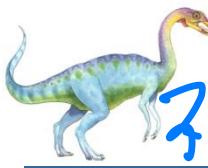
- A kernel thread now referred as a software thread is scheduled onto a logical CPU – a hardware thread (SCS!)
 - CPU or process scheduling – operating system

Third level!

- A hardware thread is scheduled to run on a CPU core
 - Each CPU core decides the scheduling, typically using RR - Chip Multithreading.

? User space to control?





Processor Affinity

Processor affinity – process has an affinity for a processor on which it runs

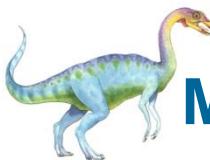
- When a thread has been running on one processor, the cache content of that processor stores the memory accesses by that thread. We refer to this as a thread having affinity for a processor (i.e., “processor affinity”) 不然多大會

不然後
多大會
diverged
 - There is a high cost of invalidating and repopulating caches, most SMP systems try to avoid migration of processes from one processor to another
 - Essentially, per-processor ready queues provide processor affinity for free!
 - Soft affinity** – the OS attempts to keep a process running on the same processor (not guaranteeing that), and it is possible for a process to migrate between processors during load balancing
 - Hard affinity** – allow a process to specify a subset of processors it may run
 - Many systems provide both soft and hard affinity
 - For example, Linux implements soft affinity, but it also provides the system call `sched_setaffinity()`, which supports hard affinity by allowing a thread to specify the set of CPUs on which it is eligible to run

Workload →
fairly, equally
distributed!

Just some concepts:





Multiple-Processor Scheduling – Load Balancing

- Load balancing attempts to keep workload evenly distributed

fairly distributed for all CPU cores

- On systems where each processor has its own private ready queue of eligible threads to execute
- There are two general approaches to load balancing

→ Race condition → only CPUs have own ready queue!

Active □ Push migration – a specific task periodically checks the load on each processor, and if it finds an imbalance, pushes task(s) from overloaded CPU to idle or less-busy CPUs

baking all the loads!!!

how to assign the ready queue, not fair

Passive □ Pull migration – idle processors pulls waiting task(s) from a busy processor

□ Push and pull migration are not mutually exclusive and in fact often both implemented in parallel on load-balancing systems. For instance, Linux CFS

only when triggered implement both techniques

□ Load balancing often counteracts the benefits of processor affinity - natural tension between load balancing and minimizing memory access times

Always wants to run in the same processor

□ Thus, scheduling algorithms for modern multicore NUMA systems have become quite complex

trying to achieve both advantages!



3 layers scheduling

1.

2. Schedule correctly from kernel space to
logical processor

3. hardware thread is scheduled to CPU core
- by RR!



Real-Time CPU Scheduling

usually

finish in specific deadline! Quite different

- Real-time scheduling demands performance guarantee – predictability
- Hard real-time systems – have stricter requirements. A task must be serviced by its deadline; service after the deadline has expired is the same as no service at all → very severe consequence!
- Soft real-time systems – provide no guarantee as to when a critical real-time process will be scheduled. They guarantee only that the real-time process will be given preference over noncritical processes *try to meet deadline!*
- The scheduler for a real-time operating system must support a priority-based algorithm with preemption

Depends on
the usage:
x care about CPU utilization

x care about time!

✓ have deadline → *不然会被扣分！*

✓ finish execution before this deadline



Real-time CPU scheduling

1. Priority based

2. Preemptive

RMS

- Fixed priority algorithm.
will not change during the whole life cycle

EDF

- Changing priority based on actual deadline!

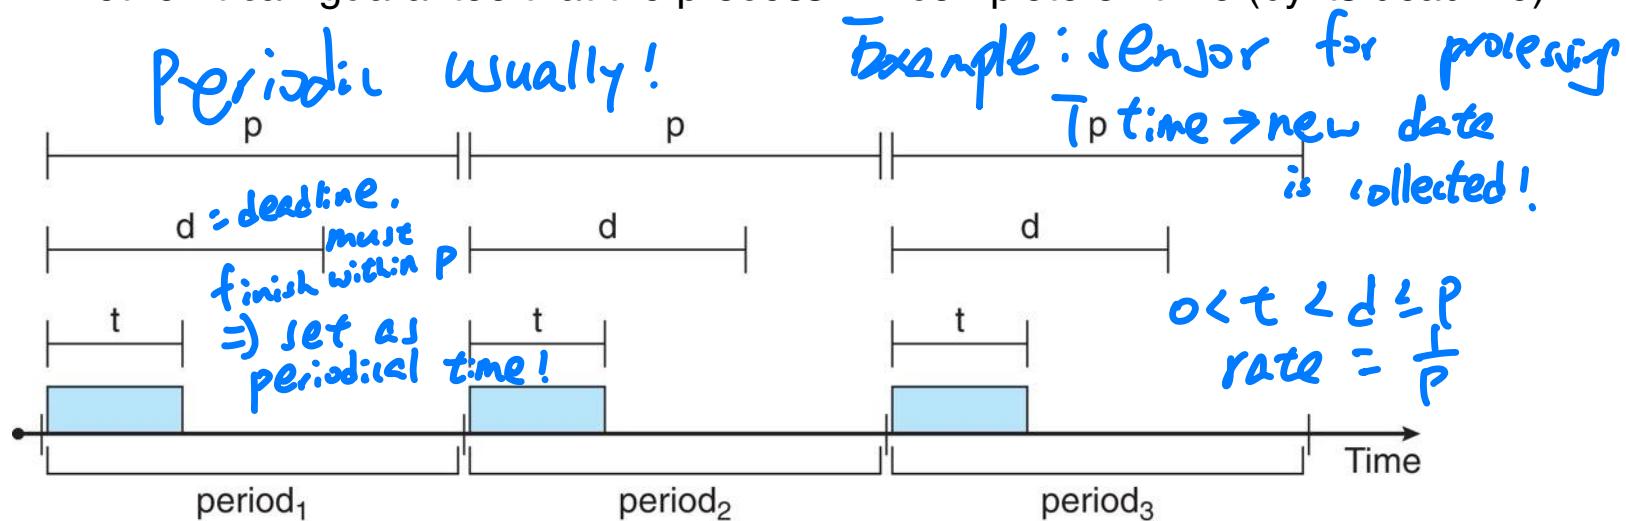


Go deep next time! → Always preemptive! Set priority based on deadline!!!

Priority-based Scheduling

finish time is predictable!

- Note that providing a preemptive, priority-based scheduler only guarantees soft real-time functionality. Processes have the characteristics: **periodic** ones require CPU at constant intervals (periods)
 - Has processing time t , deadline d , period p , in which $0 \leq t \leq d \leq p$
 - The **rate** of a periodic task is $1/p$ *task will come every t time*
 - A process may have to announce its deadline requirements to the scheduler. The scheduler decides whether to admit the process or not depending on whether it can guarantee that the process will complete on time (by its deadline)



Fixed period!

Common: $d=p$
- some process can announce the deadline!
(optional)



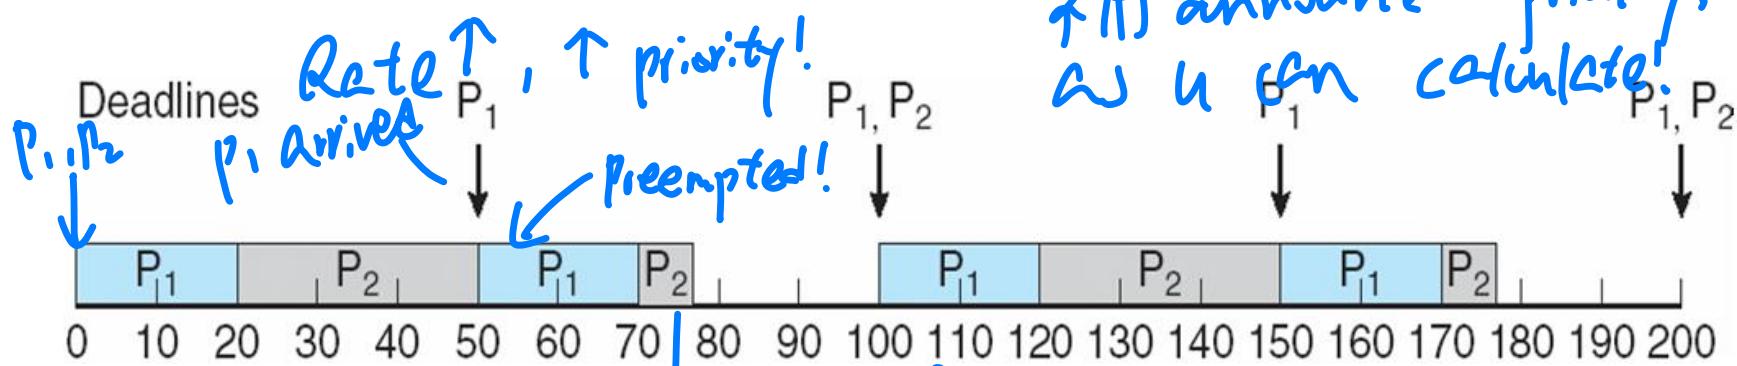


more frequently!

Priority Exchange!!!

Rate Montonic Scheduling

- A **static priority** is assigned based on the inverse of its period
 - Shorter (longer) period = higher (lower) priority; **only periodic arrival!!!*
 - The rationale is to assign a higher priority to tasks requiring CPU more often
- Suppose P_1 has a period of 50 (also deadline), and processing time 20. P_2 has a period of 100 (also deadline), and processing time 35.
 - Shorter period = higher priority!*
 - P_1 is higher priority! static priority*
 - The deadline for each process requires that it complete its CPU burst by the start of its next period. Since $50 < 100$, P_1 is assigned a higher priority than P_2
 - The CPU utilization of a process P_i is the ratio of its burst to its period t_i/P_i the CPU utilization of P_1 is $20/50 = 0.40$ and that of P_2 is $35/100 = 0.35$, so the total CPU utilization of 75 percent



↑ Announce priority,
as u can calculate!

only process 35 second is ok

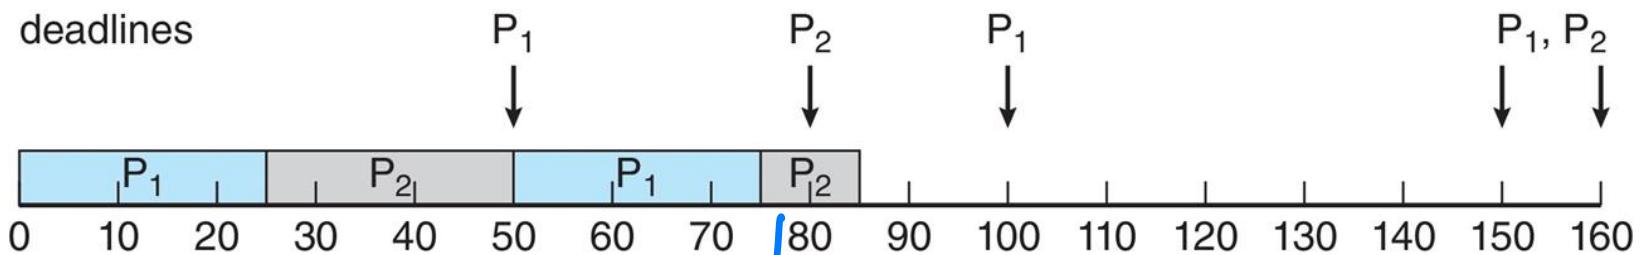
∴ But if P_2 processing time ↑↑↑
it will miss the deadline → 不可否實件





Missed Deadlines with Rate-Monotonic Scheduling

- Suppose P_1 has a period of 50 (also deadline), and processing time 25. P_2 has a period of 80 (also deadline), and processing time 35.
 - Since $50 < 80$, P_1 is assigned a higher priority than P_2 Preempted P_2 !
 - The CPU utilization of P_1 is $25/50 = 0.50$ and that of P_2 is $35/80 = 0.44$, for a total CPU utilization of 94 percent
- Process P_2 misses the deadline (80) by finishing at time 85



- Shows drawbacks of RMS
- May change to dynamic priority! ↳ impossible to miss deadline ↳ only can improve



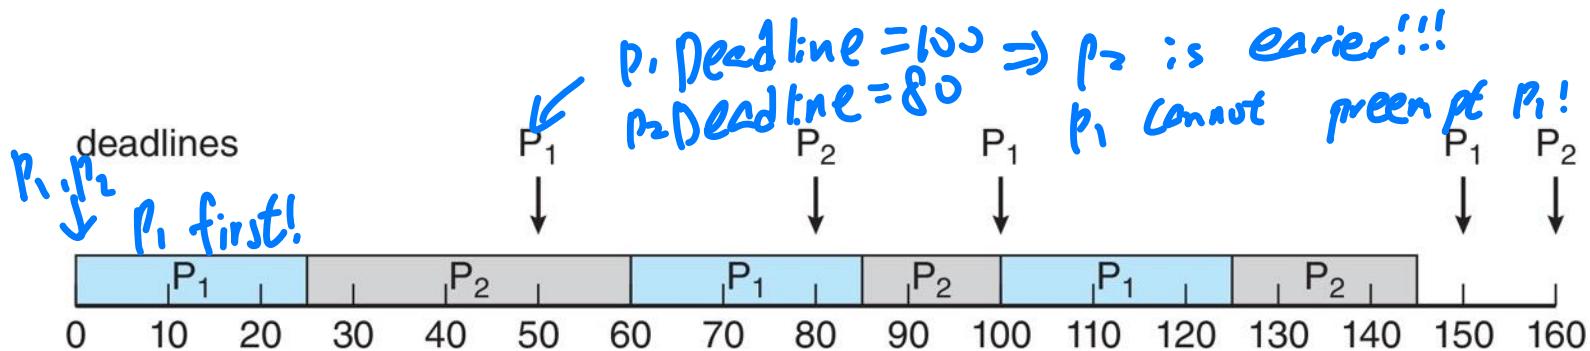


Earliest Deadline First Scheduling (EDF)

↑ priority may be changed

- Earliest-deadline-first (EDF) scheduling assigns dynamic priorities according to the deadline
 - the earlier (later) the deadline, the higher (lower) the priority
- Consider the same example, where P1 has a period of 50 (also deadline), and processing time 25. P2 has a period of 80 (also deadline), and processing time 35.

Check the deadline!



- Need to know the deadline
→ announce the deadline based on current time!





Rate-monotonic vs. EDF Scheduling

Rm

- The **rate-monotonic** scheduling algorithm schedules periodic tasks using a static priority policy with preemption \Rightarrow priority based on periodic!
不会这样 priority → optimal
- The rate-monotonic scheduling is considered to be **optimal** in that if a set of processes cannot be scheduled by this algorithm, it cannot be \Rightarrow 永远不会!
scheduled by any other algorithm that assigns static priorities. 永远不会 periodic!
- Unlike rate-monotonic algorithm, **EDF** scheduling does not require that processes be periodic, nor must a process require a constant amount of CPU time per burst. The only requirement is that a process announce its deadline to the scheduler when it becomes runnable need know deadline!
- EDF scheduling is **theoretically optimal** - it can schedule processes such that each process can meet its deadline requirements and CPU utilization will be 100 percent \Rightarrow extreme case may not work!
 - In practice, however, it is impossible to achieve this level of CPU utilization due to the cost of context switching between processes and interrupt handling

flexible, powerful
overhead!

Traditional CPU scheduling
real time scheduling
- different goals!





Methodology!

Algorithm Evaluation

long tail distribution! \rightarrow 80% of processes
time quantum!

- Selecting CPU-scheduling algorithm in practice can be difficult – as there are many scheduling algorithms, each with its own set of parameters
- The first problem is defining the criteria to be used in selecting an algorithm - often defined in terms of CPU utilization, response time, or throughput *Different w. terms?*
- Determine criteria – the criteria may include several measures with their relative importance, such as
 - Maximizing CPU utilization under the constraint that the maximum response time is 300 milliseconds
 - Maximizing throughput such that turnaround time is (on average) linearly proportional to total execution time





Deterministic Modeling

lots of concrete examples / numbers

QA also can deal with normal distribution! 7-7a | 3 methods!

- Deterministic modeling takes a particular predetermined workload and defines the performance of each algorithm for that workload
- Deterministic modeling is simple and fast. It gives us exact numbers, to compare algorithms. However, it requires exact numbers for input, and its answers apply only to those cases
- How processes run vary from time to time, so there is no static set of processes (or times) to use for deterministic modeling
- Consider 5 processes arriving at time 0:

Usually don't
know the data

- Data is dynamic!

Process	Burst Time
P_1	10
P_2	29
P_3	3
P_4	7
P_5	12

↓ to evaluate the system!

→ deterministic number!

→ have predicted time!





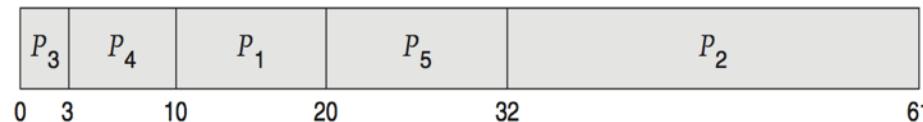
Deterministic Evaluation

Draw chart!!!

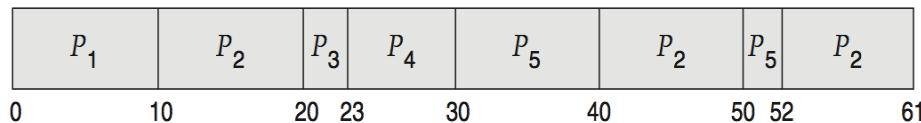
- For each algorithm, calculate the average waiting time
- Simple and fast, but requires exact numbers for input, applies only to those inputs
 - FCFS is 28ms:



- Non-preemptive SJF is 13ms:



- RR (quantum 10) is 23ms:





Queueing Analysis

Very complicated!
3/2 number, know
distribution \Rightarrow lie & M!

- Though the actual numbers (e.g., process arrival time, CPU or I/O bursts) vary from time (system) to time (system), the distributions of CPU and I/O bursts, and process arrival-time can be possibly measured and then approximated or simply estimated *(Collect lots of data \Rightarrow draw distribution usually normal!)*
- The computer system can be described as a network of servers, and each server has a queue of waiting processes.
 - The CPU is a server with its ready queue, I/O system with its device queues
 - Commonly use the exponential distribution, and described by mean
- Knowing arrival rates and service rates, we can compute the utilization, average queue length, average wait time, and so on.
- This area of study is called **queueing-network analysis**
- Queueing analysis can be useful in comparing scheduling algorithms, but the classes of algorithms and distributions that can be handled are very limited. Often the assumptions for the mathematical models to be tractable are unrealistic in practice





Queueing Models *'powerful!'*

- complicated

- Mathematical approach for handling stochastic workloads
- n = average queue length ✓
- W = average waiting time in queue ✓
- λ = average arrival rate into queue ✓
- **Little's Formula** – in steady state (there are mathematical assumptions for this to hold, e.g., arrival rate must be smaller than service rate), processes leaving queue must equal processes arriving, thus:

$$\underline{n = \lambda \times W}$$

understand!

- Valid for any scheduling algorithm and arrival distribution
- For example, if on average 7 processes arrive per second, and normally 14 processes in queue, then average wait time per process = 2 seconds

- care about long term/average performance





last way!

Always we this! Simulations

↳ distributions &
numbers =
we simulations!

- Queueing models is restricted to a few known distributions
- Running **simulations** involves programming a model of the computer system – which is more accurate
- As the clock value is increased, the simulator modifies the system state to reflect the activities of the devices, the processes, and the scheduler.
- As the simulation executes, statistics that indicate algorithm performance are gathered and printed. *Based on knowledge!*
- The data to drive the simulation can be generated in several ways
 - Random number generator according to probability distributions - distributions can be defined mathematically (uniform, exponential, Poisson) or empirically
 - **Trace files** to monitor the real system and record the sequence of actual events

write codes → collect system output!

- Brute force way:

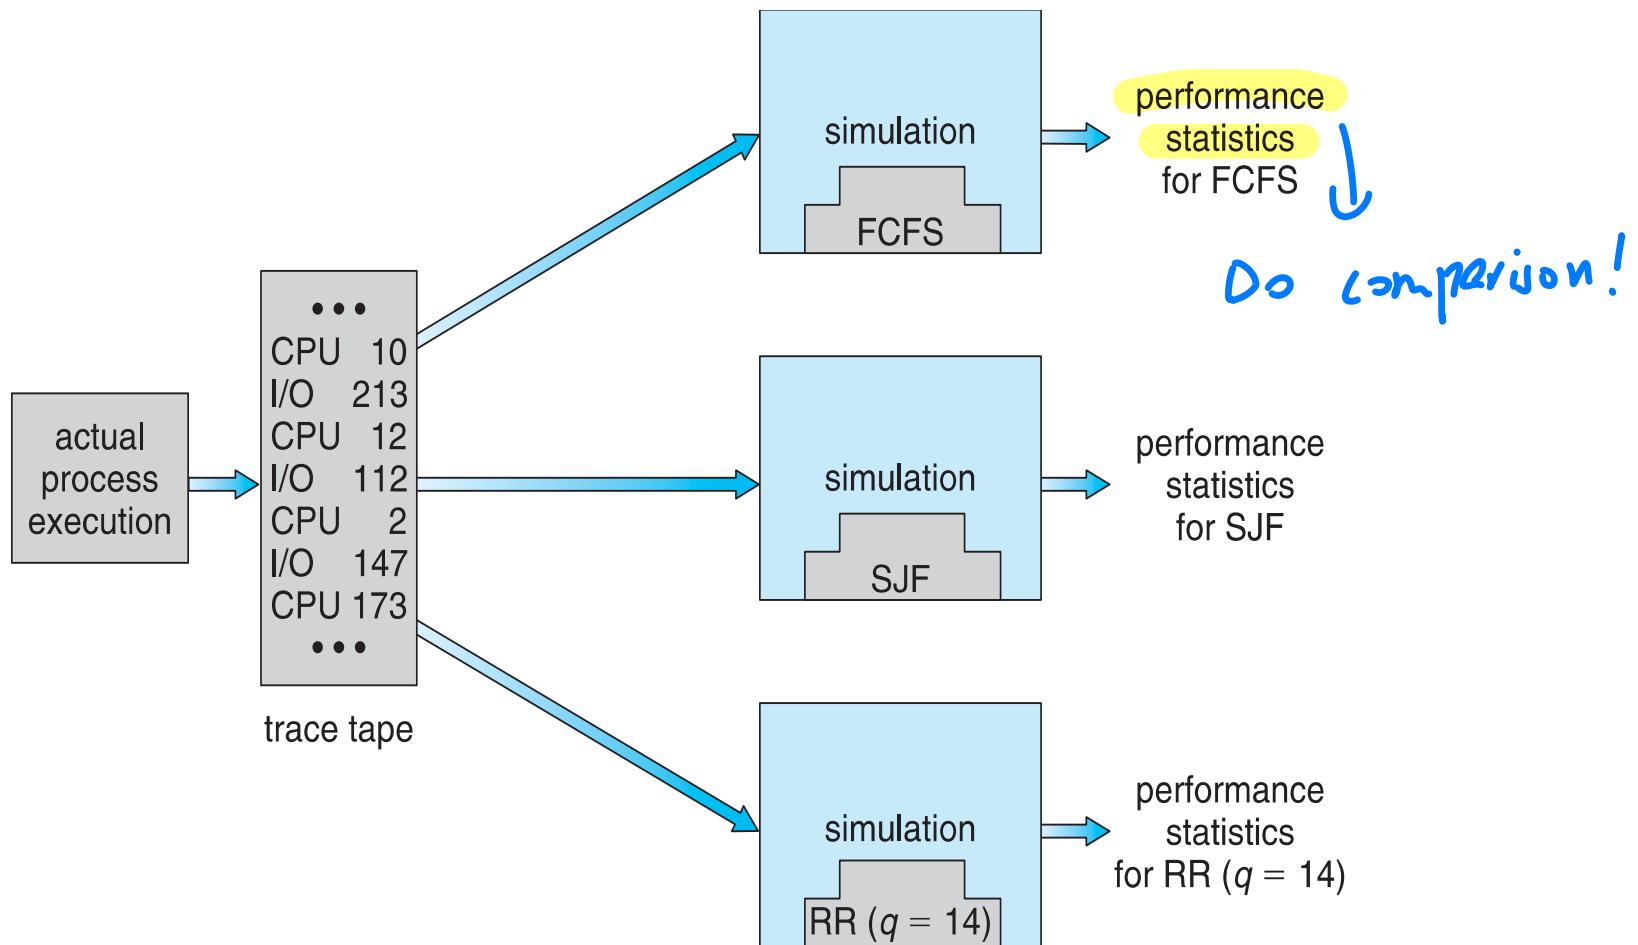
- Not accurate at all!

take as
input!





Evaluation of CPU Schedulers by Simulation





Implementation

- Even simulations have limited accuracy → *x 100% reflect:
of real life*
- Build a system which allows actual algorithms to run with real data set – more flexible and general. :)
- Implementing a new scheduler and test in real systems has difficulties:
 - This incurs high cost (coding the new scheduler), and high risk (e.g., potentially introducing new bugs)
 - Environments also changes constantly *various*
- Most flexible scheduling algorithms are those that can be altered by the system managers so that they can be tuned for a specific application
 - A system supporting graphical applications or web (file) service, for instance, may have different scheduling needs
 - Many UNIX systems allow the system manager to fine-tune the scheduling parameters for a particular system configuration
- APIs can be used to modify priority of a process or thread – improving performance of specific application, not overall application performance
 - The Java, POSIX, and Windows APIs provide such functions

Simulation first → deploy real work system



End of Chapter 5

