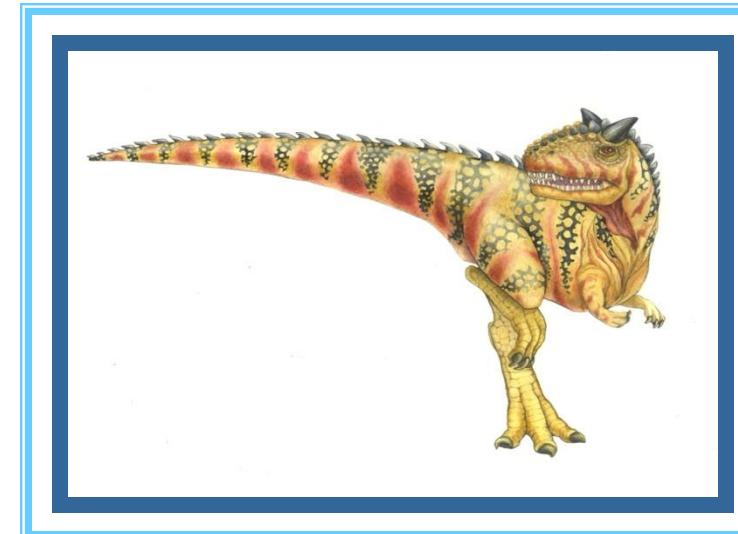


拔準

Chapter 6: Synchronization

Tools

Ring buffer
不~~用~~ index → Account

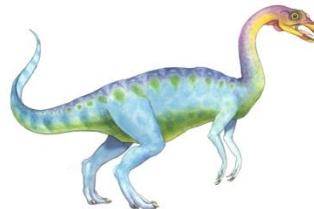




Chapter 6: Synchronization Tools

- Background
- The Critical-Section Problem
- Synchronization Hardware → function provided by hardware:
- ① Mutex Locks → lock
- ② Semaphores → Mutex Locks
- ③ Condition variables → if else





Objectives

- Describe the **critical-section problem** and illustrate the **race condition** *? Definition → just memorize!*
- Describe hardware solutions to the critical-section problem using **compare-and-swap operations**, and **atomic variables**
- Demonstrate how **mutex locks, semaphores, and condition variables** can be used to solve the critical section problem *? in ch 7!*

? software to help u to do sth.

Solution need to meet 3 requirements!

- ? define, ? to judge





Background

Processes
threads
Scheduling
may be interrupt, preemptive!
Order of execution will be changed

- Processes execute concurrently
 - Processes may be interrupted at any time, partially completing execution, due to a variety of reasons.
Do things together!
- Concurrent access to any shared data may result in data inconsistency
- Maintaining data consistency requires OS mechanisms to ensure the orderly execution of cooperating processes

*↳ Cause unexpected result! ↳
Order ≠ 一致!*





Illustration of the Problem

IPC problem → use ring buff!

- Think about the Producer-Consumer problem
 - ↳ consume data from the buff!
- An integer **counter** is used to keep track of the number of buffers occupied.
 - Initially, **counter** is set to 0
 - It is **incremented** each time by the **producer** after it produces an item and places in the buffer
 - It is **decremented** each time by the **consumer** after it consumes an item in the buffer.

Check whether full or not!

Corrected version





Producer-Consumer Problem

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++; // high level language, c code  
}
```

Producer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```

Everything is fine
but mind!!! This is high-level language
不是 instruction :^ machine code
Consumer

CPU do scheduling





Race Condition

- `counter++` could be implemented as

Addition in register!

```
register1 = counter
register1 = register1 + 1
counter = register1
```

Modern CPU will do some optimization!

process executed can be interrupted,
be preempted → 偷走

- `counter--` could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

*Addition → return value
→ assign*

Similar as addition

不可见
俩同时执行!

- Consider this execution interleaving with “count = 5” initially:

只在 blue code 后执行! 因为 OS can interrupt in one of 3 lines

```
S0: producer execute register1 = counter
S1: producer execute register1 = register1 + 1
S2: consumer execute register2 = counter
S3: consumer execute register2 = register2 - 1
S4: producer execute counter = register1
S5: consumer execute counter = register2
```

以公平 RR → switch out!

cause severe problem!!!

{register1 = 5}

{register1 = 6}

{register2 = 5}

{register2 = 4}

{counter = 6}

{counter = 4}

↓ context switch

↓ .

↓ .

66

Unlucky things happened!

*null pointer problem, seg. fault
overwrite! Change the correct result!*





Race Condition

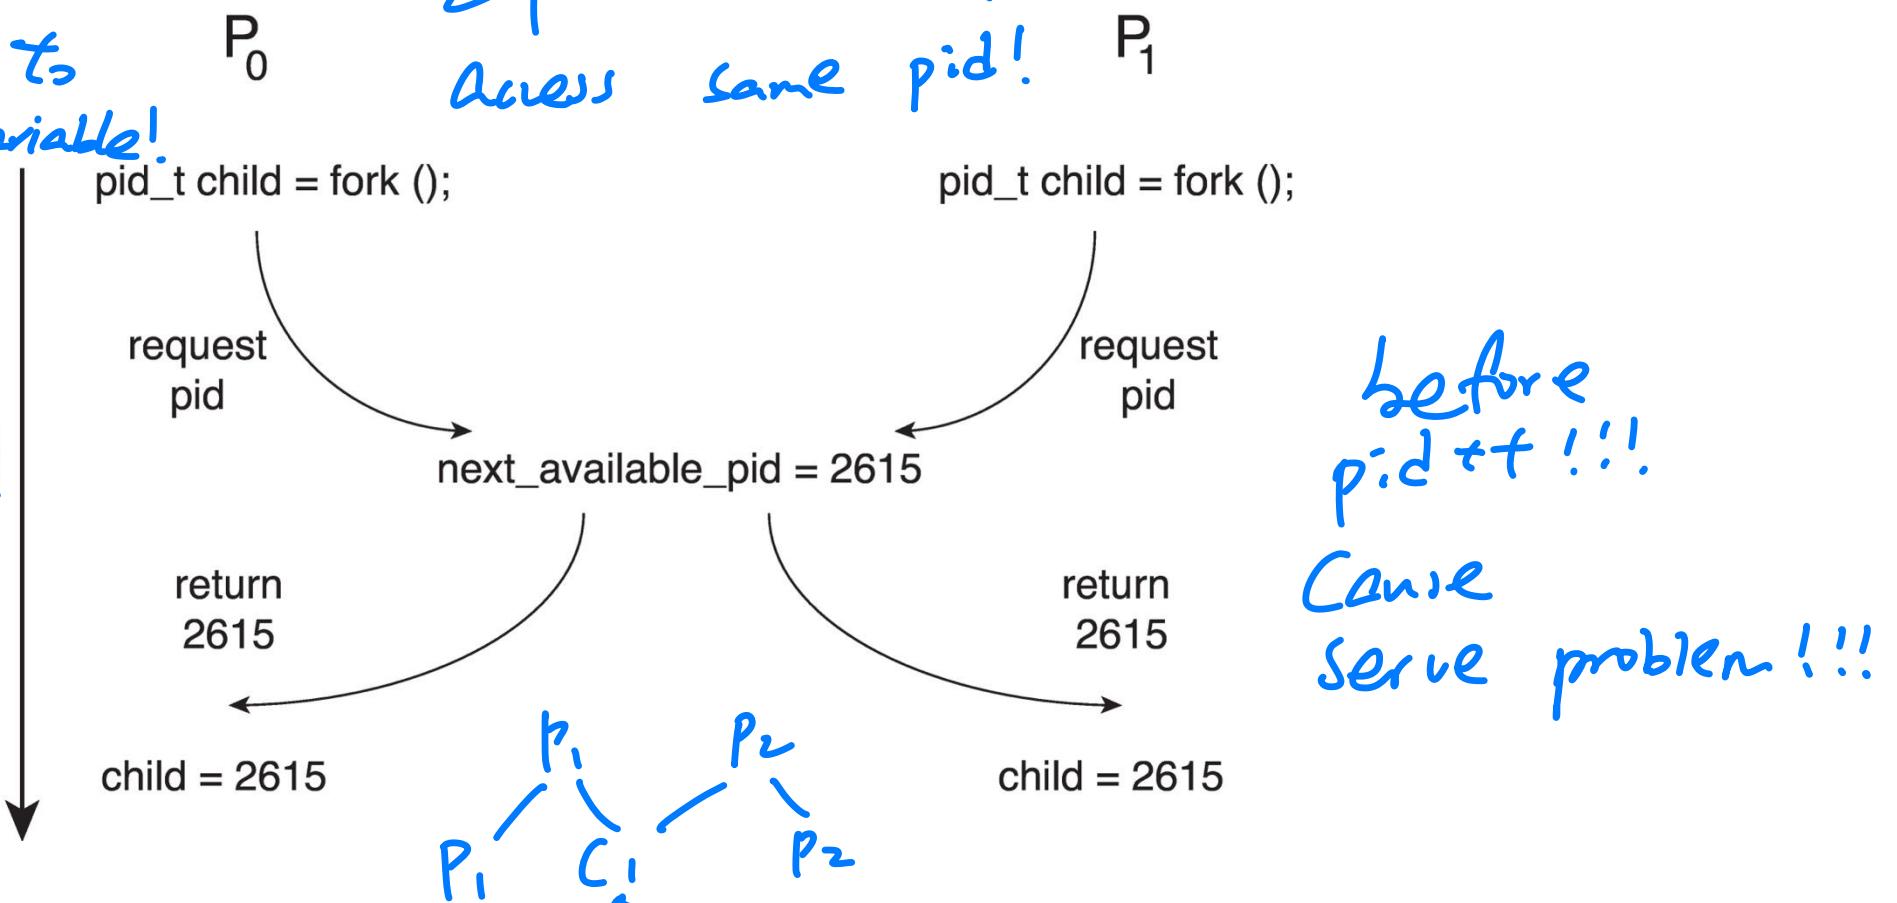
- Processes P_0 and P_1 are creating child processes using the `fork()` system call
- Race condition on kernel variable `next_available_pid` which represents the next available process identifier (pid)

- ≥ 2 process want to access the same variable!

Solution:

1. Disable interrupt
⇒ OS become low efficient

2. Make things mutually exclusive!



- Unless there is **mutual exclusion**, the same pid could be assigned to two different processes!

Partially solve this problem!

Only one access in one time!
Only allow one thread to access one variable at one time!

May cause other problems



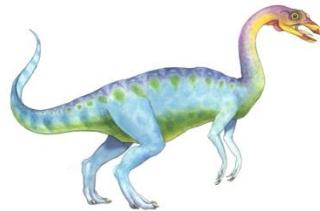
Critical Section Problem

Example: Counter!

- A **Race Condition** is an undesirable situation where several processes access or/and manipulate a **shared data concurrently** and the outcome of the executions depends on the particular order in which the accesses or executions take place - The results depend on the timing or the order of program execution. With some bad luck (i.e., context switches that occur at untimely points during execution), the result become **non-deterministic**

- Part of code* ↗ ↘ *not expected result!*
- Consider a system with n processes $\{P_0, P_1, \dots, P_{n-1}\}$
 - A process has a **Critical Section** segment of code (can be short or long), during which
 - A process or thread may be changing shared variables, updating a table, writing a file, etc.
 - We need to ensure when one process is in Critical Section, no other can be in its critical section
 - In a way, mutual exclusion and critical section imply the same thing
 - ↳ counter ++ \Rightarrow need to wait!!
 - Critical section problem is to design a protocol to solve this
 - Specifically, each process must request/ask permissions before entering a critical section in entry section, may follow critical section with exit section, then remainder section
- only allow one process + one thread to change the shared variable!





Critical Section

- The general structure of process p_i is

```
do {  
    — entry section  
    } → critical section  
    exit section  
    remainder section  
} while (true);
```

(locked room)

Countert + (其他不可访问 this)

Can be implemented by
other tools!

(key)
↓

done can give the key to others.





Solution to Critical-Section Problem

~~A~~ Meet 3 requirements at the same time

1. **Mutual exclusion** - If process P_i is executing in its critical section, no other processes can be executing in their critical sections

Most easy requirement!

*only one process!
Room with single key!*

2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical sections, the selection of a process that will enter the critical section next cannot be postponed indefinitely – selection of one process entering

~~A~~ *Very difficult to design! waiting in entry part → cannot delay it! Waiting no one executing → let that process come in!*

3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted – any waiting process

*Process try to enter the session → cannot let it
Assume speed are the same! try so many times!*

- Assume that each process executes at a nonzero speed, and there is no assumption concerning relative speed of each individual process

very unlucky → cannot enter the room

process due

*random assign the "key"
to the early session!*

unfairness issue! Starvation!

Equal opportunity

waiting trial is bounded → for fairness!





Critical-Section Problem in Kernel

- Kernel code - the code the operating system is running, is also subject to several possible race conditions
 - A kernel data structure that maintains a list of all open files can be updated by multiple kernel processes, i.e., two processes were to open files simultaneously
 - Other kernel data structures such as the one maintaining memory allocation, process lists or process table (keep track of processes in a system), interrupt handling etc.
 - Two general approaches are used to handle critical sections in operating system, depending on whether the kernel is preemptive or non-preemptive
 - Preemptive – allows preemption of process when running in the kernel mode, not free from the *race condition*, and becomes increasingly more difficult in SMP architectures.
 - Non-preemptive – runs until exiting the kernel mode, blocks, or voluntarily yields CPU. This is essentially free of race conditions in the kernel mode, possibly used in single-processor systems
- pop problem!* ↴
in kernel space
because codes also have preemptive
- Single processor
in single core ↴
no need
to have
this in
kernel
- interrupt driven! ← but modern OS
are multithreaded ↴
not practical!





Synchronization Tools

Processor IC pins

can provide the solutions

- Many systems provide hardware support for implementing the critical section code.
On uniprocessor systems – it could simply disable interrupts, currently running code would execute without being preempted or interrupted. But this is generally inefficient on multiprocessor or SMP systems
- Operating systems provide hardware and high level API support for critical section code

| Programs | Share Programs |
|--|--|
| Hardware | Load/Store, Disable Interrupts, Test&Set, Compare&Swap <i>in lectures</i> |
| High level APIs <i>include atomic levels!</i> | Locks, Semaphores <i>and. variables</i> |





Synchronization Hardware

- Modern OS provides special **atomic** hardware instructions *(Rai) interrupt*
 - ▶ **Atomic** = non-interruptible → *special instructions!*
 - ▶ This ensures the execution of atomic instruction can not be interrupted, thus, no race condition can occur
 - ▶ These also serve as building blocks for more sophisticated synchronization mechanisms

- There are two commonly used atomic hardware instructions, which can be used to construct more sophisticated synchronization *tools* *True or false*
 - Test a memory word and set a value – **Test_and_Set()** *return value*
 - Swap contents of two memory words – **Compare_and_Swap()** *] X be interrupted!*





test_and_set Instruction

□ Definition:

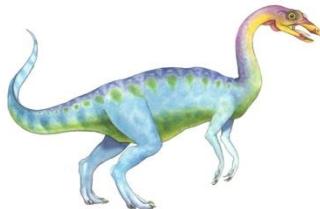
```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE; → set target to be true!
    return rv;
}
```

return bool (original value)

- 不会中断!

不用被中断





Solution using test_and_set()

?

- Shared Boolean variable lock, initialized to **FALSE**
- Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
    /* critical section */  
    lock = false;  
    /* remainder section */  
} while (true);
```

insured by hardware

- Access mutually exclusive easily!

\not satisfied the boundary!
most difficult





compare_and_swap Instruction

hardware!

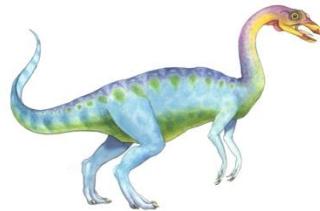
- Definition:

(called CAS)

```
int compare_and_swap(int *value, int expected, int new_value) {  
    int temp = *value;  
    if (*value == expected)  
        *value = new_value;           → same as expected true ⇒  
    return temp;                  *value = new value  
}  
}
```

Cannot be interrupted!





Solution using compare_and_swap

- Shared Boolean variable lock initialized to FALSE; Each process has a local Boolean variable key
- Solution:

```
do {  
    while (compare and swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
} while (true);
```

change test and set to this

X ensure bounded condition!





Bounded-waiting Mutual Exclusion with test_and_set

Shared by all processes!
& simplified!

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    /* through the while loop */ → enter
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false; /* no one is waiting, so release the lock */
    else
        waiting[j] = false; /* Unblock process j */
    /* remainder section */
} while (true);
```

only can access this
region

circular buffer!
Circular seen

no one waiting there!

When no one executing!
Everyone try to guess!
if nobody waiting → unlock!





Sketch Proof

- **Mutual-exclusion:** P_i enters its critical section only if either `waiting[i]==false` or `key==false`. The value of `key` can become false only if `test and set()` is executed. Only the first process to execute `test and set()` will find `key==false`; all others must wait. The variable `waiting[i]` can become `false` only if another process leaves its critical section; only one `waiting[i]` is set to false, thus maintaining the mutual-exclusion requirement. *ensure only one!*
- **Progress:** since a process exiting its critical section either sets `lock to false` or sets `waiting[j] to false`. Both allow a process that is waiting to enter its critical section to proceed.
- **Bounded-waiting:** when a process leaves its critical section, it scans the array `waiting` in cyclic order $\{i+1, i+2, \dots, n-1, 0, 1, \dots, i-1\}$. It designates the first process in this ordering that is in the entry section ($\text{waiting}[j]==\text{true}$) as the next one to enter the critical section. Any process waiting to enter its critical section will thus do so within $n-1$ turns.

One chance within $n-1$
Satisfied ALL conditions'!!





int / float

Atomic Variables

1. Mutually exclusive
 2. No waiters
 3. Boundary
- TAS CAS

- Typically, instructions such as compare-and-swap are used as building blocks for other more sophisticated synchronization tools.
- One tool is an **atomic variable** that provides atomic (uninterruptible) updates on basic data types such as integers and Booleans.
- For example, the increment() operation on the atomic variable sequence ensures sequence is incremented without interruption - increment(&sequence);
- The increment() function can be implemented as follows:

this increment
 cannot be
 interrupted!!!

```
void increment(atomic_int *v)
{
    int temp;
    do {
        temp = *v;
    } while (temp !=
        (compare_and_swap(v, temp, temp+1)));
}
```

false!

so no can't fit

count++ → no problem!

memorize!

if interrupted!
 not interruptable!

if CAS didn't return true!
 if change original value!

if interrupted!





true or
false!

Mutex Locks

software!

lock!

- OS builds a number of software tools to solve the Critical Section problem *True or false!*
- The simplest tool that most OSes use is mutex lock ↗ functions ↓
- To access the critical regions with it by first acquire() a lock then release() it afterwards – both atomic operations *Acquire the key* ↗
release the key! When leaves the room!
 - Boolean variable indicating if lock is available or not
(其他 異常 critical session!!)
- Calls to acquire() and release() must be atomic (non-interruptible)
 - Usually implemented via hardware atomic instructions
- But this solution requires busy waiting. This lock therefore called a spinlock ↗ similar term with *busy waiting!*
waste CPU cycles!
 - while(true){
 ; } ↑ CPU always checks condition! Do the check!
 □ Spinlock wastes CPU cycles due to busy waiting, but it has one distinct advantage in that no context switch is required when a process must wait on a lock, and a context switch may take considerable time. Thus, when locks are expected to be held for short times, spinlock is useful
 - Spinlocks are often used in multiprocessor systems where one thread can “spin” on one processor while another thread performs its critical section on another processor

Always consume CPU Spinlock ↗ no context switch!

100% utilization
하지만 다른 일을!

Low latency!
3ms ↗ Still in running status!!!





initialize available to be T/F

acquire() and release()

consume CPU cycles!!! As it continuously checks!

```
acquire() {  
    while (!available)  
        ; /* busy wait */ → no context switch  
    available = false;  
}  
  
release() {  
    available = true;  
}  
  
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

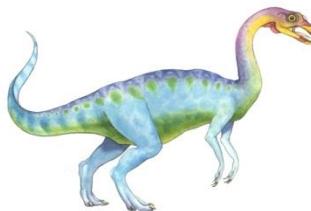
↑ consume time!!!
- Just wait for a short moment, acquire a lock do things in short delay!

Solutions based on the idea of **lock** to protect critical section

- Always in running status!!! + OS !!! hardware instruction
- Operations are **atomic** (non-interruptible) – at most one thread acquires a lock at a time
- Lock before entering critical section for accessing share data
- Unlock upon departure from critical section after accessing shared data
- Wait if locked - all synchronization involves busy waiting, should “sleep” or “block” if waiting for a long time

Busy wait
stuck in while loop!
CPU check again and again!
waste CPU cycles!





Semaphore

indicate available resources! (Permissions)

More powerful than mutex lock!

Positive zero

- Semaphore S – non-negative integer variable, can be considered as a generalized lock
 - First defined by Dijkstra in late 1960s. It can behave similarly as mutex lock, but it has more sophisticated usage - the main synchronization primitive used in original UNIX
- Two standard operations modify S : **wait()** and **signal()**
 - Originally called $P()$ and $V()$, where $P()$ stands for “proberen” (to test) and $V()$ stands for “verhogen” (to increment) in Dutch

- It is essential that semaphore operations are executed **atomically**, which guarantees that no more than one process can execute **wait()** and **signal()** operations on the same semaphore at the same time – serialization

- The semaphore can only be accessed via these two atomic operations except initialization

```

initial value = 1
wait (S) {
  while (S <= 0)
    no check! ; // busy wait
    S--;
  }
  no quota
signal (S) {
  S++;
  leave the room!

```

Fairly exclusive!
Allow ≥ 1 people to enter this room!

- Use lot of low level instructions and be b.i. cannot interrupted!





Semaphore Usage

Useful tools!

□ Counting semaphore – An integer value can range over an unrestricted domain

- Counting semaphore can be used to control access to a given set of resources consisting of a finite number of instances; semaphore value is initialized to the number of resources available

□ Binary semaphore – integer value can range only between 0 and 1

- This can behave like mutex locks, can also be used in different ways

$= 2 \rightarrow$ still race

condition!

- This can also be used to solve various synchronization problems

- Consider P_1 and P_2 that share a common semaphore synch, initialized to 0; it ensures that P_1 process executes S_1 before P_2 process executes S_2

$P_1:$
in parallel
 $S_1;$
 $\quad\quad\quad$ init. to 0

$\boxed{\quad}$
 $\quad\quad\quad$ signal(synch);

$P_2:$
 $\quad\quad\quad$ wait(synch);

$S_2;$

Always S_1 before S_2 !

Semaphore + mutex!!!

Do prove for 3 threads!!!





need context switch!
just website in the CPU hierarchy → save CPU cycles!

Semaphore Implementation with no Busy waiting

- Each semaphore is associated with a **waiting queue**
 - Each entry in a waiting queue has two data items:
 - value (of type integer)
 - pointer to next record on the queue

沒有 while loop!
Put in waiting
steps, or put back
to the waiting queue!
Context switch!

- Two operations:
 - still needs ~~you~~ member **block** – place the process invoking the operation on the appropriate waiting queue
 - wakeup – remove one of processes in the waiting queue and place it on the ready queue

Fig overwrite for the scheduling! no busy waiting \rightarrow do not CPU / consume CPU cycle!

- If a semaphore value is negative, its magnitude is the number of processes currently waiting on the semaphore.

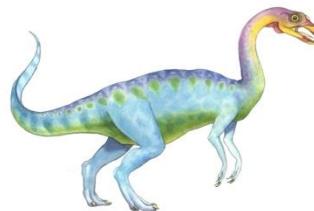
Save CPU cycles, CPU can be allocated to other threads.



Only running states can consume CPU cycle!

-! : only one thread in the waiting queue.

-x : x threads in the waiting queue



Semaphore Implementation with no Busy waiting (Cont.)

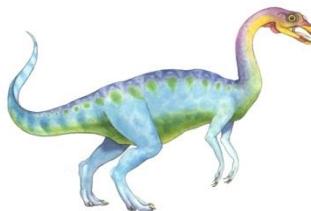
```
typedef struct{
    int value;
    struct process *list;
} semaphore;
wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}
signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```

less be used in modern OS.
Noticing that
usually we in class!

- Increment and decrement are done before checking the semaphore value, unlike the busy waiting implementation
- The **block()** operation suspends the process that invokes it.
- The **wakeup (P)** operation resumes the execution of a suspended process P.

wake up one by one!





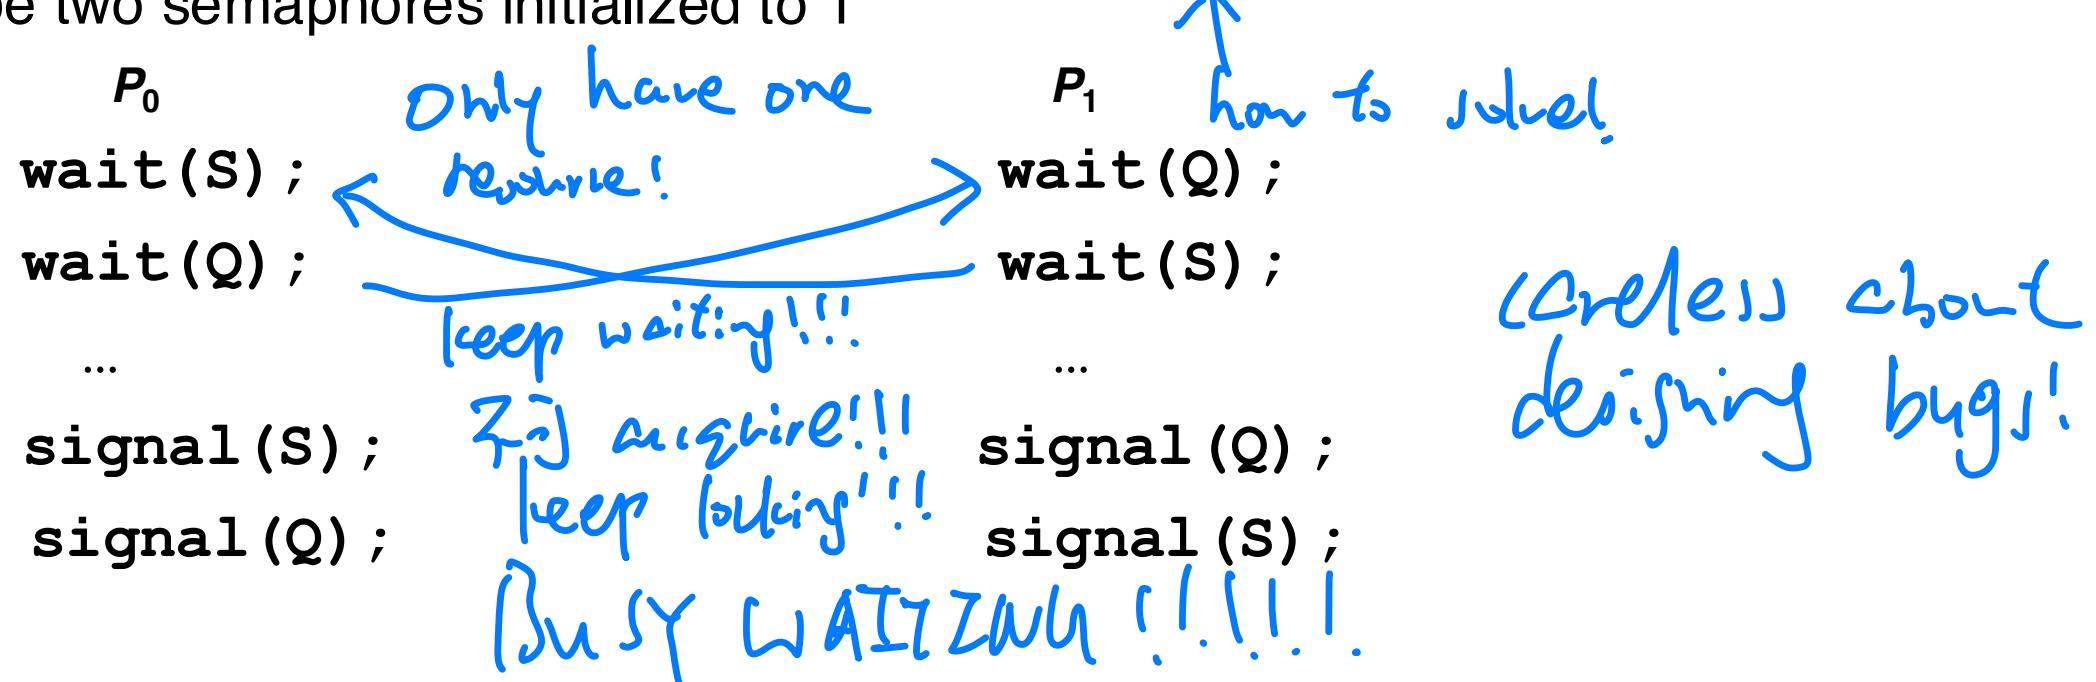
Deadlock and Starvation

classic!

Extremely common!

Difficult!!!

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes (to be examined in Chapter 8)
- Let S and Q be two semaphores initialized to 1



- Consider if P_0 executes `wait(S)` and P_1 executes `wait(Q)`. When P_0 executes `wait(Q)`, it must wait until P_1 executes `signal(Q)`. However, P_1 is waiting until P_0 executes `signal(S)`. Since these `signal()` operations will never be executed, P_0 and P_1 are **deadlocked**. This is extremely difficult to debug → waiting in while loop.
- **Starvation – indefinite blocking**

- A process may never be removed from the semaphore queue, in which it is suspended. For instance, if we remove processes from the queue associated with a semaphore using LIFO (last-in, first-out) order or based on certain priorities.



End of Chapter 6

