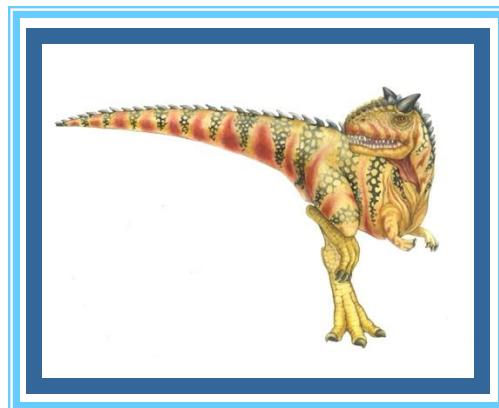


Chapter 17: Protection

whether process
is permitted

How to implement?

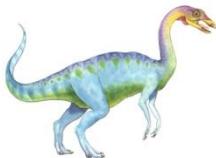




Chapter 17: Protection

- Goals of Protection Why we need ?
- Principles of Protection ✗ grab too much
- Protection Rings ✓ very first, 2 rings
- Domain of Protection more general, 2nd hierarchy structure
- Access Matrix How to represent a domain of protection!
- Implementation of Access Matrix



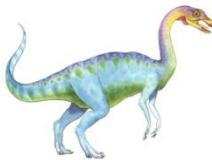


Objectives

- Discuss the goals and principles of protection in modern computer systems
- Explain how **protection domains** combined with **an access matrix** are used to specify the resources that a process may access
- Examine capability-based protection system

↙
有2/5 meaning!

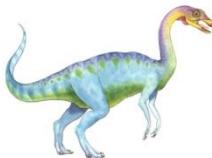




Goals of Protection

- In a protection model, computer system consists of a collection of **objects**, hardware or software
 - Hardware objects: CPU, memory segments, printers, disks, and tape
 - Software objects: files, programs, and semaphores *→ process can use objects*
 - Each object has a unique name and can be accessed through a well-defined set of operations *_objs only need , Rj write, execution.*
 - **Protection problem** is to ensure that each object is accessed correctly and only by those processes allowed to do so *what should be done, how something will be done*
 - **Mechanisms** are distinct from **policies**, in which mechanisms determine how something will be done, and policies decide what will be done.
implemented by OS
 - The separation is important for flexibility, as policies are likely to change *不同地太子同时向商可改!*
(Matrix)
 - The separation ensures that not every change in policy would require a change in the underlying mechanism.
what we place in matrix
- policies is defined user!*





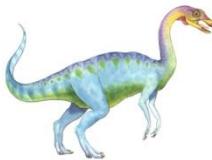
Principles of Protection

- The **guiding principle** – **principle of least privilege**
 - Programs, users and systems should be given **just enough privileges** to perform their tasks – as **more privileges for users, higher chance for attack**
 - In file permissions, for instance, this principle dictates if a user **only have read access but not write or execute access to a file**. **The principle of least privilege** would require that the OS provides a mechanism to **only allow read access but not write or execute access**
- Properly set **permissions** (i.e., the access rights to an object) **can limit damage if entity has a bug or gets abused**

↑
Good !!!

13. XP read only
DOS involve other files,
Only grab files
privilege
↓
relatively safe!





Protection Rings

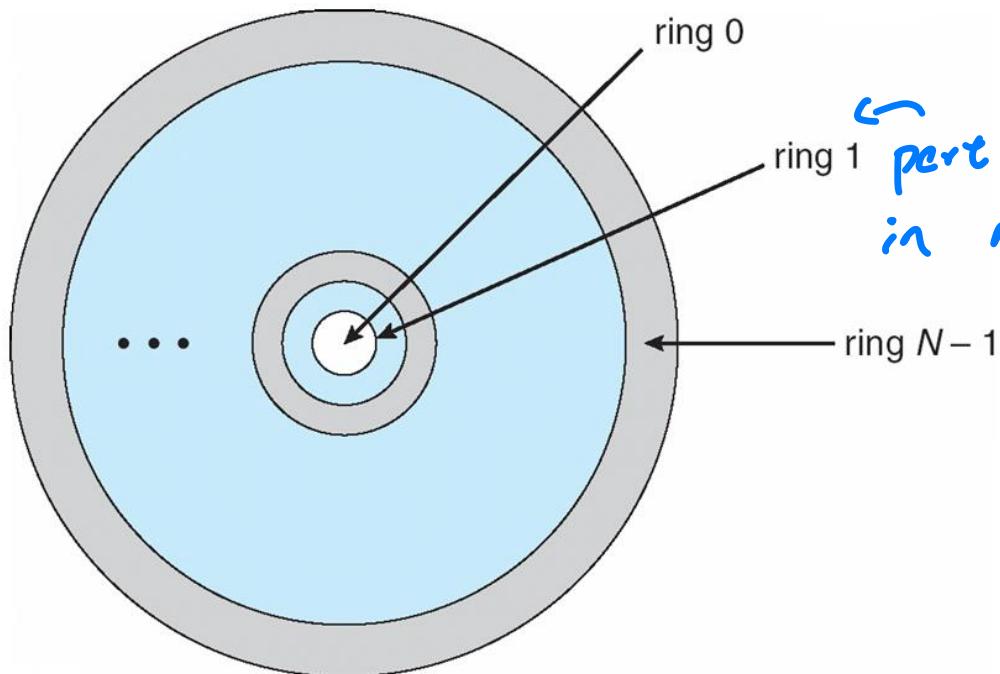
Revisit!

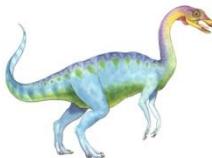
→ can do everything! (~~like API~~)

- Extension beyond **user mode** and **kernel mode** – **privilege separation**
- Hardware support required to support the notion of separate execution
- Let D_i and D_j be any two domain rings
- If $j < i \Rightarrow D_i \subseteq D_j$ *mean something (process/user) in modern OS!*
- The innermost ring, ring 0, provides the full set of privileges

Important concept

part of privileges
in ring 0!



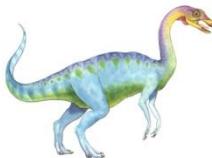


Protection Rings (Cont.)

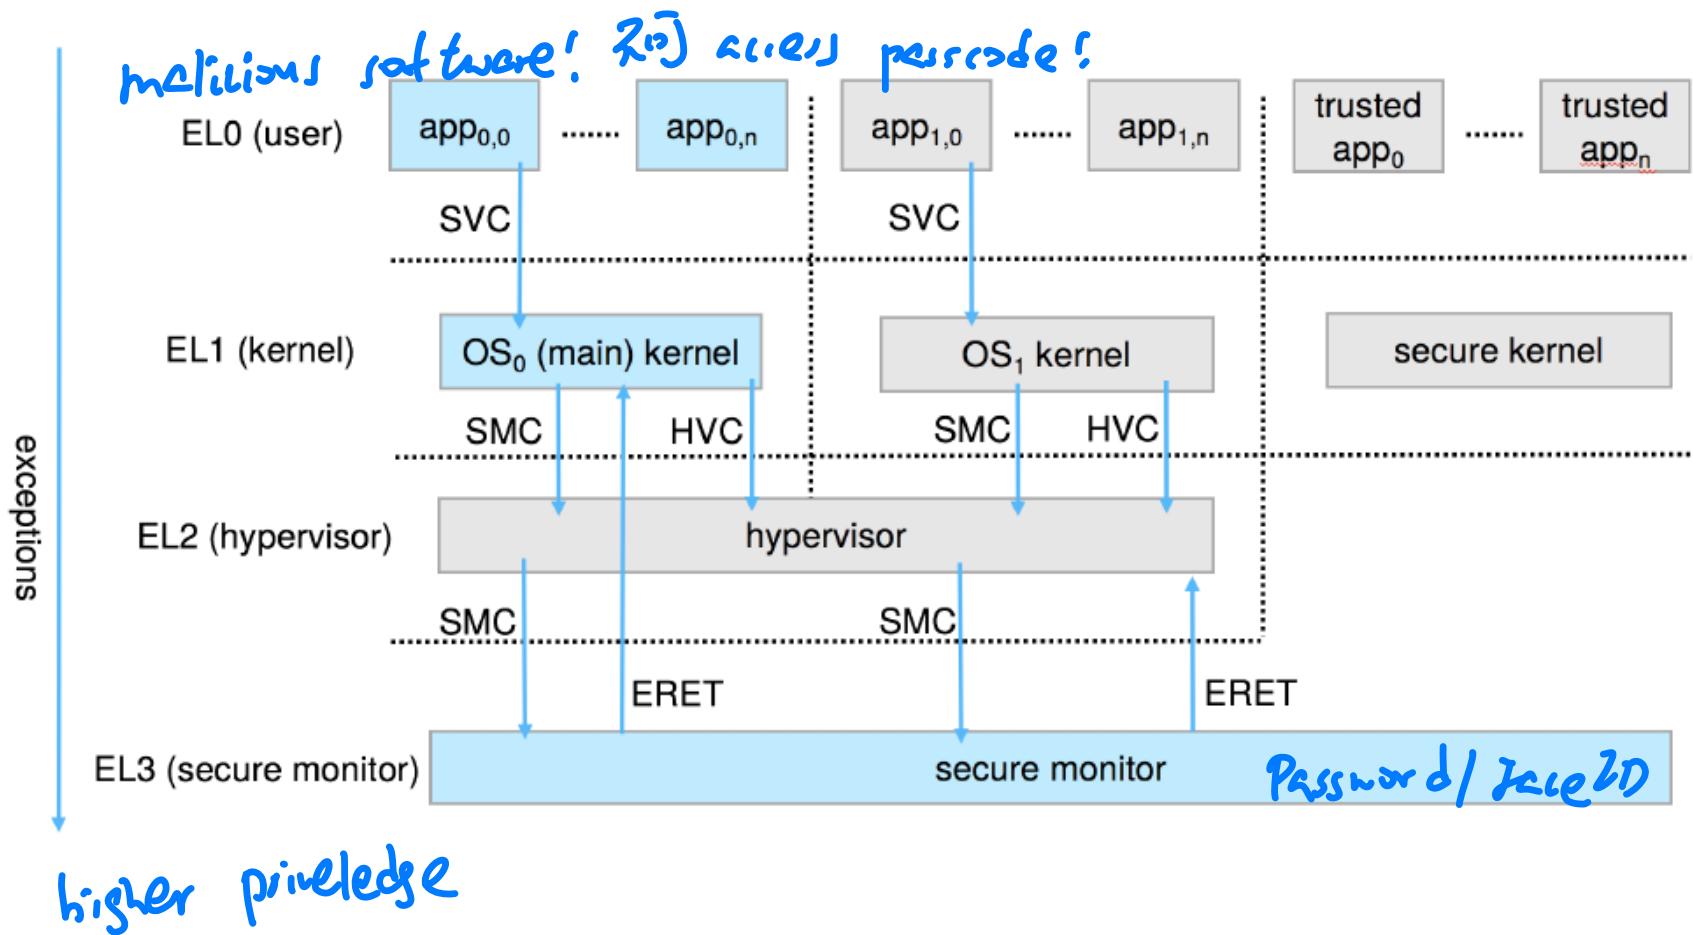
- Components ordered by amount of privilege and protected from each other
 - For example, the kernel is in one ring and user applications in another
 - This privilege separation requires hardware support
 - “Gates” used to transfer between rings, for example the **syscall** Intel instruction, also **traps** and **interrupts** → involve **space change!**
higher permission than kernel!
- **Hypervisors** (Intel) is introduced (another ring) - virtual machine managers, which create and run virtual machines, and have more capabilities than the kernels of the guest operating systems
- ARM processors added **TrustZone** or **TZ** ring to protect crypto functions with access (more privileged than kernel)
 - This most privileged execution environment has exclusive access to hardware-backed cryptographic features, such as the NFC Secure Element and an on-chip cryptographic key, that make handling passwords and sensitive information more secure.

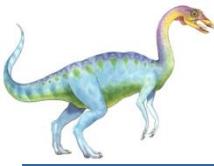
*APPLE : use this to store password, Face ID
actions apps' user space*





ARM CPU Architecture





← More general!

Domain of Protection

- Protection rings separate functions into different domains and order them hierarchically ← *some times* no need!
- **Domain** can be considered as a generalization of rings without a hierarchy
- A computer system can be treated as processes and objects
 - **Hardware objects** (such as CPU, memory, disk) and **software objects** (such as files, programs, semaphores)
- Process for example should only have access to objects it currently requires to complete its task – the need-to-know principle (policy) ← *at least knowledge*
- Implementation can be via process operating in a protection domain *Artificially define some domains!*
 - Protection domain specifies the set of resources a process may access
 - Each domain specifies set of objects and types of operations may be invoked on each object

*↑
domain I can access only files!*

Implement very well: flexible!

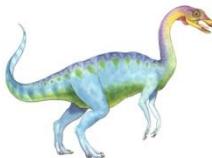




Domain of Protection (Cont.)

- The ability to execute an operation on an object is an **access right**
- A **domain** is a collection of access rights, each of which is an ordered pair **<object-name, rights-set>**
 - An example: if domain D has the access right **<file F, {read,write}>**, then a process executing in domain D can both read and write file F. It cannot, however, perform any other operation on that object.
file name & what is allowed
- Domains may share access rights *next slide has figure*
- Associations between processes and domains can be static if the set of resources available to the process is fixed throughout the process's lifetime, or can be **dynamic** *→ switch process in different domain!*
- If the association is dynamic, a mechanism is available to allow **domain switching**, enabling the process to switch from one domain to another during different stage of execution
protected



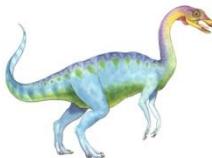


Domain of Protection (Cont.)

Domain can be realized in a variety of ways:

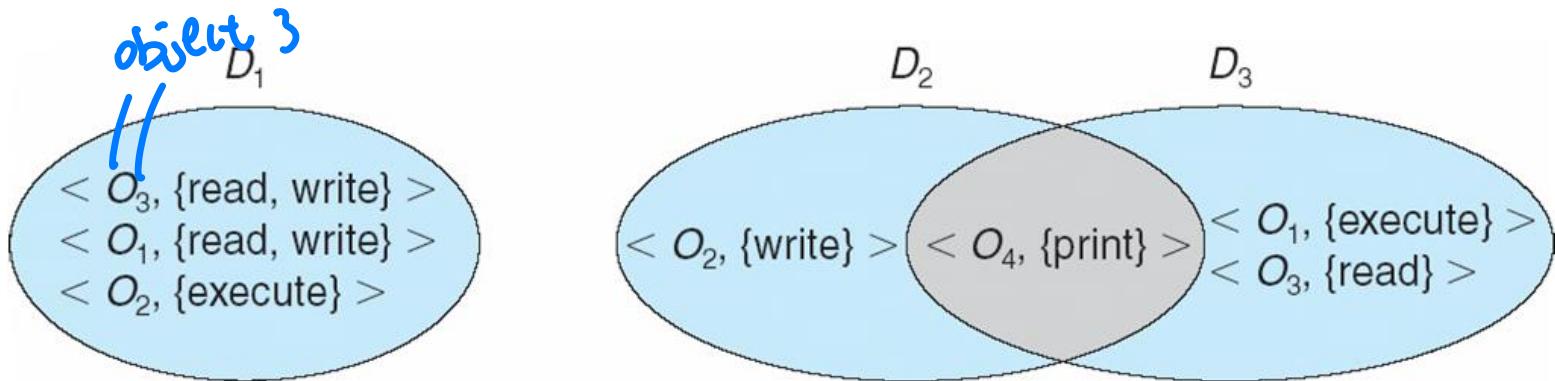
- Each **user** may be a domain - the set of objects that can be accessed depends on the identity of the user. Domain switching occurs when the user is changed
- Each **process** may be a domain - the set of objects that can be accessed depends on the identity of the process. Domain switching occurs when one process sends a message to another process and then waits for a response.
- Each **procedure** may be a domain - the set of objects that can be accessed corresponds to the local variables defined within the procedure. Domain switching occurs when a procedure call is made

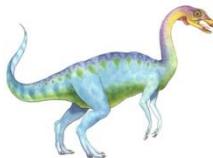




Domain Structure

- **Access-right** = <object-name, rights-set>
where rights-set is a subset of all valid operations that can be performed on the object
- **Domain** = set of access-rights *can be shared* *2之有1 domain mutually exclusive!*
- The access right <O4, {print}> shared by domains D2 and D3, thus, a process executing in either of these two domains can print object O4.





Access Matrix

- View protection as a matrix (**access matrix**) *Mechanism!*
- **Rows** represent domains, and **columns** represent objects
- **Access(i,j)** consists of a set of access rights - the set of operations that a process executing in Domain_i can invoke on Object_j

Domain Object owner of T₂

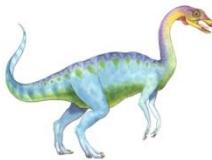
object domain	F ₁ F ₂ F ₃ printer				
	D ₁	read		read	
D ₂					print
D ₃			read	execute	
D ₄	read write			read write	

Annotations:
- A blue arrow points from the text "owner of T₂" to the column header "F₂".
- A blue box highlights the "read" entry in the cell (D₃, F₂). Handwritten text "Initialization" is written next to this box.
- A blue arrow points from the text "Initialization" to the highlighted cell.
- A blue arrow points from the text "D₁ can read object F₁" to the cell (D₁, F₁).

D₁ can
read object
F₁

Policy: define
which domain
have permission!

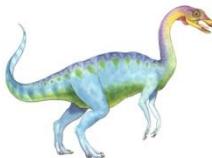




Use of Access Matrix

- The **access matrix** scheme provides the mechanism for specifying a variety of policies - mechanism and policy separation
 - The **mechanism** consists of implementing the access matrix and ensuring that the semantic properties hold.
 - To ensure that a process executing in domain D_i can access only those objects specified in row i .
 - The **policy** decisions specify which rights should be included in the (i,j) th entry, and determine the domain in which each process executes
 - If a process in Domain D_i tries to do “op” on object O_j , then “op” must be in the access matrix *ensure permission/ protection*
 - User who creates an object can define access column for that object
 - When a user creates a new object O_j , the column O_j is added to the access matrix with the appropriate initialization entries, as dictated by the creator. The user may decide to enter some rights in some entries in column j and other rights in other entries, as needed.





Use of Access Matrix (Cont.)

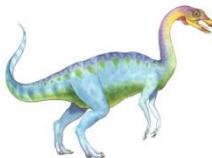
This can be expanded to dynamic protection

very flexible tools!

- Operations to add, delete access rights
- Special access rights:
 - **owner** of O_i - can add and remove any right in any entry in column
 - **copy** op from O_i to O_j (denoted by “*”) - only within the column (that is, for the object) *Example: Copy read only!*
 - **control** – D_i can modify D_j access rights – modify domain objects (a row)
↳ **transfer** – switch from domain D_i to D_j *can change*
 - **Copy** and **Owner** applicable to an object - change the entries in a column
 - **Control** applicable to domain object - change the entries in a row
 - New objects and new domains can be created dynamically and included in the access-matrix model
 - In a **dynamic** protection system, we may sometimes need to revoke access rights to objects shared by different users – **revocation** of access right

interesting complicated





Access Matrix of Figure A with Domains as Objects

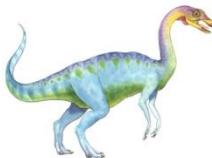
owner

control

object domain	F_1	F_2	F_3	laser printer	D_1	D_2	D_3	D_4
D_1	read		read			switch		
D_2				print			switch	switch
D_3		read	execute					
D_4	read write		read write		switch			

can allow to switch to D_2 !





Access Matrix with Copy Rights

object domain	F_1	F_2	F_3
D_1	execute		write*
D_2	execute	read*	execute
D_3	execute		

(a)

read can copy!
Only can copy in same column

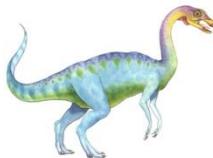
object domain	F_1	F_2	F_3
D_1	execute	read	write*
D_2	execute	read*	execute
D_3	execute	read	

(b)

non-copyable
after copied

copy To other column!





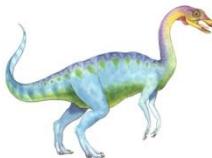
Access Matrix With *Owner* Rights

object domain	F_1	F_2	F_3
D_1	owner execute		write
D_2		read* owner	read* owner write
D_3	execute	write	write

object domain	F_1	F_2	F_3
D_1	owner execute		write
D_2		owner read* write*	read* owner write
D_3		write	write

(b)





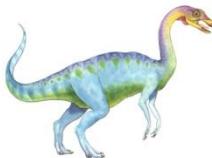
Modified Access Matrix of Figure B

其他不可switch完後 control!

object domain	F_1	F_2	F_3	laser printer	D_1	D_2	D_3	D_4
D_1	read		read			switch		
D_2				print			switch	switch control
D_3		read	execute					
D_4	write		write		switch			

add
control!
↓
 $D_2 \rightarrow D_4$
control D_4 !
modify
the role
of D_4





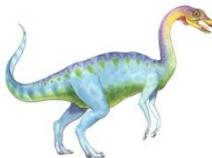
Implementation of Access Matrix

- 2D array
Complete access matrix 太 big files \Rightarrow memory $\uparrow\uparrow$
- In general, the access matrix is sparse; that is, most of the entries will be empty
 - **Option 1 – Global Table**
Object-oriented: 1D array!, 1维空间,
empty \Rightarrow X store!
 - Store ordered triples \langle domain, object, rights-set \rangle in table
 - A requested operation M on object O_j within domain $D_i \rightarrow$ search table for $\langle D_i, O_j, R_k \rangle$ with $M \in R_k$
 - But the table could be large \rightarrow might not fit in main memory, requires additional I/O – virtual memory techniques are often used
 - Difficult to group objects - For example, if everyone can read a particular object, this object must have a separate entry in every domain.

if matrix is full \Rightarrow table is also large

difficult group F_i together!
resource
multiplexing!





Implementation of Access Matrix (Cont.)

- Each column = **Access-control list** for one object
Defines who can perform what operation

Domain 1 = Read, Write
Domain 2 = Read
Domain 3 = Read

like a key!

- Each row = **Capability List** (like a key)
For each domain, what operations allowed on what objects

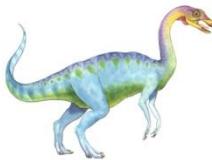
Object F1 – Read
Object F4 – Read, Write, Execute
Object F5 – Read, Write, Delete, Copy

For object

- make different domain, different access right

ALL apply in domain, we domain to group!





Implementation of Access Matrix (Cont.)

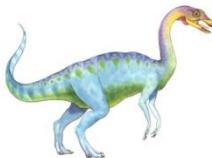
□ Option 2 – **Access lists for objects**

ACL

- Each column implemented as an access list for one object
- Resulting per-object list consists of ordered pairs <domain, rights-set> defining all domains with non-empty set of access rights for the object
- Obviously, the empty entries can be discarded.
- This can be easily extended to define **default** set of access rights -> If $M \in$ default set, also allow access (for all domains)

can file default to be read!
Default permission!





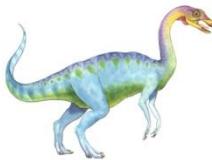
Implementation of Access Matrix (Cont.)

□ Option 3 – Capability list for domains

domain I can access . . .

- Instead of object-based, list is domain-based
- A capability list for domain is a list of objects together with operations allowed on them
→ whether object can do . . .
- An object represented by its name or address, called a capability
- To execute operation M on object O_j , a process requests operation M, specifying the capability (or pointer) for object O_j as a parameter
want to pass object to
- Possession of capability means access is allowed
capability is what should be protected
- Capability list associated with a domain, but never directly accessible by a process executing in that domain
capability list itself is a protected object, maintained by OS and accessed by users only indirectly
- This avoids the possibility of capability list modification by users
- If all capabilities are secure, the object they protect is also secure against unauthorized access





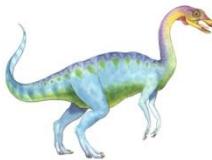
Implementation of Access Matrix (Cont.)

□ Option 4 – Lock-key

- Compromise between access lists and capability lists
- Each object has list of unique bit patterns, called **locks**
- Each domain has list of unique bit patterns called **keys**
- Process in a domain can only access object if domain has key that matches one of the locks of the object *key → open lock*
- As with capability lists, the list of keys for a domain must be managed by the operating system on behalf of the domain. *Complicated, not in modern OS!*
- Users are not allowed to examine or modify the list of keys (or locks) directly.

long array, ALL capability list is enough!





Comparison of Implementations

: pros
: cons

Choosing a technique for implementing an access matrix involves various trade-offs.

- five triple! - ve*
- Global table is simple, but large, lack of grouping of objects or domains
 - Access lists *good!* correspond directly to the needs of users *can define the access list*
 - An access list on an object is specified when a user creates the object
 - Determining set of access rights for each domain is difficult - every access to the object must be checked, *requiring a search of the access list. overhead! :)*
 - Capability lists *useful for localizing information for a given process*
 - But revocation capabilities can be inefficient
 - Lock-key *new! secret the table!* can be effective and flexible depending on the length of the keys
 - Keys can be passed freely from domain to domain, easy revocation
 - Most systems use combination of access lists and capabilities

→
lists of keys
lots of overhead



End of Chapter 17

